Hadoop Distributed File System

Big Data Management

UNIVERSITAT POLITÈCNICA
DE CATALUNYA
BARCELONATECH

DTIM
www.essi.upc.edu/dtim

Knowledge objectives

1. Recognize the need of persistent storage
2. Enumerate the design goals of GFS
3. Explain the structural components of HDFS
4. Name three file formats in HDFS and explain their differences
5. Recognize the importance of choosing the file format depending on workload
6. Explain the actions of the coordinator node in front of chunkserver failure
7. Explain a mechanism to avoid overloading the master node in HDFS
8. Explain how data is partitioned and replicated in HDFS
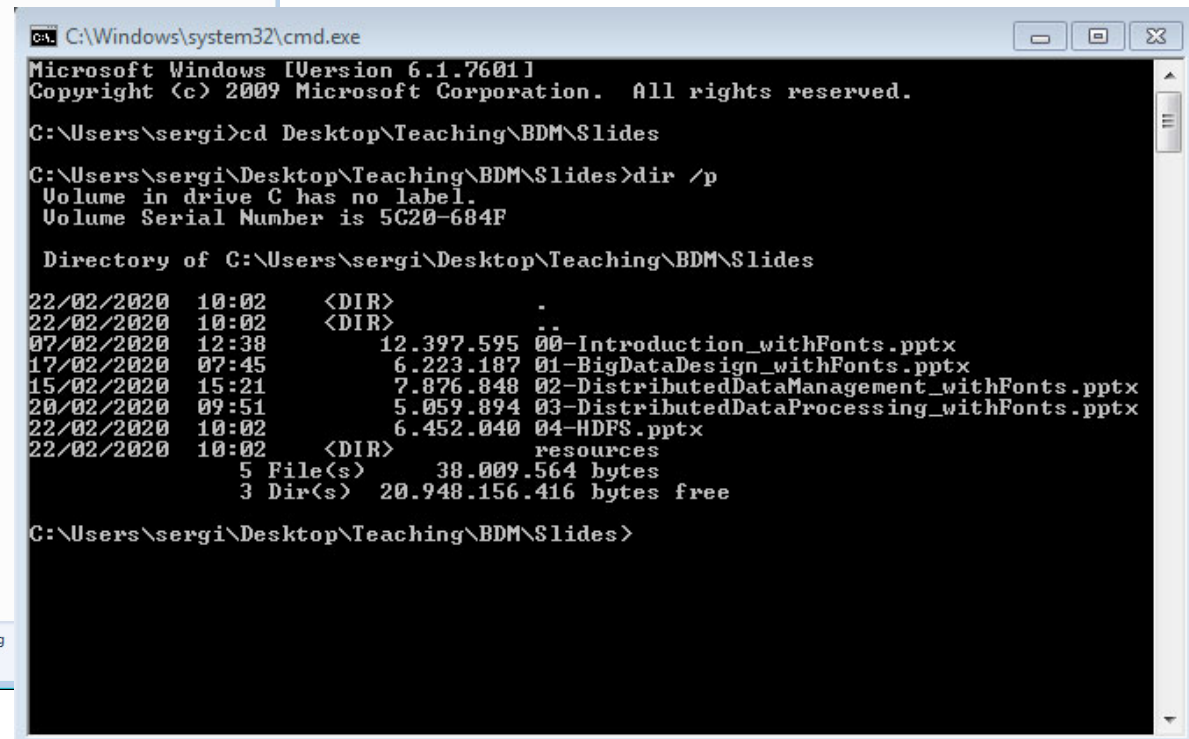9. Recognize the relevance of sequential read

Understanding objectives

1. Choose the format for an HDFS file based on heuristics
2. Estimate the data retrieved by scan, projection and selection operations in SequenceFile, Avro and Parquet
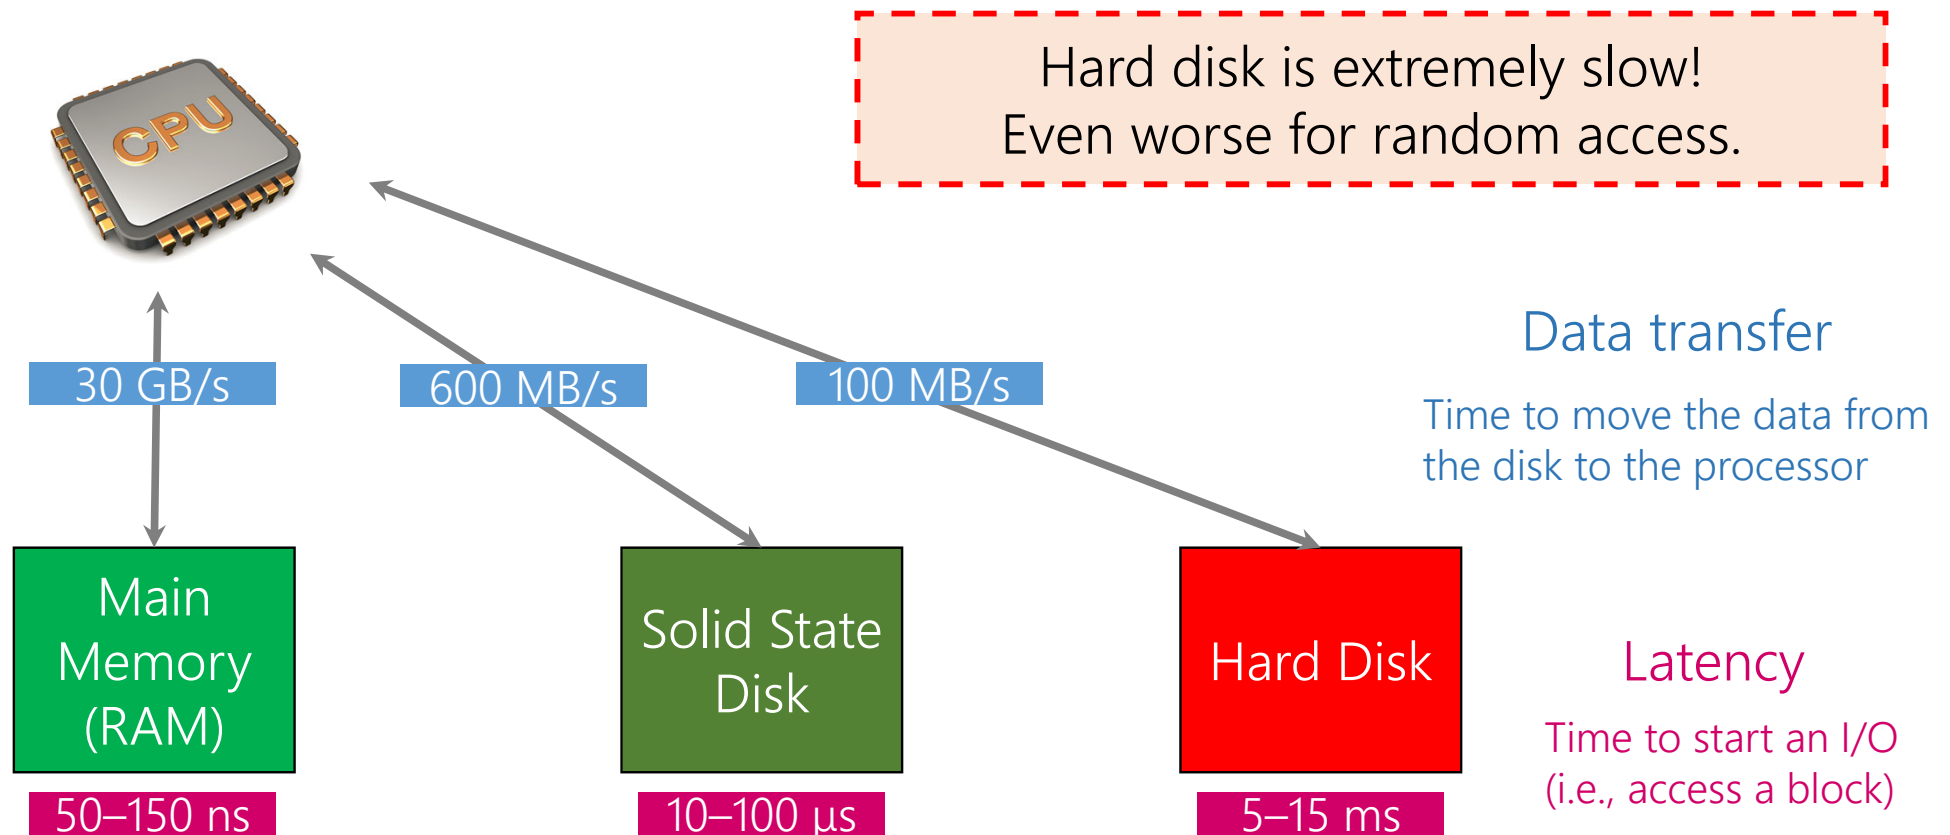
(Distributed) File Systems

Google File System

What is a file system?

Time to bring data (approximations)

Hard disk is extremely slow!
Even worse for random access.

30 GB/s   600 MB/s   100 MB/s

Main Memory (RAM)   Solid State Disk   Hard Disk

50–150 ns   10–100 μs   5–15 ms

Data transfer

Time to move the data from the disk to the processor

Latency

Time to start an I/O (i.e., access a block)

A. Hogan

Reasons to keep using HDDs

Main
Memory

Faster ✔

✔ Cheaper
✔ Persistant

Hard
Drive Disk

Functionalities provided by a FS

- Creates a hierarchical structure of data
- Splits and stores data into files and blocks
- Maintains directories/files metadata
  - Size, date of creation, permissions, …
- Provides interfaces to read/write/delete

Directory

File

Block

Distributed File Systems

- Same requirements, different setting
    1. Files are huge for traditional standards
    2. Most files are updated by appending data rather than overwriting
        - Write once and read many times
    3. Component failures are the norm rather than the exception

- Google File System (GFS)
    - The first large-scale distributed file system
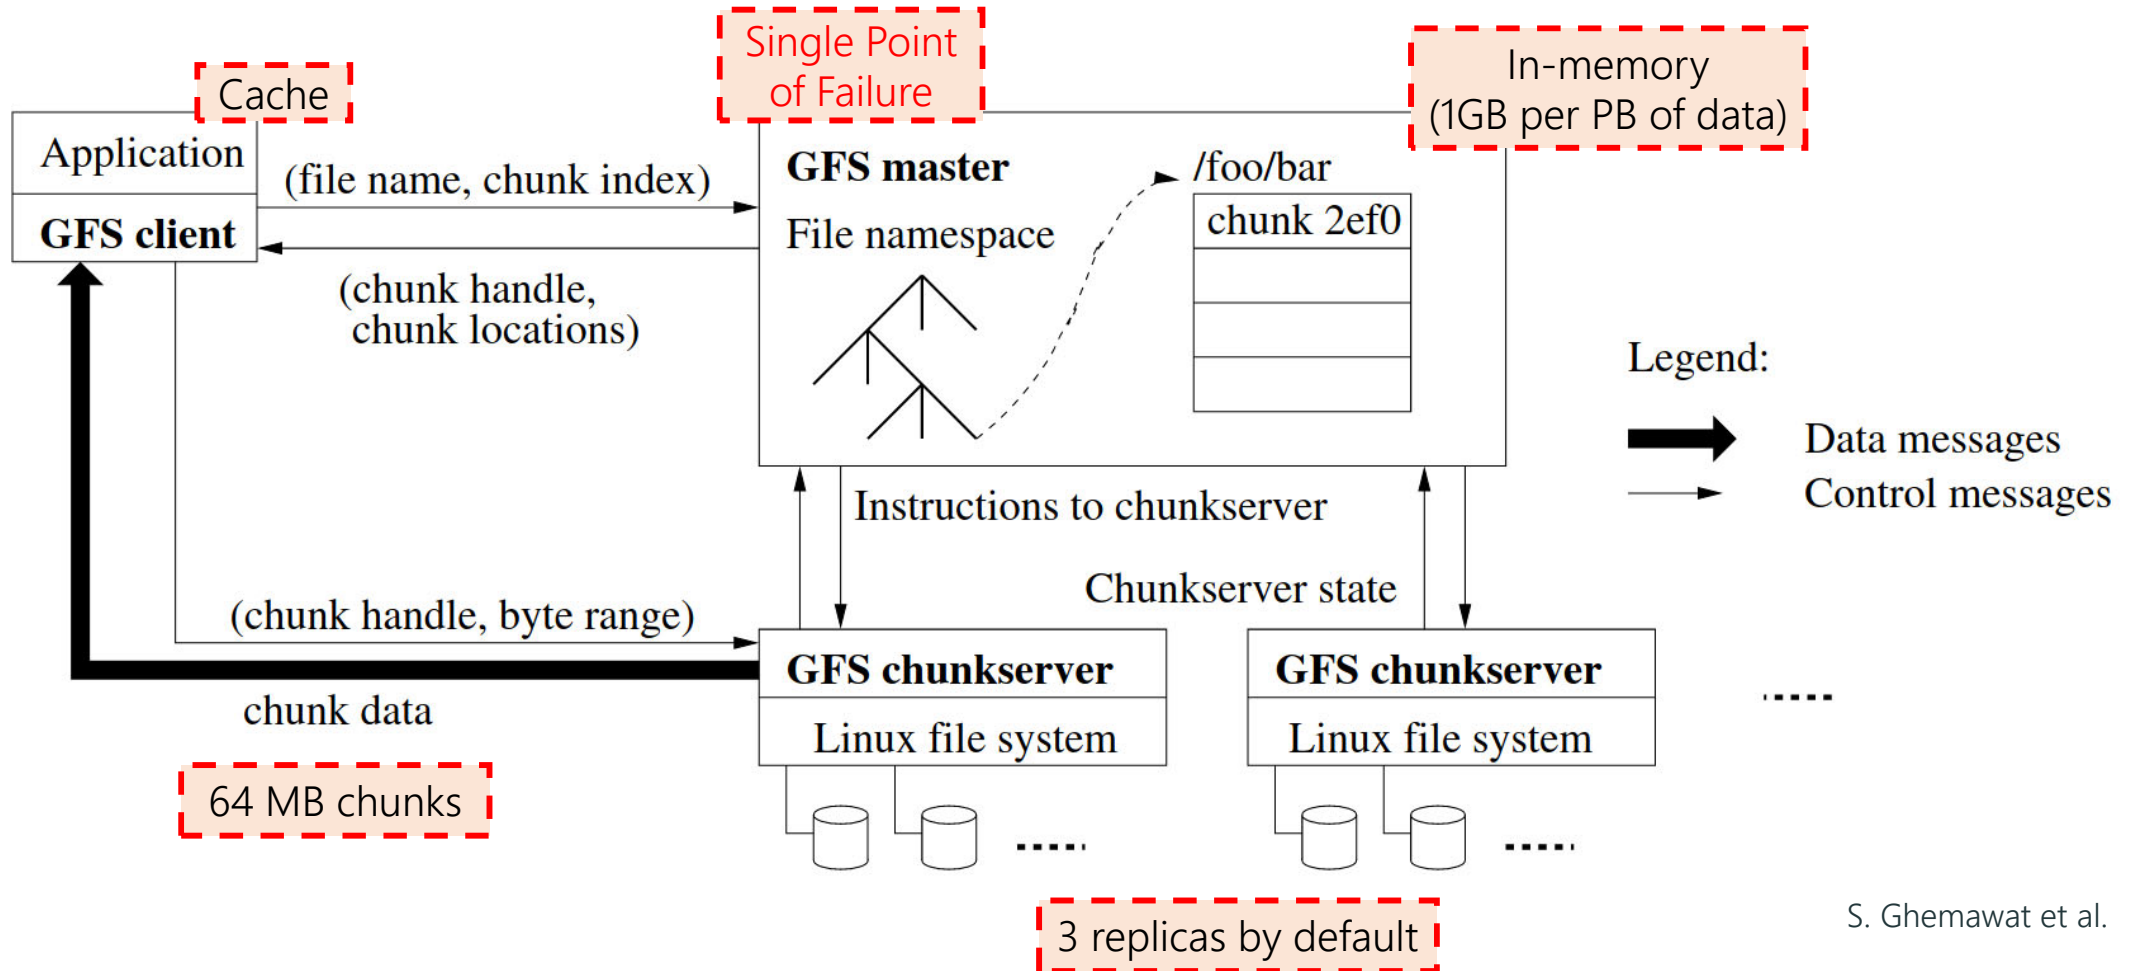   - Capacity of a GFS cluster

Capacity	Nodes	Clients	File
10 PB	10.000	100.000	100.000.000

Design goals of GFS

- Efficient management of files
  - Optimized for very large files (GBs to TBs)
- Efficiently append data to the end of files
  - Allow concurrency
- Tolerance to failures
  - Clusters are composed of many inexpensive machines that fail often
    - Failure probability (2-3/1.000 per day)
- Sequential scans optimized
  - Overcome the high latency of HDDs (5-15ms) compared to main memory (50-150ns)

UNIVERSITAT POLITÈCNICA
DE CATALUNYA
BARCELONATECH

DTIM
www.essi.upc.edu/dtim

GFS Architecture

HDFS Architecture

Cache

Single Point of Failure

In-memory (1GB per PB of data)

Application

GFS client

(file name, chunk index)

(chunk handle, chunk locations)

GFS master

File namespace

/foo/bar

chunk 2ef0

Instructions to chunkserver

Chunkserver state

Legend:

Data messages

Control messages

(chunk handle, byte range)

chunk data

GFS chunkserver

Linux file system

GFS chunkserver

Linux file system

64 MB chunks

3 replicas by default

S. Ghemawat et al.

Other features

- Rebalance
  - Avoids skewness in the distribution of chunks
- Deletion
  - Moves a file to the trash (hidden)
    - Kept for 6h
  - *expunge* to force the trash to be emptied
- Management of stale replicas
  - Coordinator maintains versioning information about all chunks

Data Design

Challenge I

Storage layouts

"Jack of all trades, master of none"

- Different workloads require different layout
  - Horizontal
    - For scan-based workloads
  - Vertical
    - For projection-based workloads (reads a subset of columns)
  - Hybrid
    - For projection- and predicate-based workloads (reads a subset of columns or rows)

UNIVERSITAT POLITÈCNICA
DE CATALUNYA
BARCELONATECH
UPC

DTIM
www.essi.upc.edu/dtim

Horizontal layout – Sequence File

- Records of binary key-value pairs

- Compression
  - Uncompressed
  - Record-compressed
  - Block-compressed
    - "block" is the compression unit – a block of records (not a chunk)
      - 1 MB default

Horizontal layout - Avro

- Binary encoding of (compressed) rows
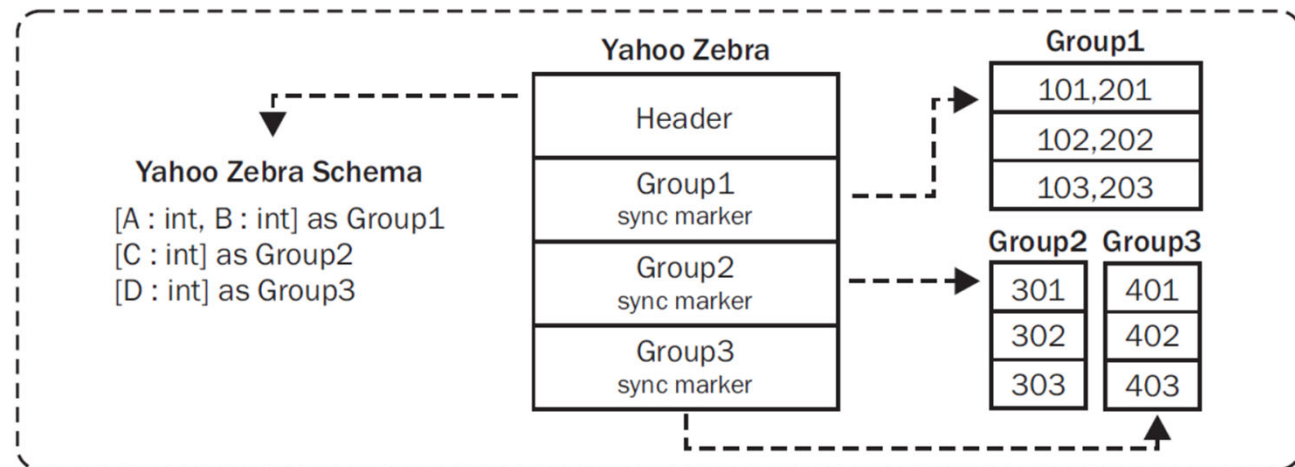- The header contains a schema encoded in JSON
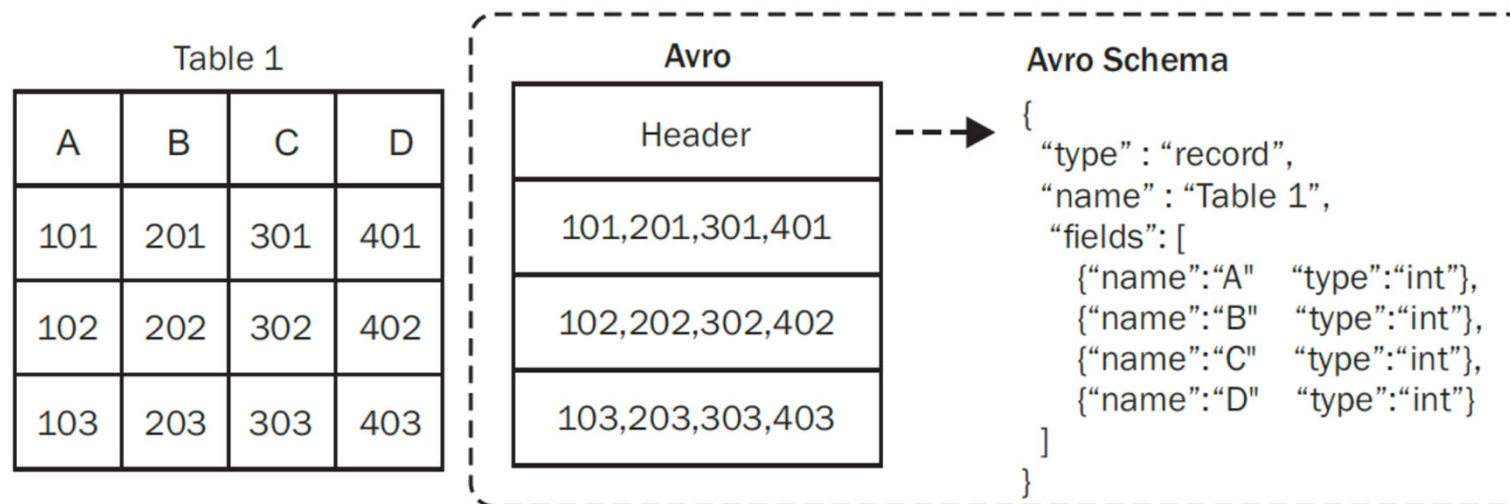
Vertical layout - Zebra

- The header contains the definition of groups
  - Each group contains a set of columns
 - Widely benefits from compression
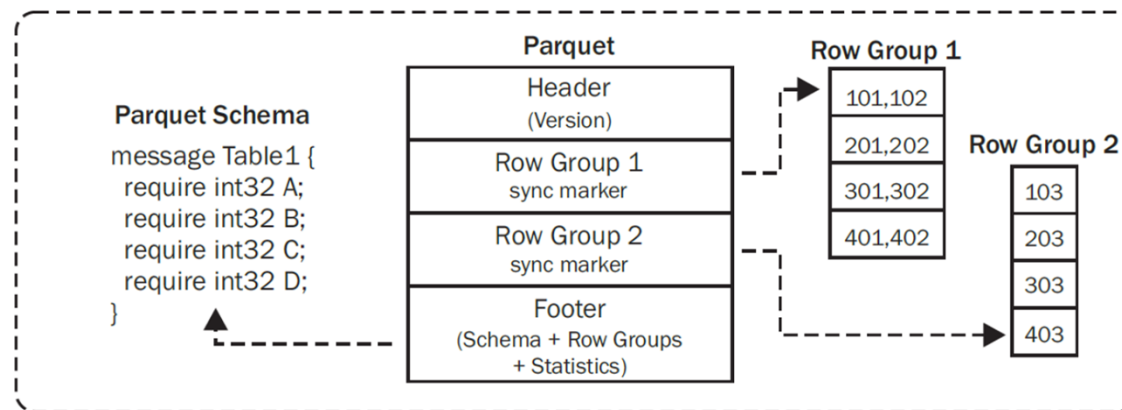- Not really used in practice

Hybrid layout - Parquet

- Row groups (RG) - horizontal partitions
  - Data vertically partitioned within RGs
- Statistics per row group (aid filtering)
  - E.g., min-max

Parquet encoding definitions

- Plain

- Valid for any data type
  - Dictionary encoding
  - Run-length encoding with dictionary

- Specific for some kind of data types
  - Delta encoding
  - Delta byte array / Delta-length byte array
  - Byte Stream Split

Comparison of data formats

Features	Horizontal		Vertical	Hybrid	
	Sequence Files	Avro	Yahoo Zebra	ORC	Parquet
Schema	No	Yes	Yes	Yes	Yes
Column Pruning	No	No	Yes	Yes	Yes
Predicate Pushdown	No	No	No	Yes	Yes
Indexing Information	No	No	No	Yes	Yes
Statistics Information	No	No	No	Yes	Yes
Nested Records	No	No	Yes	Yes	Yes

Rule-based choice (heuristic)

Given a flow represented as DAG(V, E)

- SequenceFile
  - size(getCol(v)) = 2
- Parquet
  - ∃e ∈ O(v), getType(e) = {AggregationOps}
  - ∃e ∈ O(v), getCol(getOP (e)) ⊂ getCol(v)
- Avro
  - ∀e ∈ O(v), getCol(getOP (e)) = getCol(v)
  - ∃e ∈ O(v), getType(e) ∈ {Join, CartesianProduct, GroupALL, Distinct}

Cost-based choice

- Helps in choosing the right storage layout based on the workloads
- Costs to Consider
  - Write cost
  - Read Cost
    - Scan Operation
    - Projection Operation
    - Selection Operation
- Costs ignored
  - Block compression
  - Dictionary encoding (in Parquet)

General formulas for size estimation

$$Size(Layout) = Size(Header_{Layout})$$
$$+Size(Body_{Layout})$$
$$+Size(FooterLayout)$$

$$UsedChunks(Layout) = \frac{Size(Layout)}{Size(chunk)}$$
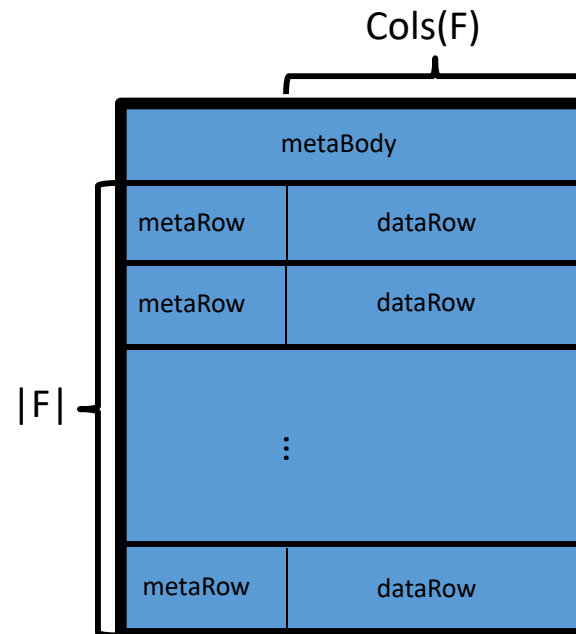
$$Seeks(Layout) = \lceil UsedChunks(Layout) \rceil$$
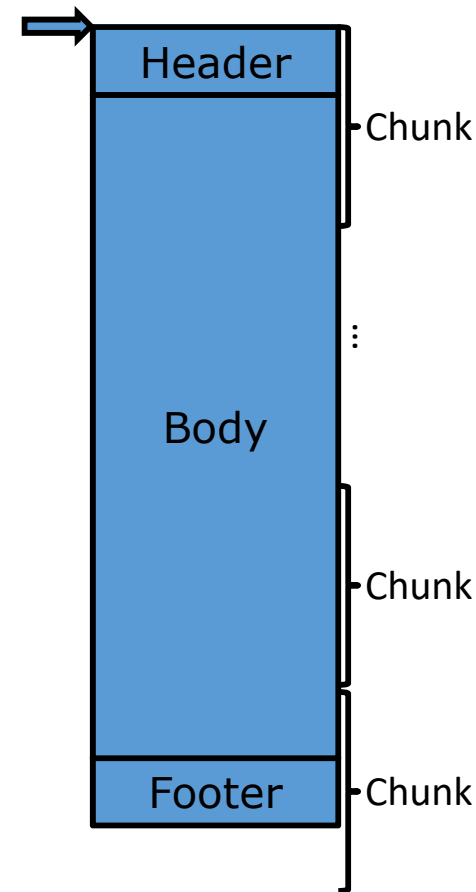
"Layout" can be either "Horizontal", "Vertical" or "Hybrid"

Horizontal layout size estimation

$$Size(dataRow) = Cols(F) * Size(cell)$$

$$Size(Body_{Horizontal}) = Size(metaBody) + |F| * (Size(metaRow) + Size(dataRow))$$
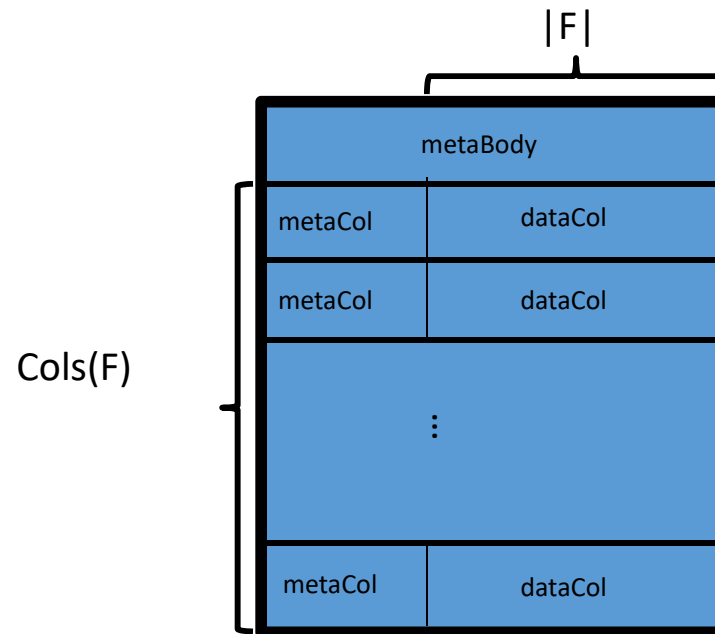

"F" is the content of the file

Vertical layout size estimation

$$Size(dataCol) = |F| * Size(cell)$$

$$Size(Body_{Vertical}) = Size(metaBody) + Cols(F) * (Size(metaCol) + Size(dataCol))$$
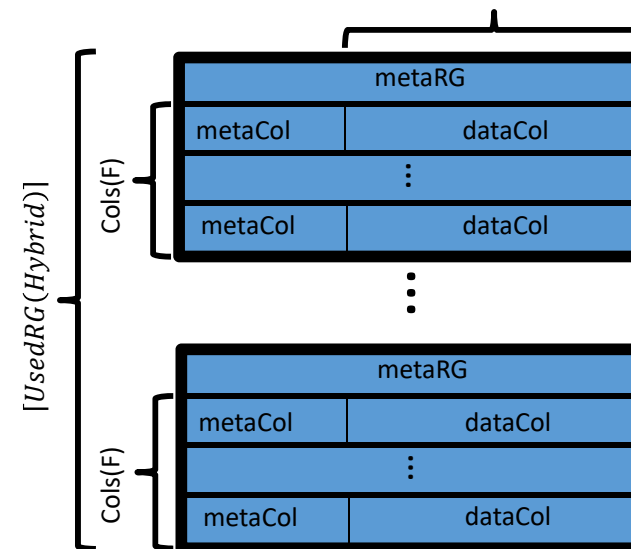

"F" is the content of the file

Hybrid layout size estimation

$$UsedRG(Hybrid) = \frac{Cols(F) * |F| * Size(cell)}{(Size(RG) - Size(metaRG) - Cols(F) * Size(metaCol))}$$

$$Size(Body_{Hybrid}) = \lceil UsedRG(Hybrid) \rceil * (Size(metaRG) + Cols(F) * Size(metaCol))$$
$$+ Cols(F) * |F| * Size(cell)$$

$$UsedRows(RG) = \frac{|F|}{UsedRG(Hybrid)}$$


"F" is the content of the file

General formulas for cost estimation

$$Cost(Write_{Layout}) = UsedChunks(Layout) * W_{WriteTransfer}$$
$$+Seeks(Layout)*(1-W_{WriteTransfer})$$

$$Size(Scan_{Layout}) = Size(Layout)$$
$$+(\lceil UsedChunks(Layout) \rceil * Size(Meta_{Layout}))$$

$$UsedChunks(Scan_{Layout}) = \frac{Size(Scan_{Layout})}{Size(chunk)}$$

$$Cost(Scan_{Layout}) = UsedChunks(Scan_{Layout}) * W_{ReadTransfer}$$
$$+Seeks(Layout)*(1-W_{ReadTransfer})$$

"Layout" can be either "Horizontal", "Vertical" or "Hybrid"

Projection in hybrid layouts

$$Size(projCols) = proj(F) * UsedRows(RG) * Size(cell)$$

$$Size(Project_{Hybrid}) = Size(Header_{Hybrid}) + Size(Footer_{Hybrid})$$
$$+ \lceil UsedRG(Hybrid) \rceil * (Size(metaRG) + proj(F) * Size(metaCol))$$
$$+ UsedRG(Hybrid) * Size(projCols)$$

$$Cost(Project_{Hybrid}) = UsedChunks(Project_{Hybrid}) * W_{ReadTransfer}$$
$$+ Seeks(Hybrid) * (1 - W_{ReadTransfer})$$

$$UsedRows(RG) = \frac{|F|}{UsedRG(Hybrid)}$$
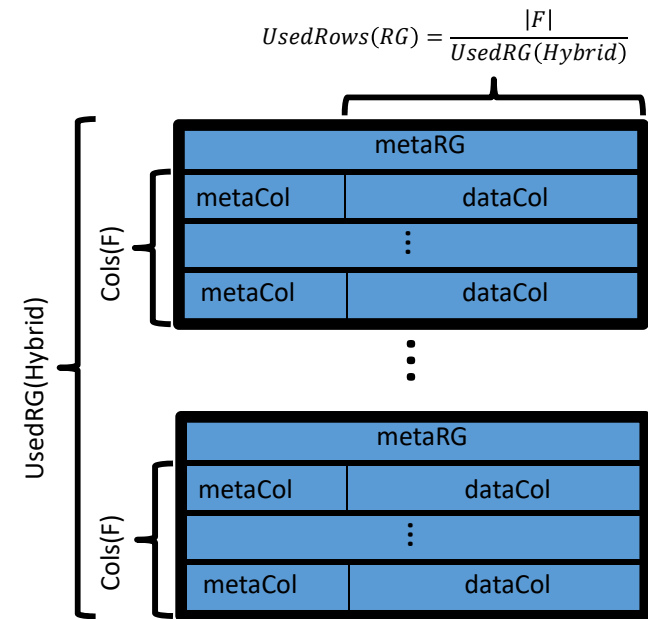
"proj(F)" is the number of projected columns

Probability of retrieving a RowGroup

- Probability of a row fulfilling P
  SF

- Probability of a row NOT fulfilling P
  1-SF

- Probability of none of the rows in a RowGroup fulfilling P
  $(1-SF) \cdot (1-SF) \cdot \ldots \cdot (1-SF) = (1-SF)^{UsedRows(RG)}$

- Probability of some row in a RowGroup fulfilling P
  $1-(1-SF)^{UsedRows(RG)}$

Selection in hybrid layouts

$$P(RGSelected) = 1 - (1 - SF)^{UsedRows(RG)}$$

$$Size(RowsSelected) = \left\lceil \frac{SF * |F|}{UsedRows(RG)} \right\rceil \big(Size(metaRG) + Cols(F) * Size(metaCol)\big)$$
$$+ SF * |F| * Cols(F) * Size(cell)))$$

$$UsedRG(Select_{Hybrid}) = \begin{cases} if\ unsorted:\ P(RGSelected) * UsedRG(Hybrid) \\ if\ sorted: \left\lceil \dfrac{Size(RowsSelected)}{Size(RG)} \right\rceil \end{cases}$$

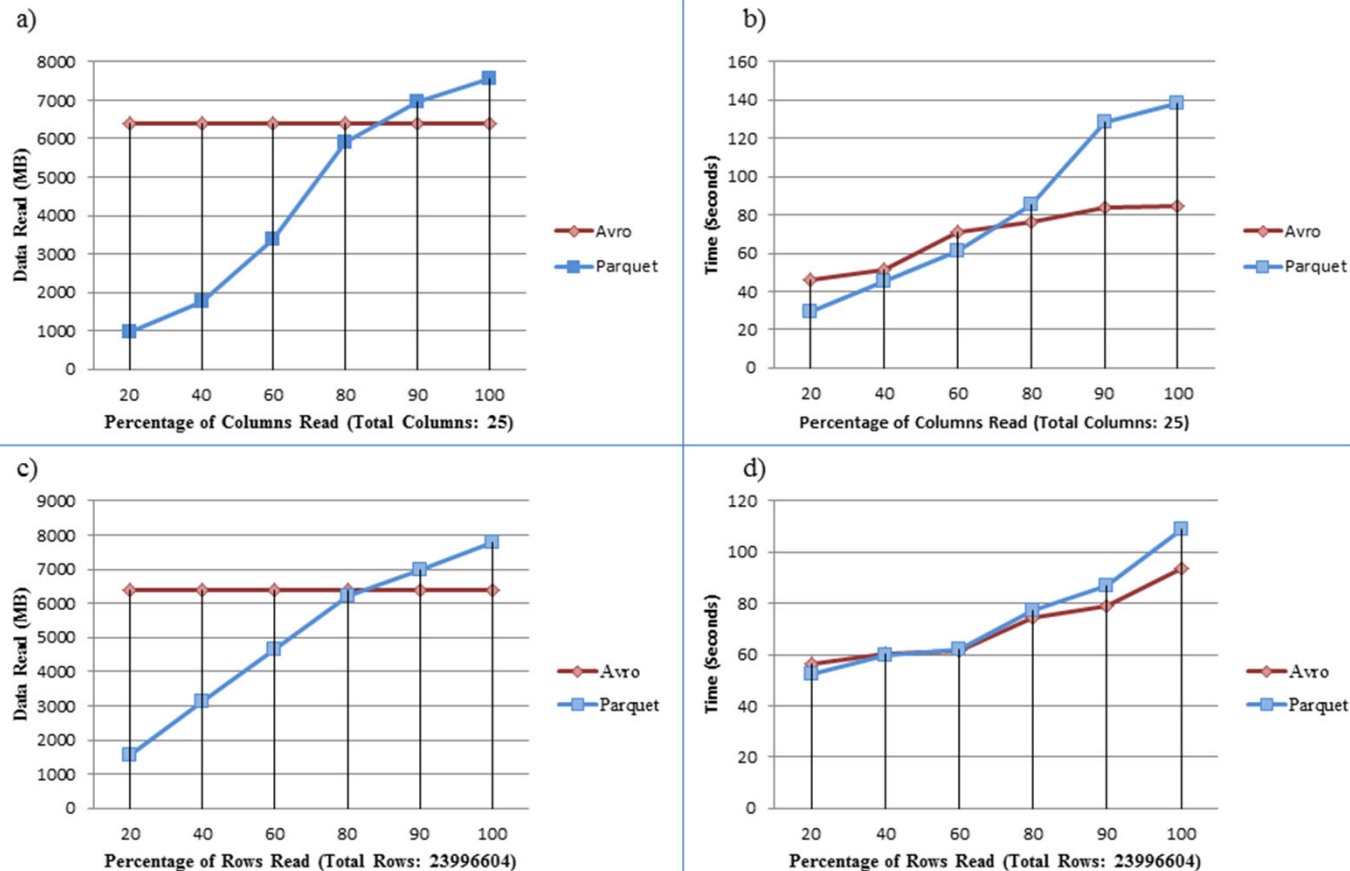$$UsedRows(RG) = \frac{|F|}{UsedRG(Hybrid)}$$

$$Size(Select_{Hybrid}) = Size\big(Header_{Hybrid}\big) + Size\big(Footer_{Hybrid}\big)$$
$$+ UsedRG(Select_{Hybrid}) * Size(RG)$$

$$Cost(Select_{Hybrid}) = UsedChunks(Select_{Hybrid}) * W_{ReadTransfer}$$
$$+ Seeks(Select_{Hybrid}) * (1 - W_{ReadTransfer})$$
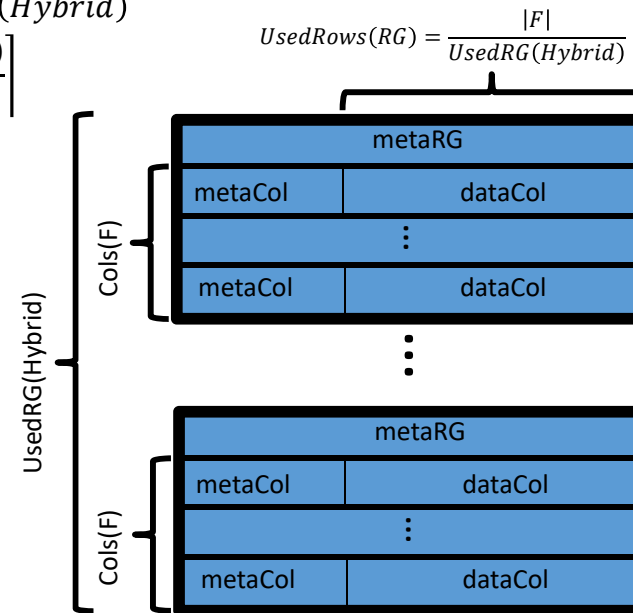
"F" is the content of the file

Comparison of selection and projection

Fault tolerance

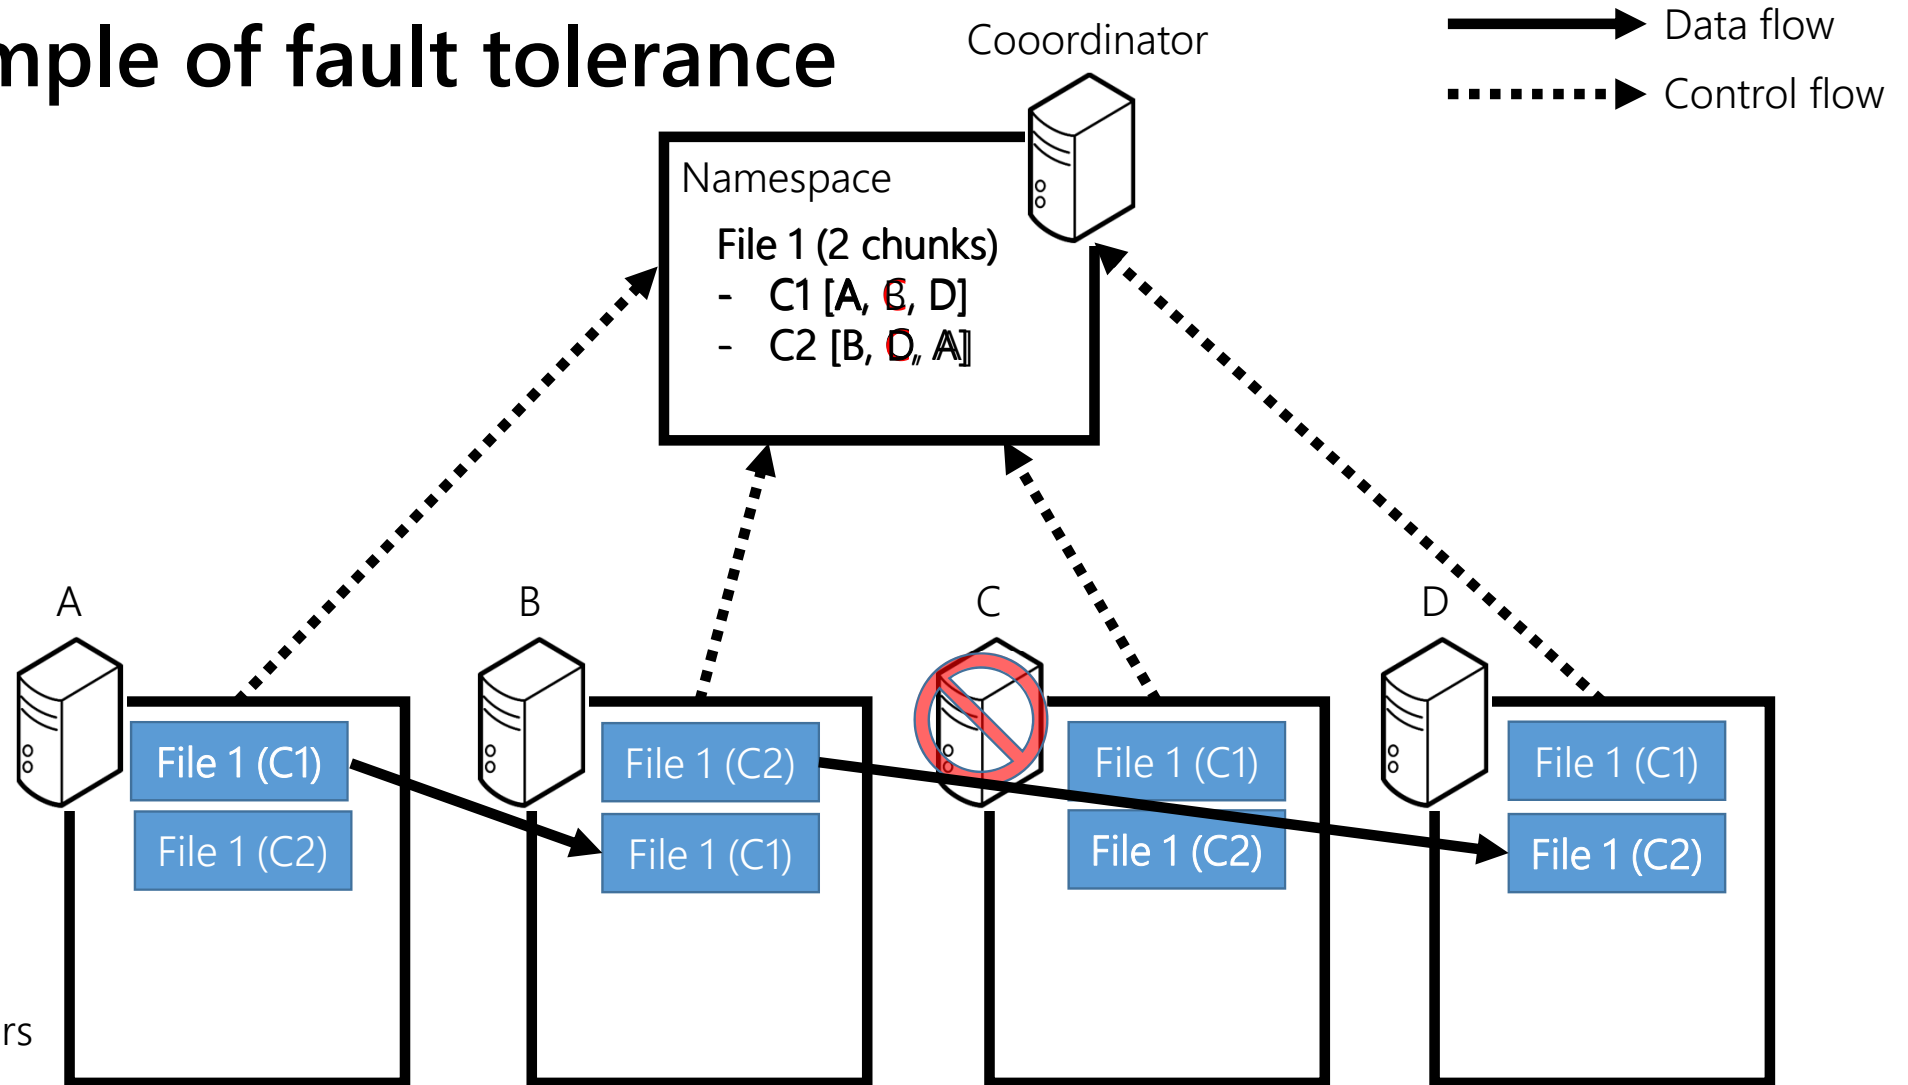- Managed from the coordinator
  - It expects to receive every 3 seconds a *heartbeat* message from chunkservers
- If a chunkserver does not send a heartbeat for 60 seconds a fault is declared
- Corrective actions
  - Update the namespace
  - Copy one of the replicas to a new chunkserver
    - Potentially electing a new primary replica

Example of fault tolerance

Catalog Management

Challenge II

Client caching

Cache miss

1. The client sends a READ command to the coordinator
2. The coordinator requests chunkservers to send the chunks to the client
   - Ranked according to the closeness in the network
3. The list is cached in the client
   - Not a complete view of all partitions

Cache hit

1. The client reads the cache and requests the chunkservers to send the chunks

> Avoid coordinator bottleneck
> +
> One communication step is saved

Example of client cache miss

Coordinator

Data flow
Control flow

File 1 (2 chunks)
- C1 [?, C, ?]
- C2 [?, ?, A]

File 1 (C1)

File 1 (C2)

Namespace

File 1 (2 chunks)
- C1 [A, C, D]
- C2 [B, C, A]

A   B   C   D

File 1 (C1)   File 1 (C2)   File 1 (C1)   File 1 (C1)

File 1 (C2)   File 1 (C2)
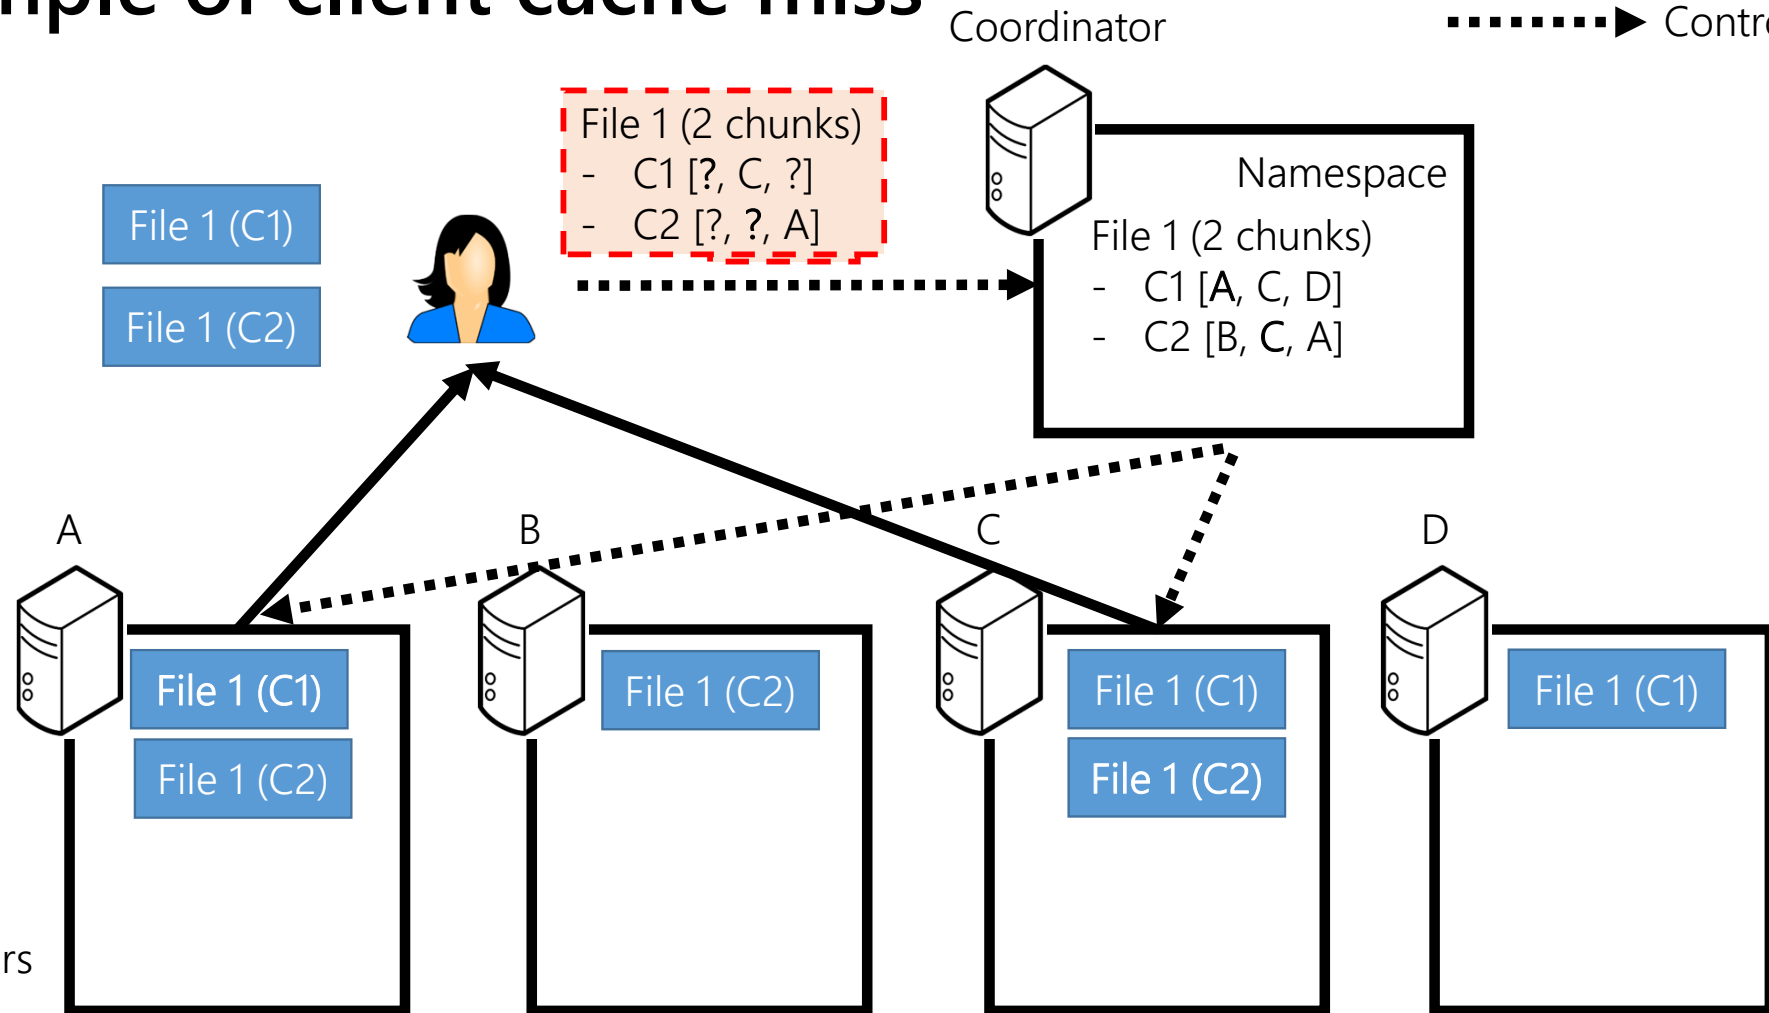
ChunkServers

Example of client cache hit

Data flow

Control flow

Coordinator

File 1 (2 chunks)
- C1 [?, C, ?]
- C2 [?, ?, A]

Namespace

File 1 (2 chunks)
- C1 [A, C, D]
- C2 [B, C, A]

File 1 (C1)

File 1 (C2)

File 1 (C1)?

File 1 (C2)?

A

B

C

D

File 1 (C1)

File 1 (C2)

File 1 (C1)

File 1 (C1)

File 1 (C2)

File 1 (C2)

ChunkServers

UNIVERSITAT POLITÈCNICA
DE CATALUNYA
BARCELONATECH
UPC

DTIM
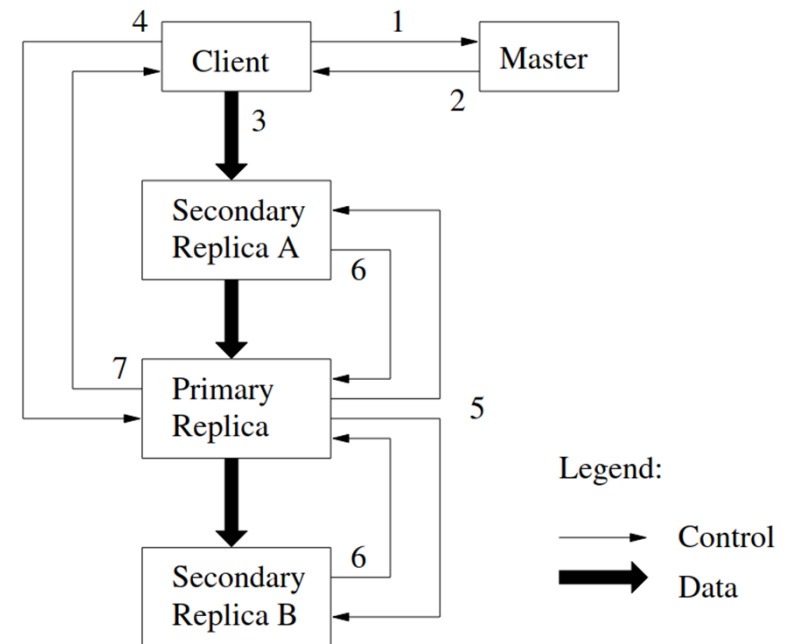www.essi.upc.edu/dtim

38
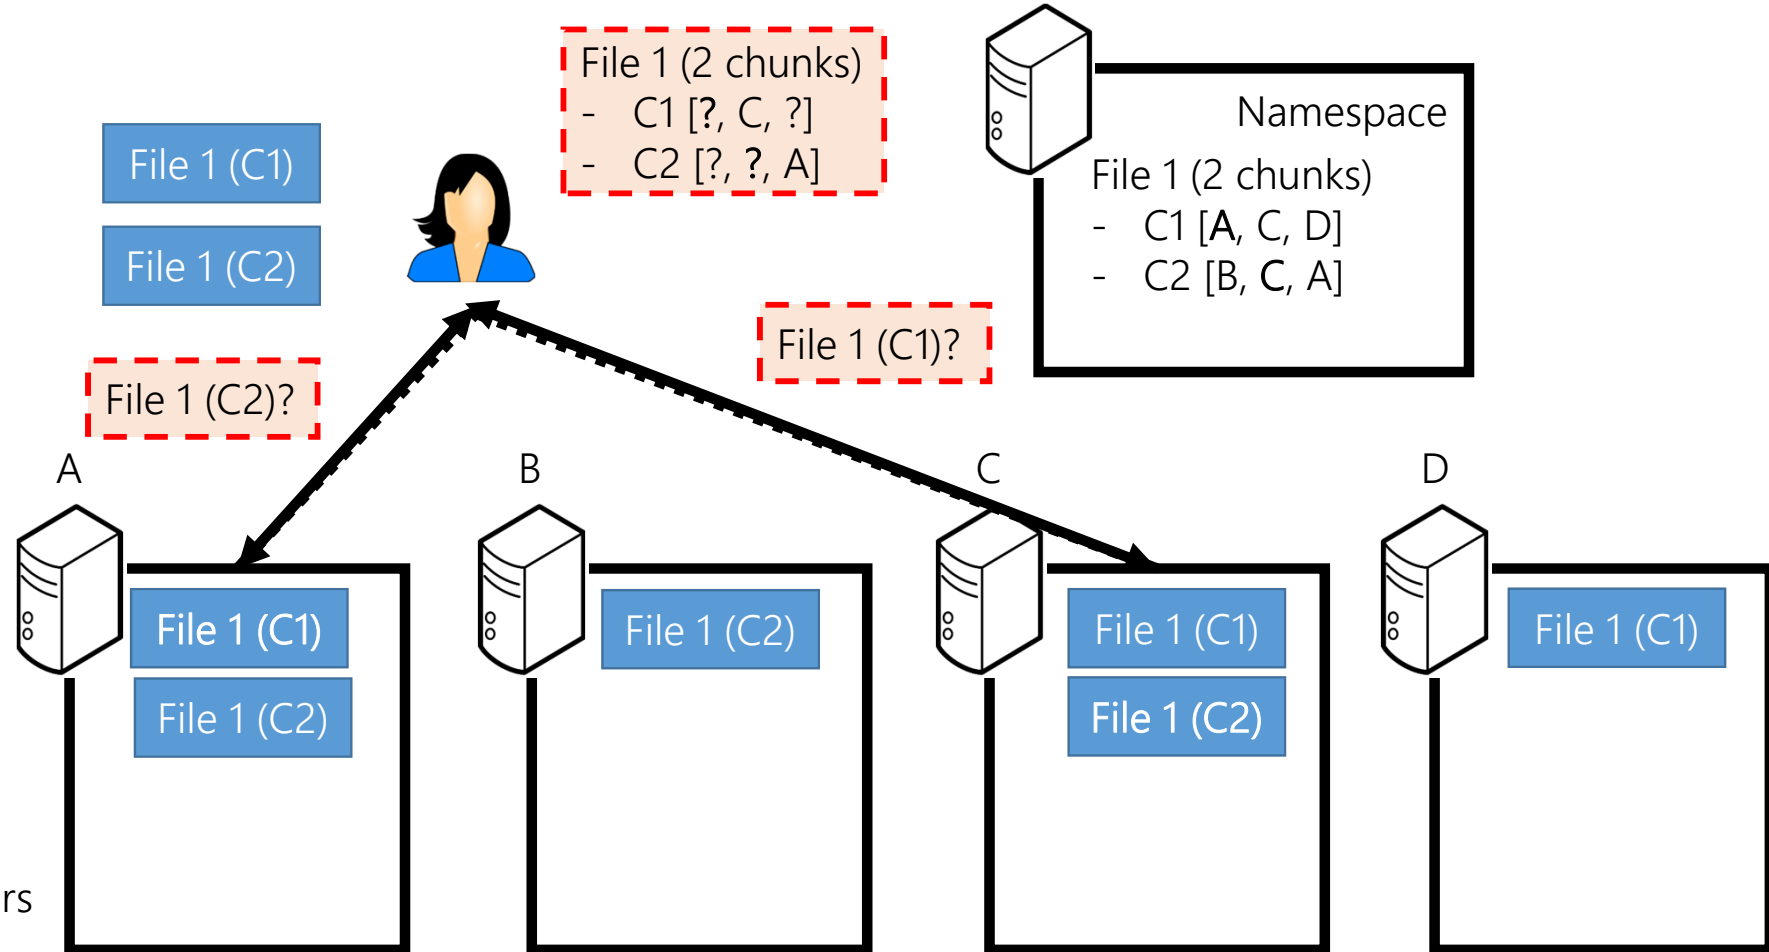
# Transaction Management

Challenge III

Writing replicas

1. The client requests the list of the replicas of a file
2. Coordinator returns metadata
3. Client sends a chunk to the closest chunkserver in the network
   - This chunk is pipelined to the other chunkservers in the order defined by the master (leases)
4. Client sends WRITE order to primary replica
5. Primary replica sends WRITE order to secondary replicas
6. Secondaries confirm to primary the change
7. Primary confirms to the client

S. Ghemawat et al.

Example of writing replicas

Data flow
Control flow

Coordinator

Namespace
File 1 (2 chunks)
- C1 [A, C, D]
- C2 [B, C, A]

File 1 (100 MB)

A   B   C   D

File 1 (C1)   File 1 (C2)   File 1 (C1)   File 1 (C1)

File 1 (C2)   File 1 (C2)

ChunkServers

UNIVERSITAT POLITÈCNICA
DE CATALUNYA
BARCELONATECH
UPC

DTIM
www.essi.upc.edu/dtim

Query processing

Challenge IV

HDDs costs

Read\Write Head

Actuator

Seek Time

Platters Spin

Rotational Latency

The amount of time it take for the platters to spin the data under the head (measured in RPM)

Seek Time

Time it takes for the ReadWrite head (mechanical arm) to move between sectors on the disk

Transfer Time

Time it takes for requests to get from the system to the disk (depends on the block size, e.g., 4KB)

Sequential vs. Random access

Sequential Access

Disk blocks: | 1 | 2 | 3 | 4 | 5 | 6 | ... |

Random Access

Disk blocks: | 1 | 2 | 3 | 4 | 5 | 6 | ... |

Cost of accessing data (approximations)

- Random Access
  - Requires indexing structures
  - Ignores data locality

  Cost = n*(seek+rotation+transfer)

- Sequential reads
  - Option to maximize the effective read ratio
    - Depends on DB design
  - Enables pre-fetching

  Cost = seek+rotation+n*transfer

seek	~12ms
rotation	~3ms
transfer (8KB)	~0.03ms

CPU

L1 cache

L2 cache

L3 cache

Page cache

Disk buffer

Closing

Summary

- GFS architecture and components
- GFS main operations
  - Fault tolerance
  - Writing files and maintenance of replicas
  - Reading files
- HDFS file formats
  - Horizontal
  - Vertical
  - Hybrid

References

- S. Ghemawat et al. *The Google File System*. OSDI'03
- K. V. Shvachko. *HDFS scalability: the limits to growth*. 2010
- S. Abiteboul et al. *Web data management*. Cambridge University Press, 2011
- A. Jindal et al. *Trojan data layouts: right shoes for a running elephant.* SOCC, 2011
- F. Färber et al. *SAP HANA database: data management for modern business applications.* SIGMOD, 2011
- V. Raman et al. *DB2 with BLU Acceleration: So Much More than Just a Column Store*. VLDB, 2013
- F. Munir et al. *ResilientStore: A Heuristic-Based Data Format Selector for Intermediate Results*. MEDI, 2016
- D. Abadi, et al. *Column-stores vs. row-stores: how different are they really?* SIGMOD Conference, 2008
- M. Stonebraker et al. C-Store: A Column-oriented DBMS. VLDB, 2005
- G. Copeland and S. Khoshafian. *A Decomposition Storage Model*. SIGMOD Conference, 1985
- *F. Munir*: PhD Thesis (UPC)
- *A*. Hogan: Procesado de Datos Masivos (U. Chile)

UNIVERSITAT POLITÈCNICA DE CATALUNYA BARCELONATECH
UPC

DTIM
www.essi.upc.edu/dtim