

Key-Value Stores

Big Data Management

Knowledge objectives

1. Give the definition of the BigTable data model
2. Explain what a map structure is
3. Explain the difference between a KeyValue and a Wide-Column store
4. Enumerate the main schema elements of Hbase
5. Explain the main operations available in HBase
6. Enumerate the main functional components of HBase
7. Explain the role of the different functional components in Hbase
8. Explain the tree structure of data in HBase
9. Explain the 3 basic algorithms of Hbase
10. Explain the main components and behaviour of an LSM-tree
11. Compare a distributed tree against a hash structure of data
12. Justify the need of dynamic hashing
13. Explain the structure of HBase catalog
14. Explain the mistake compensation mechanism of the cache in HBase client
15. Enumerate the ACID guarantees provided by HBase
16. Explain the execution flow of an HBase query both at global and local levels

Understanding objectives

1. Given few queries, define the best logical structure of a table considering its physical implications in terms of performance
2. Given the data in two leafs of a Log-Structured Merge-tree, merge them
3. Given the current structure of a Linear Hash, modify it according to insertions potentially adding buckets
4. Given the current structure of a Consistent Hash, modify it in case of adding a bucket
5. Calculate the number of round trips needed in case of mistake compensation of the tree metadata

Application objectives

1. Use HBase shell to create a table and access it
2. Use HBase API to create a table and access it

Key-Value

BigTable

BigTable Data Model

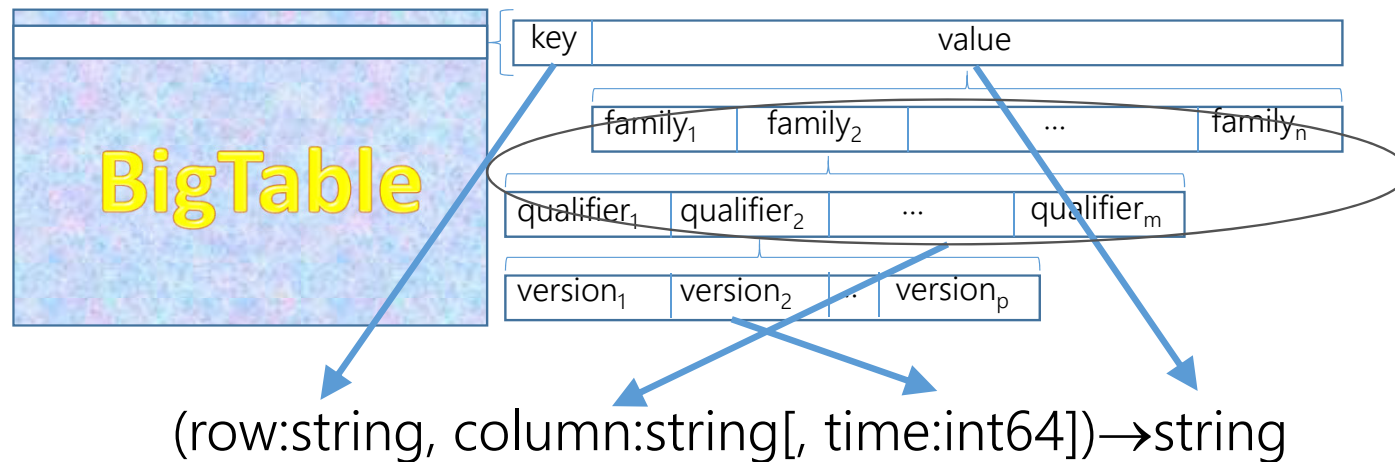
"A Bigtable is a sparse, distributed, persistent, multi-dimensional, sorted map."

F. Chang et al.

- Sparse: most elements are unknown
- Distributed: cluster parallelism
- Persistent: disk storage (HDFS)
- Multi-dimensional: values with columns
- Sorted: sorting lexicographically by primary key
- Map: lookup by primary key

BigTable schema elements

- Stores tables (collections) and rows (instances)
 - Data is indexed using row and column names (arbitrary strings)
- Treats data as uninterpreted strings (without data types)
- Each cell of a BigTable can contain multiple versions of the same data
 - Stores different versions of the same values in the rows
 - Each version is identified by a timestamp
 - Timestamps can be explicitly or automatically assigned




Key-Value

- Key-value stores
 - Entries in form of key-values
 - One key maps only to one value
 - Query on key only
 - Schemaless



- Bigtable (Wide-column key-value stores)
 - Entries in form of key-values
 - But now values are split in wide-columns
 - Typically query on key
 - May have some support for values
 - Schemaless within a wide-column



HBase

- Apache project
 - Based on Google's Bigtable
- Designed to meet the following requirements
 - Access specific data out of petabytes of data
 - It must support
 - Key search
 - Range search
 - High throughput file scans
 - It must support single row transactions

HBase Architecture

HBase shell


- ALTER <tablename>, <columnfamilyparam>
- COUNT <tablename>
- CREATE TABLE <tablename>
- DESCRIBE <tablename>
- DELETE <tablename>, <rowkey>[, <columns>]
- DISABLE <tablename>
- DROP <tablename>
- ENABLE <tablename>
- EXIT
- EXISTS <tablename>
- GET <tablename>, <rowkey>[, <columns>]
- LIST
- PUT <tablename>, <rowkey>, <columnid>, <value>[, <timestamp>]
- SCAN <tablename>[, <columns>]
- STATUS [{summary|simple|detailed}]
- SHUTDOWN

<https://learnhbase.wordpress.com/2013/03/02/hbase-shell-commands>

Functional components (I)

- Zookeeper - Quorum of servers that stores HBase system config info
- HMaster - Coordinates splitting of regions/rows across nodes
 - Controls distribution of HFile chunks
- Region Servers (HRegionServer)
 - Serves HBase client requests
 - Manage stores containing all column families of the region
 - Logs changes
 - Guarantees “atomic” updates to one column family
- HFiles - Consist of large (e.g., 128MB) chunks
 - 3 copies of one chunk for availability (default)
- HDFS - Stores all data including columns and logs
 - NameNode holds all metadata including namespace
 - DataNodes store chunks of a file
 - HBase uses two HDFS file types
 - HFile: regular data files (holds column data)
 - HLog: region’s log file (allows flush/fsync for small append-style writes)


Functional components (II)



StoreFiles

- When the MemStore is full (128MB), data are flushed to HDFS
- A StoreFile is generated
 - Format HFile
- An HFile stores data in blocks
 - Hbase blocks (64 KB)

Storefile
(HFile format)



HBase processes

a) Flush

- On memory structure reaching threshold
- Take memory content and stores it into an SSTable
- Generates different disk versions of the same record

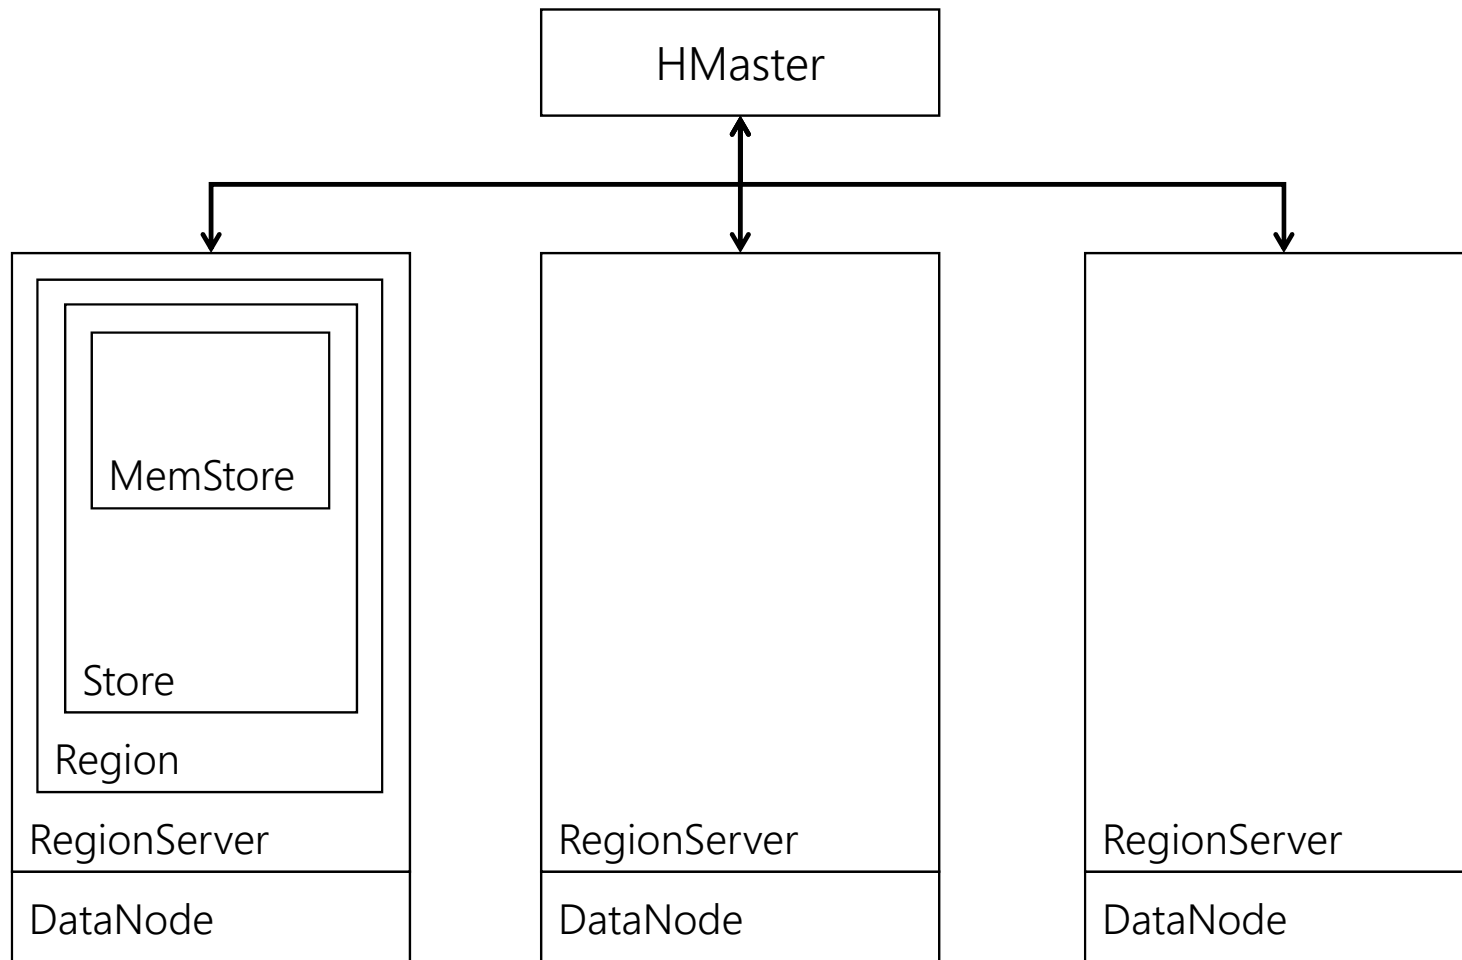
b) Minor compactation

- Run regularly in the background
- Merge a given number of equal size SSTables into one
- Does not remove all record versions (only some)

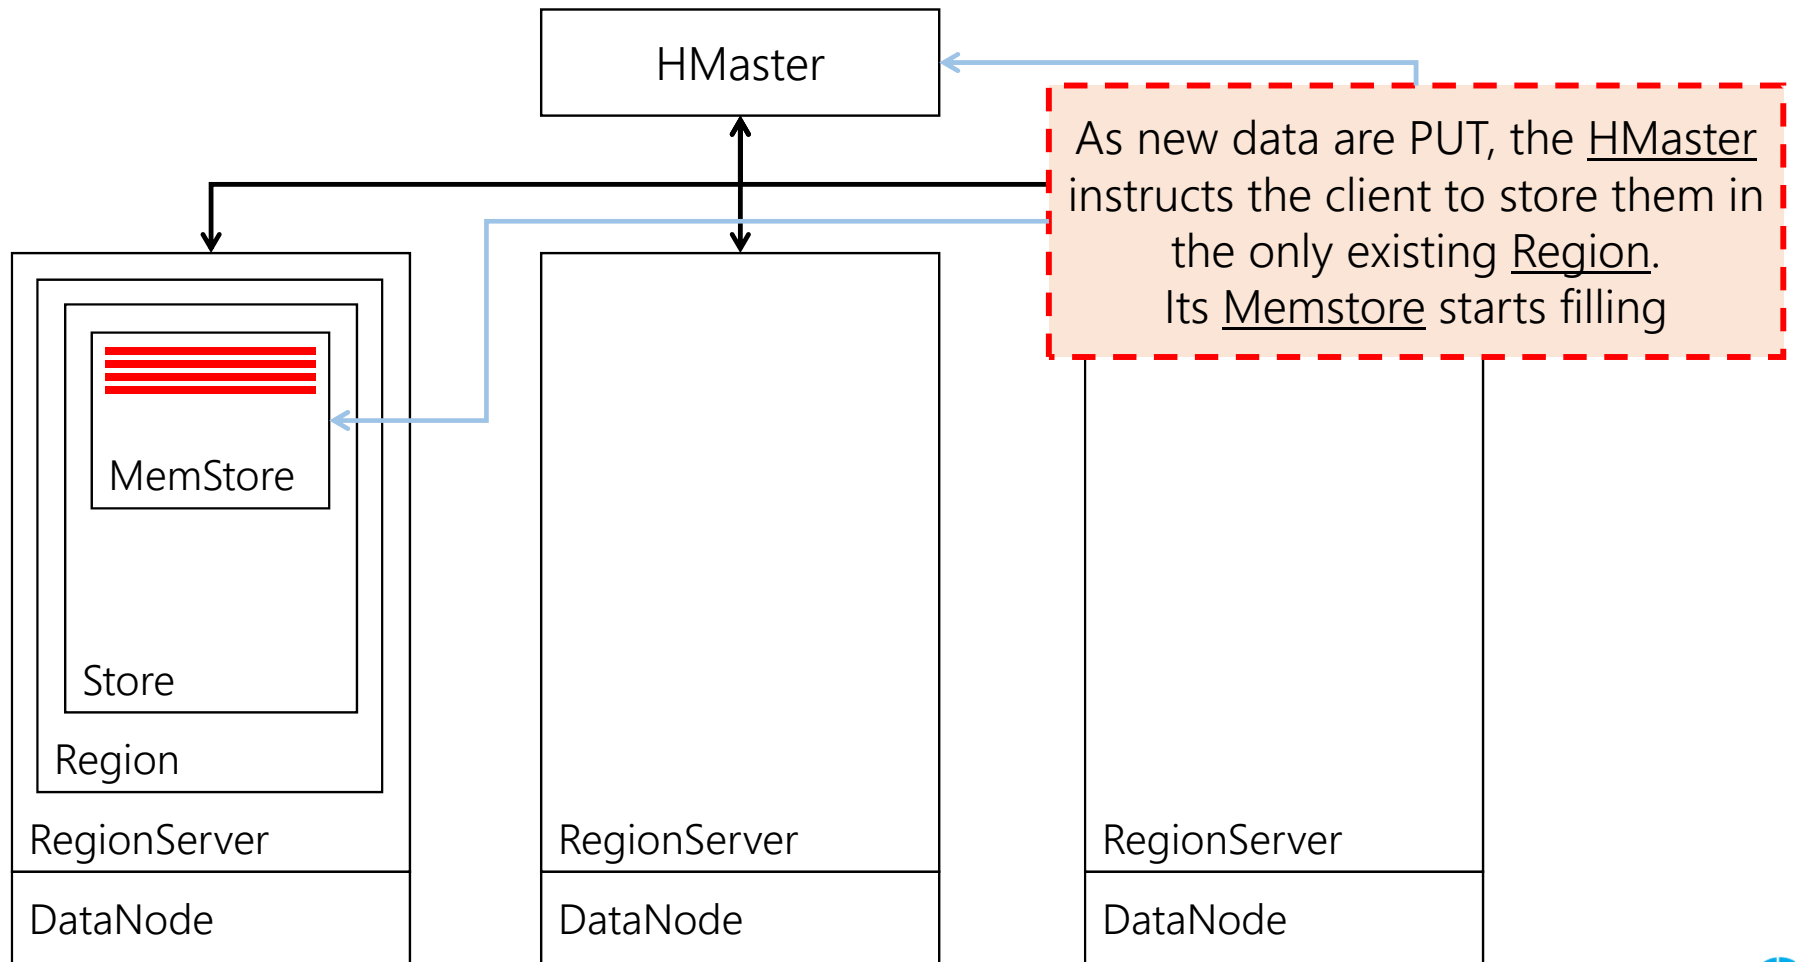
c) Major compactation

- Triggered manually
- Merge all SSTables
- Leaves one single SSTable
- All versions of a record are merged into one

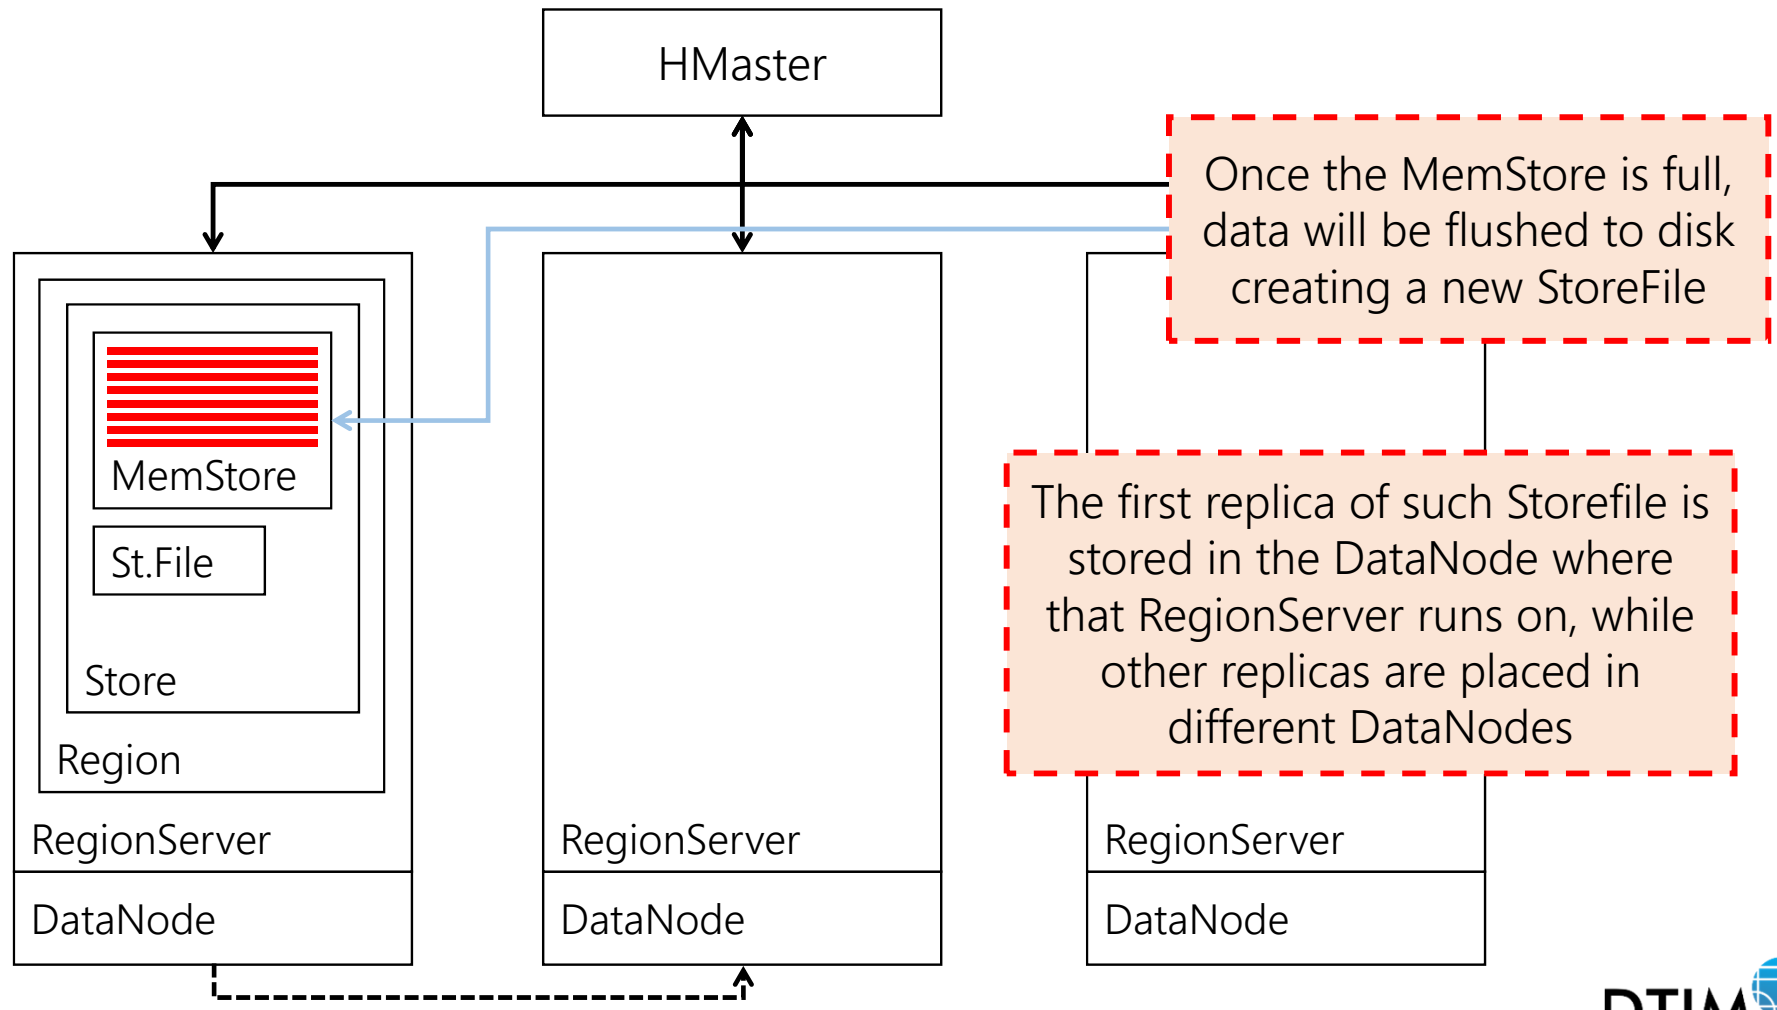
Example of Flush



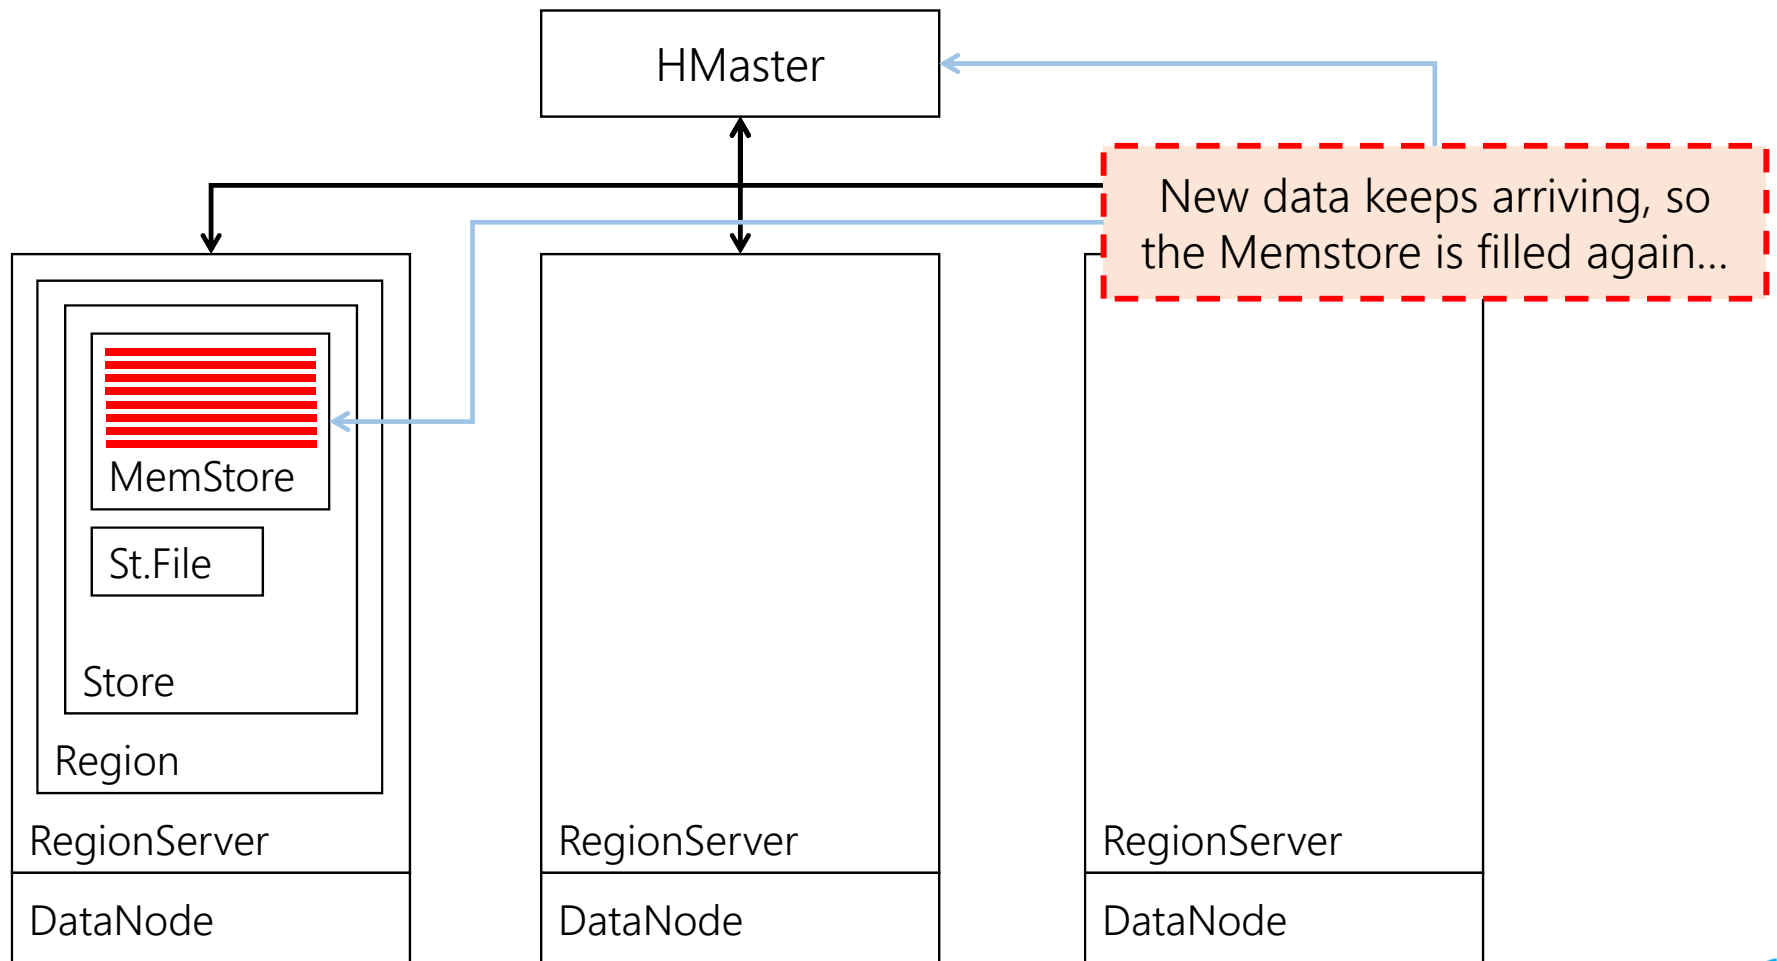
Example of Flush



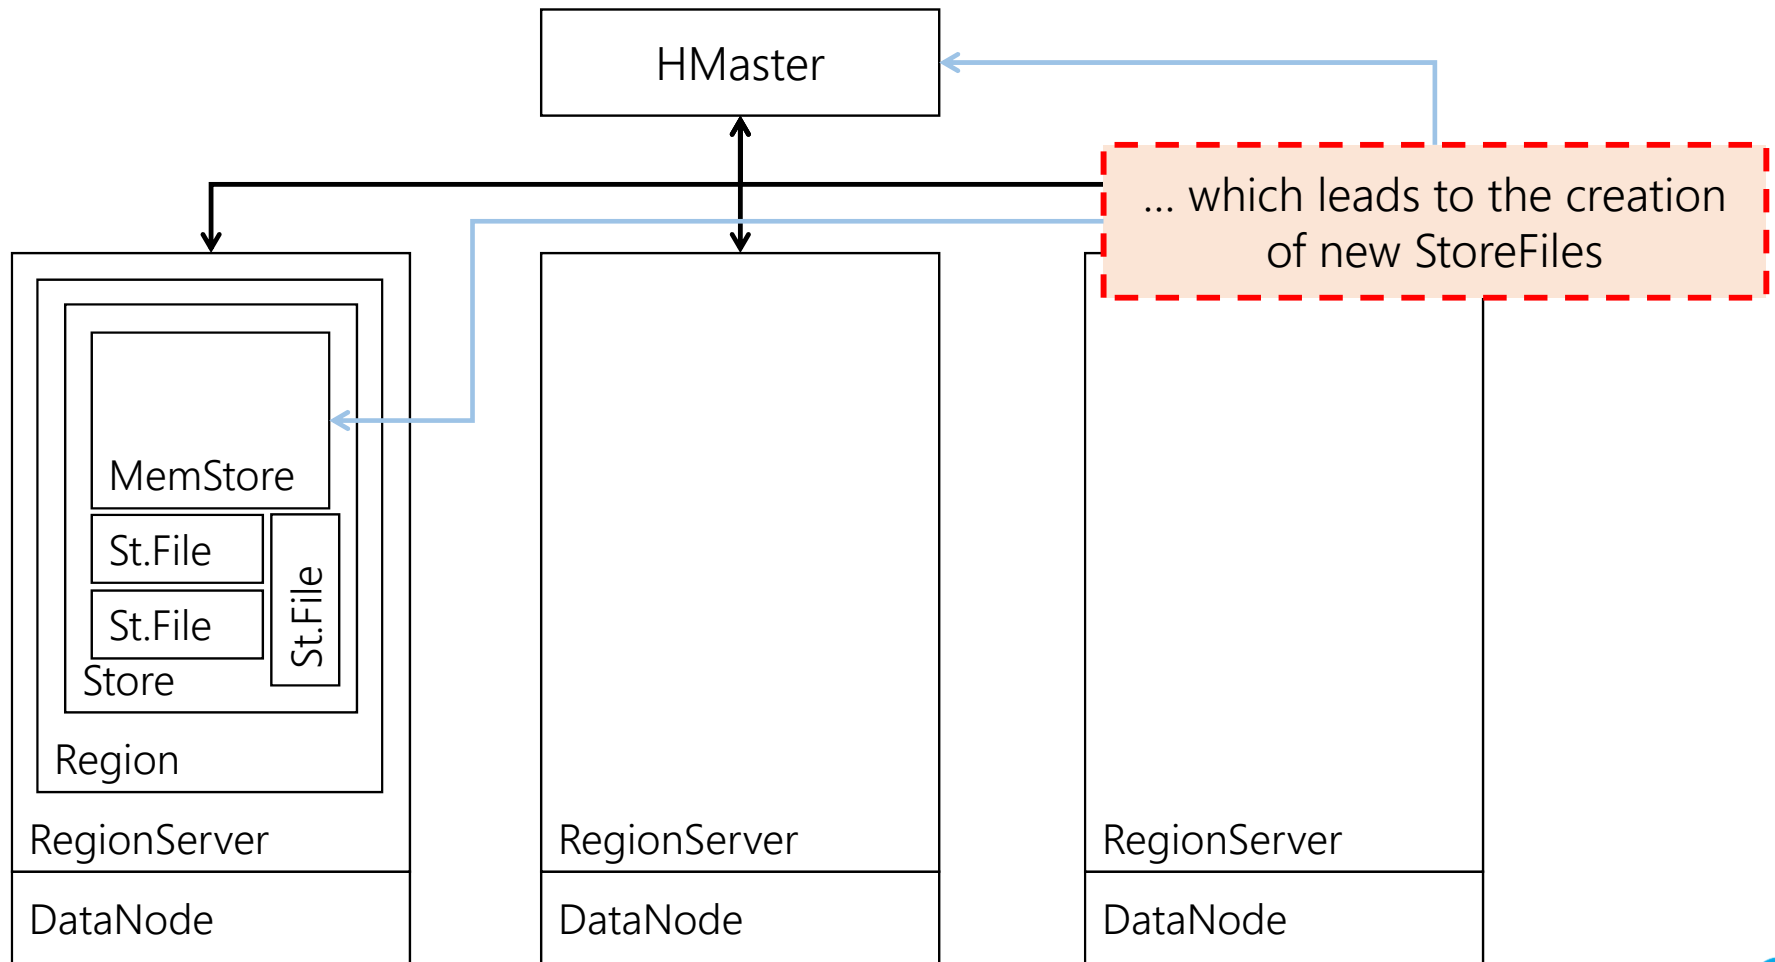
Example of Flush



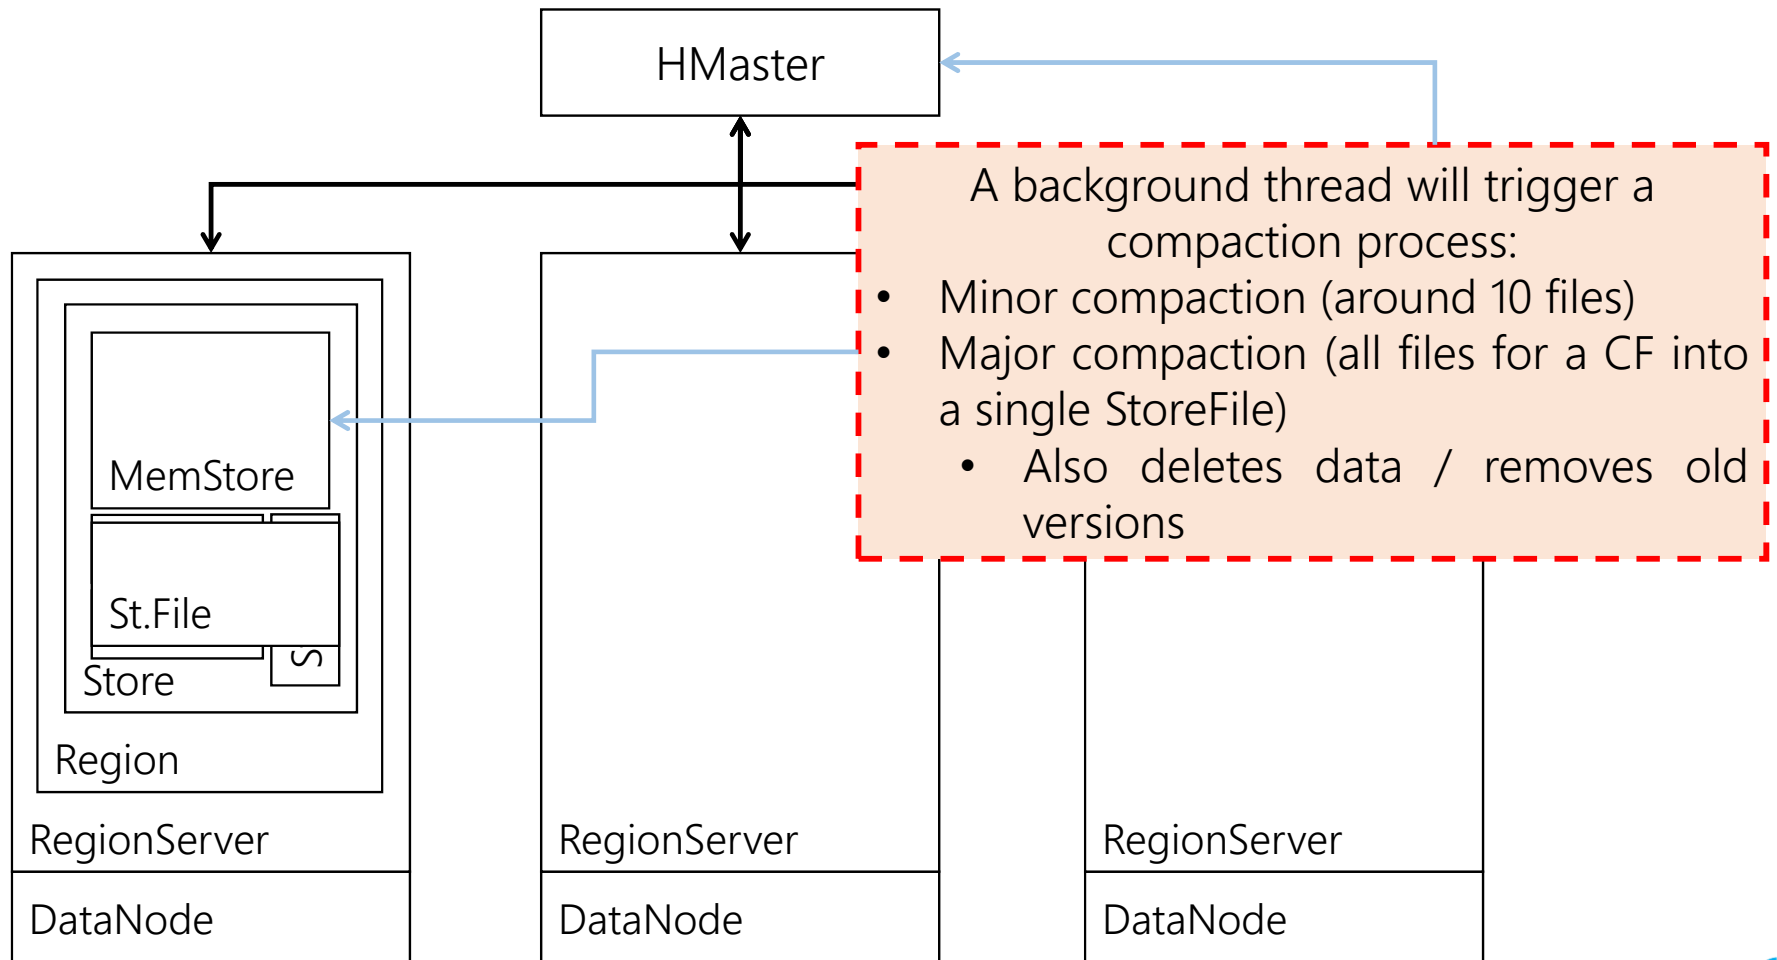
Example of Compactation



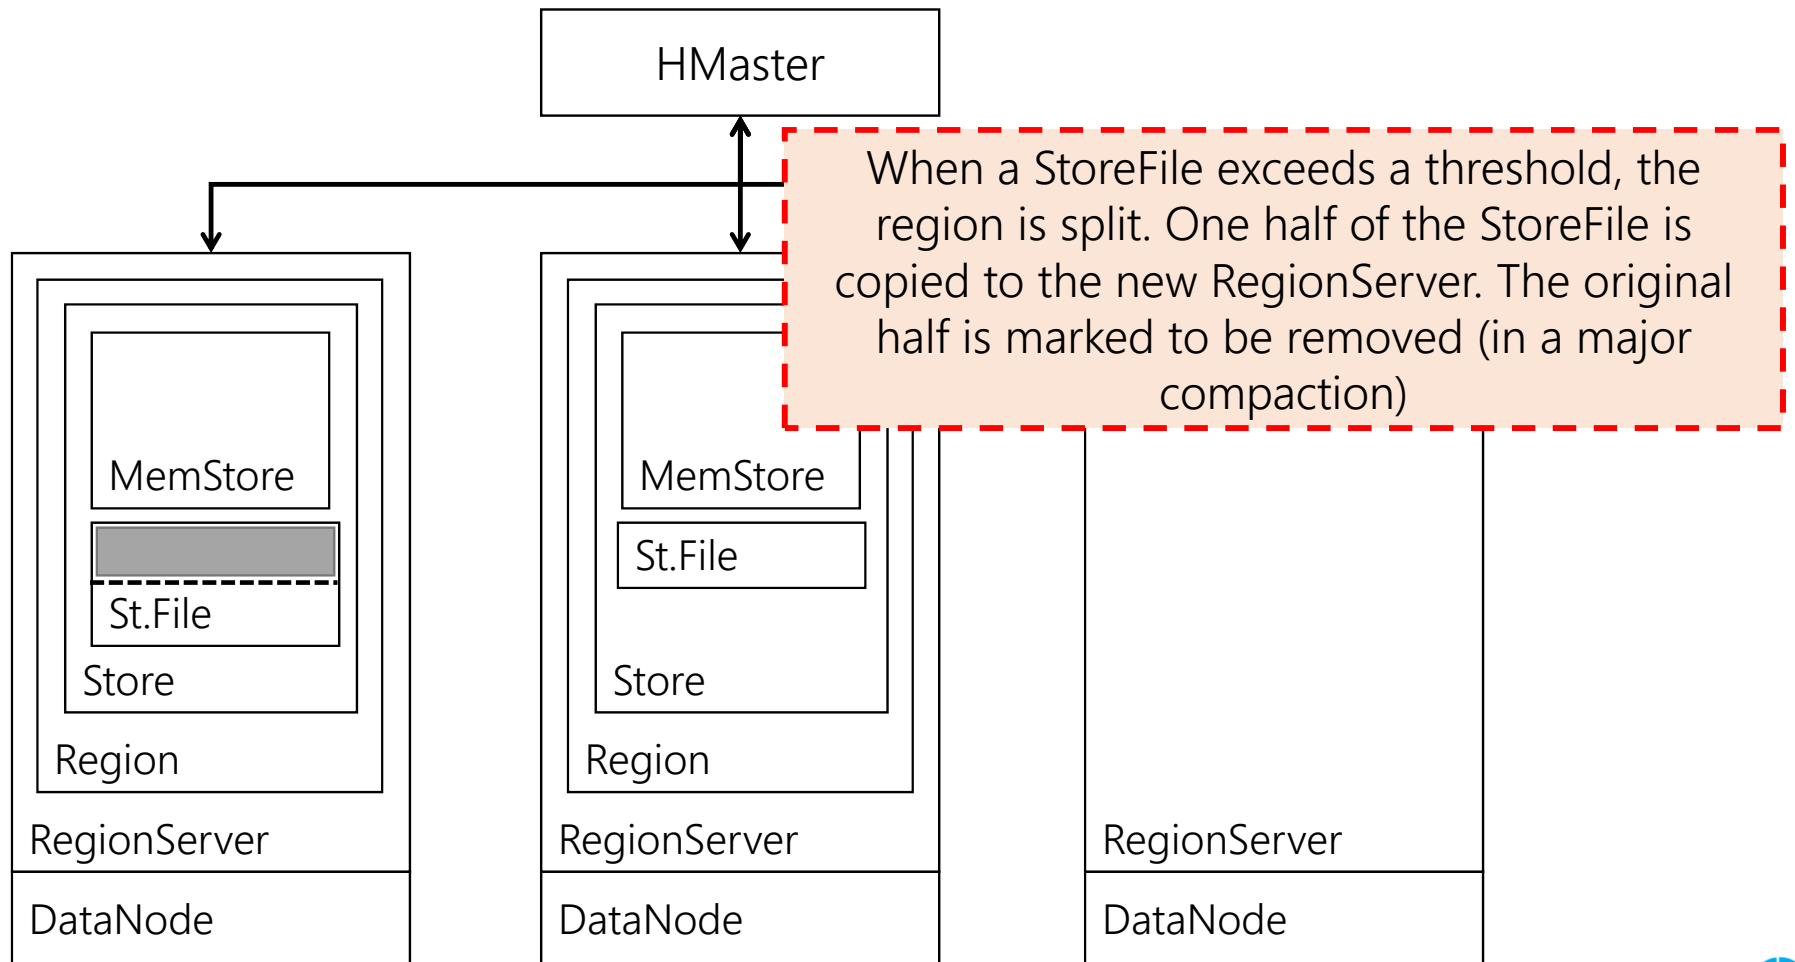
Example of Compactation



Example of Compaction



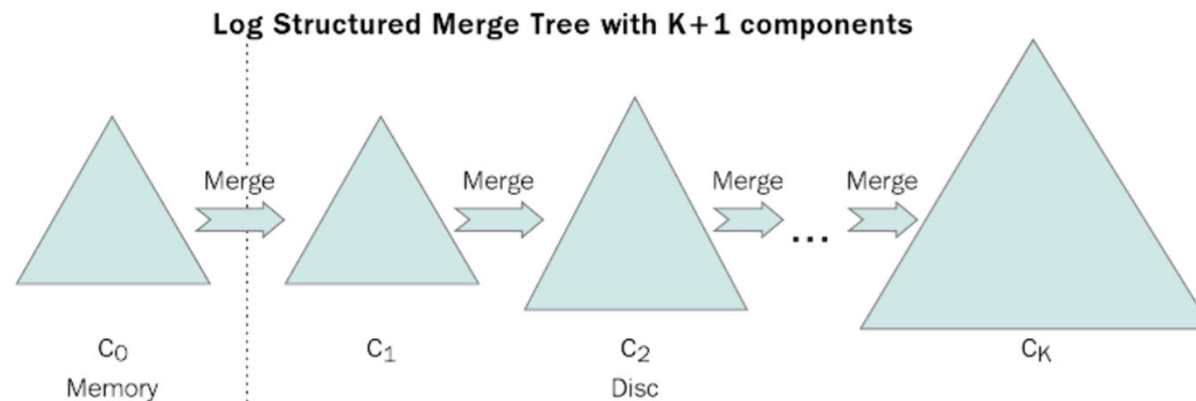
Example of Split



Underlying structures


Log Structured Merge-trees

- Defers and batches index changes
- Consists on two structures
 - In Memory
 - Not necessarily a B-Tree
 - Node different from disk blocks
 - In Disk
 - Multiple components
 - Blocks 100% full
- Regularly merges trees in one level into the next one



HBase LSM-tree implementation


- In Memory (MemStore)
 - Holds the most recent updates for fast lookup
 - Sorted by key
- In Disk (StoreFiles)
 - Immutable Sorted-String Tables (with Bloom Filters and Indexes)
 - May contain different versions of the same row
 - All of them need to be accessed at query time
 - Regularly performs a size-tiered merge process
 - Old SSTables not overwritten (available for recovery)



LSM-tree maintenance algorithms

On memory structure reaching threshold:

1. Take next in memory leafs
2. Flush them to an SSTable




On triggering a compactation

1. Take n SSTables and merge them
2. Put the merge in an in-memory buffer
3. If buffer size exceeds chunk size
 1. Write one chunk to disk
 2. Purge buffer
 3. Keep exceeds in the buffer

Distributed Hashing (alternative to a tree)

- Hash does neither support range queries nor nearest neighbours search
- Distributed hashing challenges
 - Dynamicity:
 - Typical hash function
$$f(x) = g(x) \% \# \text{servers}$$
 - Adding a new server implies modifying hash function
 - Massive data transfer
 - Communicating the new function to all servers
 - Location of the hash directory:
 - Any access must go through the hash directory



Adding a node in static hash

3 Nodes

Bucket/Node 0

Key 0
Key 3
Key 6
Key 9

Bucket/Node 1

Key 1
Key 4
Key 7
Key 10

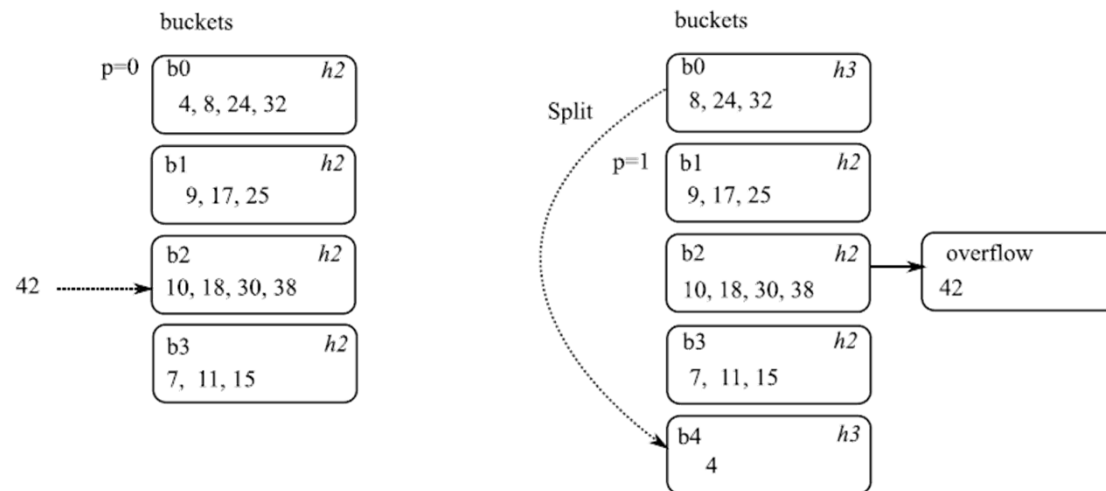
Bucket/Node 2

Key 2
Key 5
Key 8
Key 11

Bucket/Node 3

Distributed Linear Hashing (LH*)

- Maintains an efficient hash in front of dynamicity
 - A split pointer is kept (next bucket to split)
 - A pair of hash functions are considered
 - $\%2^n$ and $\%2^{n+1}$ (being $2^n \leq \text{\#servers} < 2^{n+1}$)
 - Overflow buckets are considered
 - When a bucket overflows the pointed bucket splits (not the overflowed one)

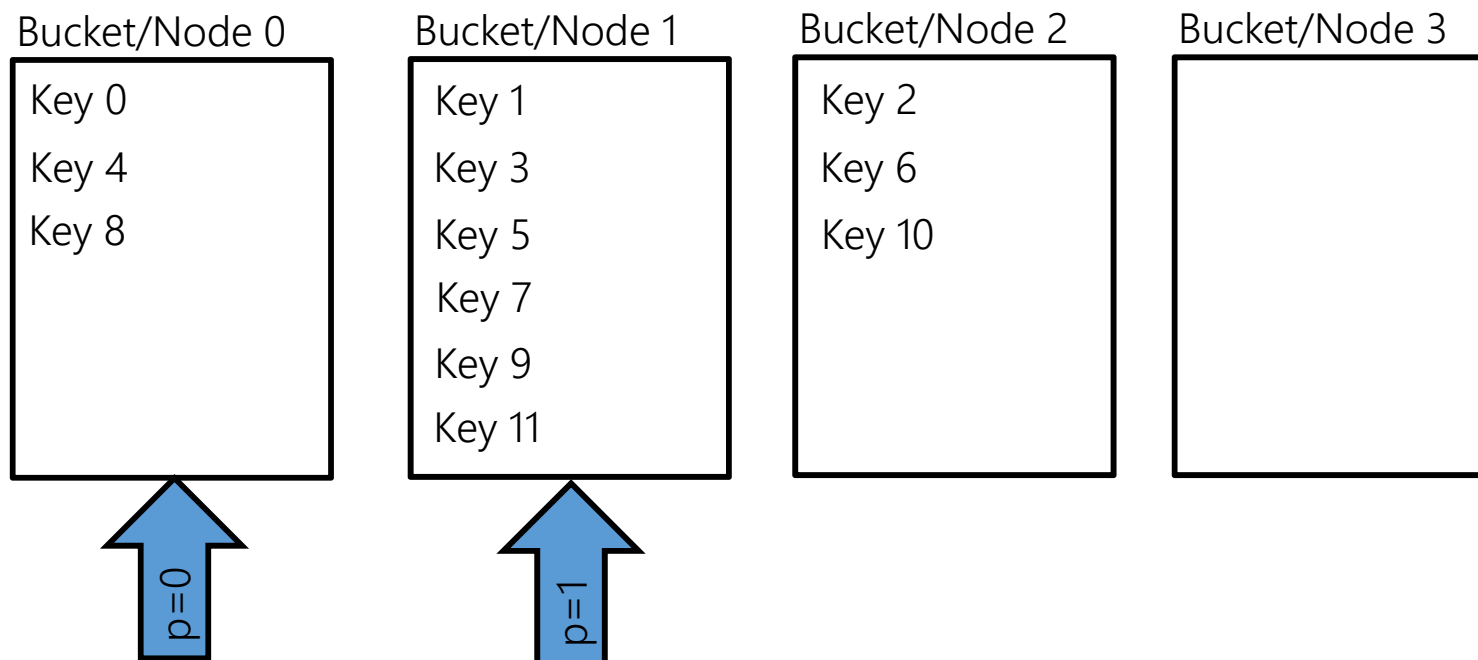


Bucket b2 receives a new object

Bucket b0 splits; bucket b2 is linked to a new one


Adding a node in linear hash

3 Nodes




Consistent Hashing


- The hash function never changes
 - Choose a very large domain D
 - Map server IP addresses and object keys to such domain
 - Organize D as a ring in clockwise order
 - Each node has a successor
 - Objects are assigned as follows:
 - For an object O , $f(O) = D_O$
 - Let $D_{O'}$ and $D_{O''}$ be the two nodes in the ring such that
 - $D_{O'} < D_O \leq D_{O''}$
 - O is assigned to $D_{O''}$
- Further refinements:
 - Assign to the same server several hash values (virtual servers) to balance load
 - Same considerations for the hash directory as for LH*



Adding a new server in Consistent Hashing



Mapping of objects to servers



Server IP3-2 is added, with local re-hashing


S. Abiteboul et al.

- Adding a new server is straightforward
 - It is placed in the ring and part of its successor's objects are transferred to it

Distributed Data Design

Challenge I

Logical structure



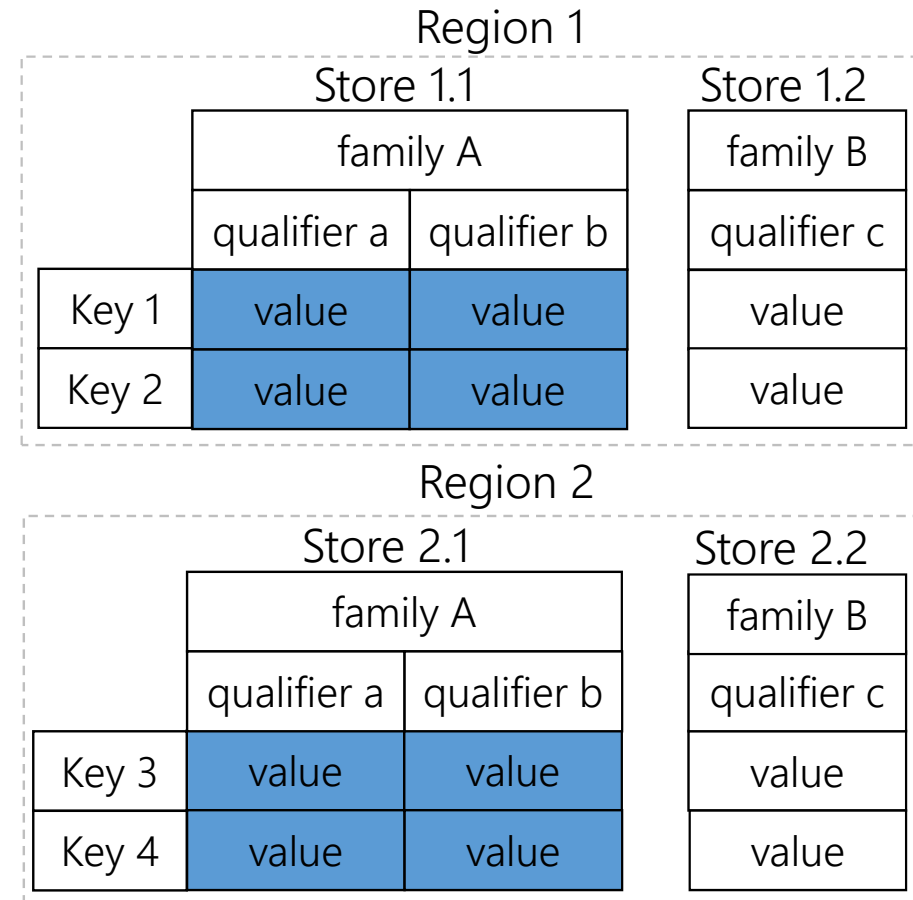

Region 1

	family A		family B
	qualifier a	qualifier b	qualifier c
Key 1	value	value	value
Key 2	value	value	value

Region 2

	family A		family B
	qualifier a	qualifier b	qualifier c
Key 3	value	value	value
Key 4	value	value	value

Physical structure



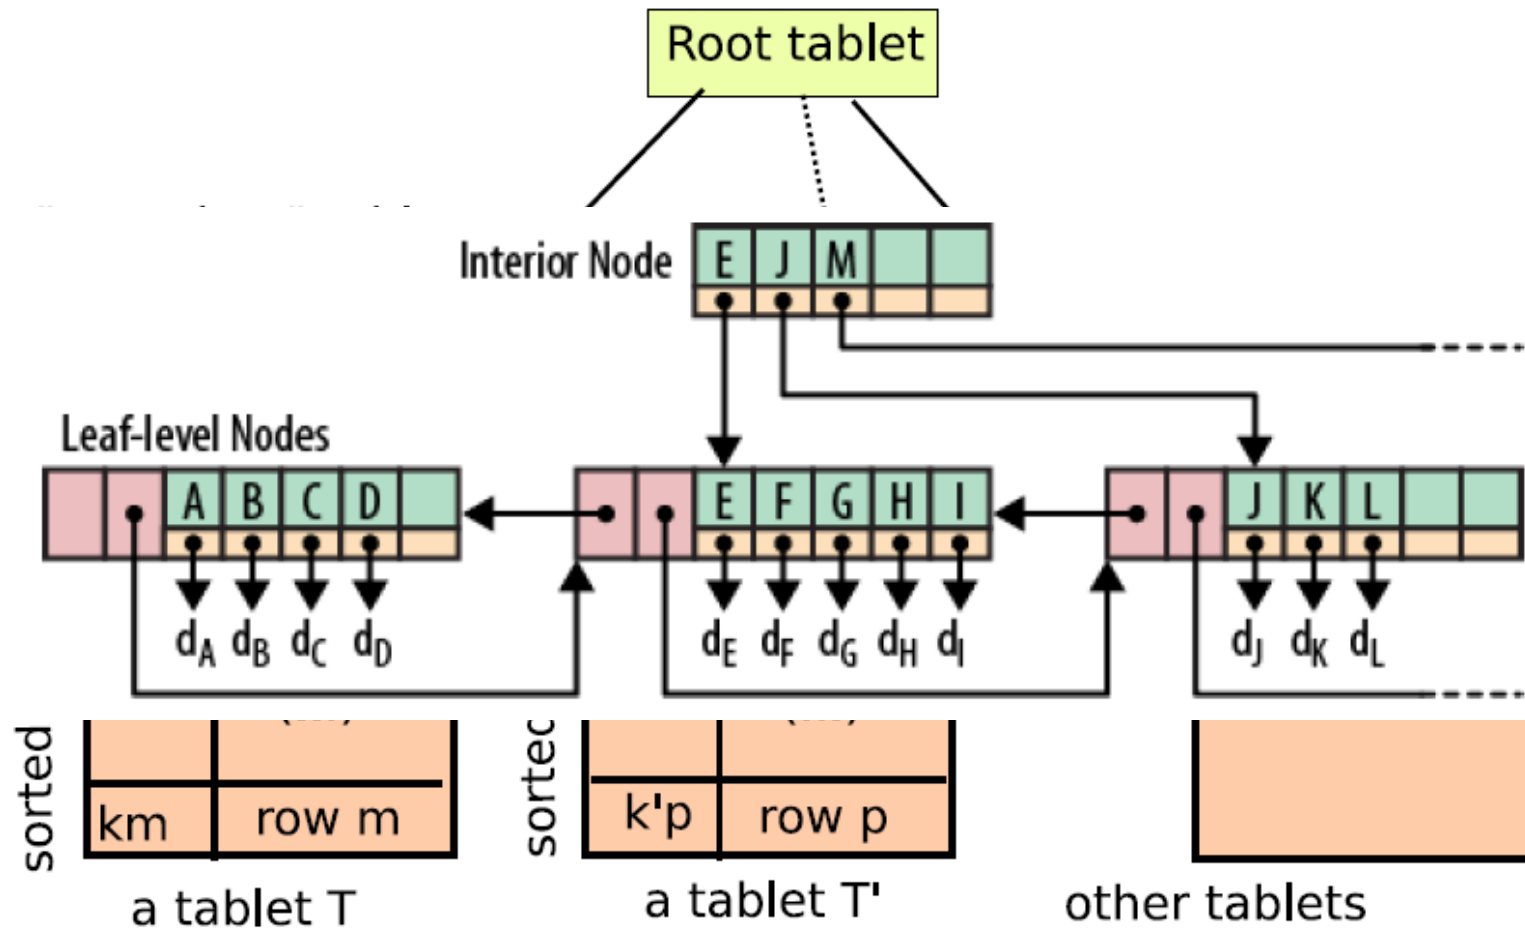
Distributed Catalog Management

Challenge II

HBase Component Roles

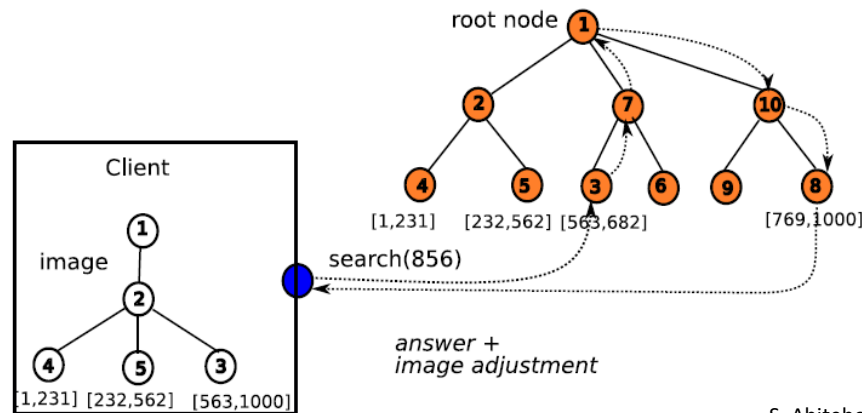
- One master server
 - Maintenance of the table schemas
 - Root region
 - Monitoring of services (heartbeating)
 - Assignment of regions to servers
- Many region servers
 - Each handling around 100-1.000 regions
 - Apply concurrency and recovery techniques
 - Managing split of regions
 - Regions can be sent to another server (load balancer)
 - Managing merge of regions
- Client nodes
 - Cache the metadata sent by the region servers

Metadata hierarchical structure



Metadata synchronization

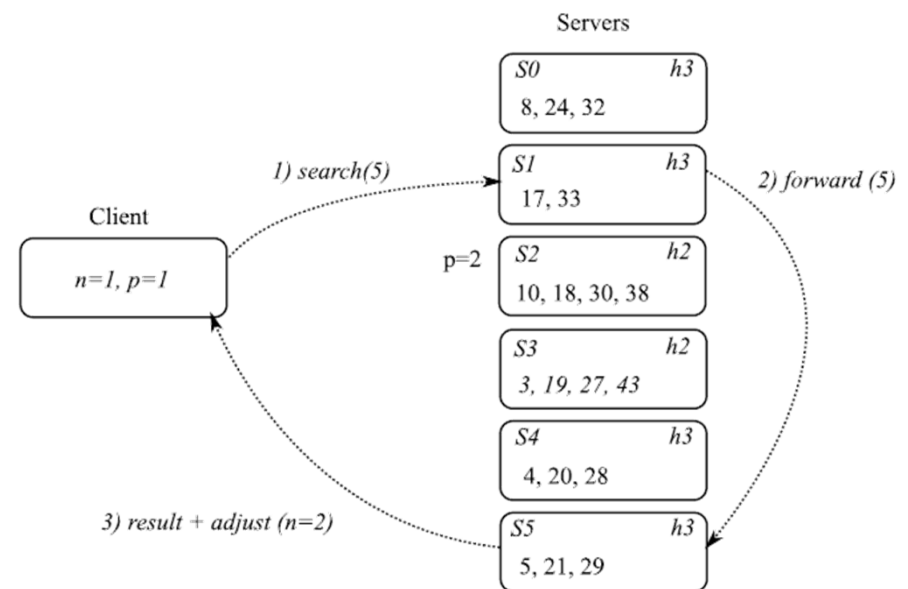
- Split and merge invalidate the cached metadata
 - a) Gossiping
 - b) Lazy updates
 - Discrepancies may cause out-of-range errors, which trigger a stabilization protocol (i.e., **mistake compensation**)
 - Apply forwarding path
 - If an out-of-range error is triggered, it is forwarded upward
 - In the worst case (i.e., reaching the root), 6 network round trips



S. Abiteboul et al.

Updating the Hash Directory in LH*

- Traditionally, each participant has a copy of the hash directory
 - Changes in the hash directory (either hash functions or splits) imply gossiping
 - Including clients nodes
 - It might be acceptable if not too dynamic
- Alternatively, they may contain a partial representation and assume lazy adjustment
 - Apply forwarding path



Distributed Transaction Management

Challenge III

Single row guarantees

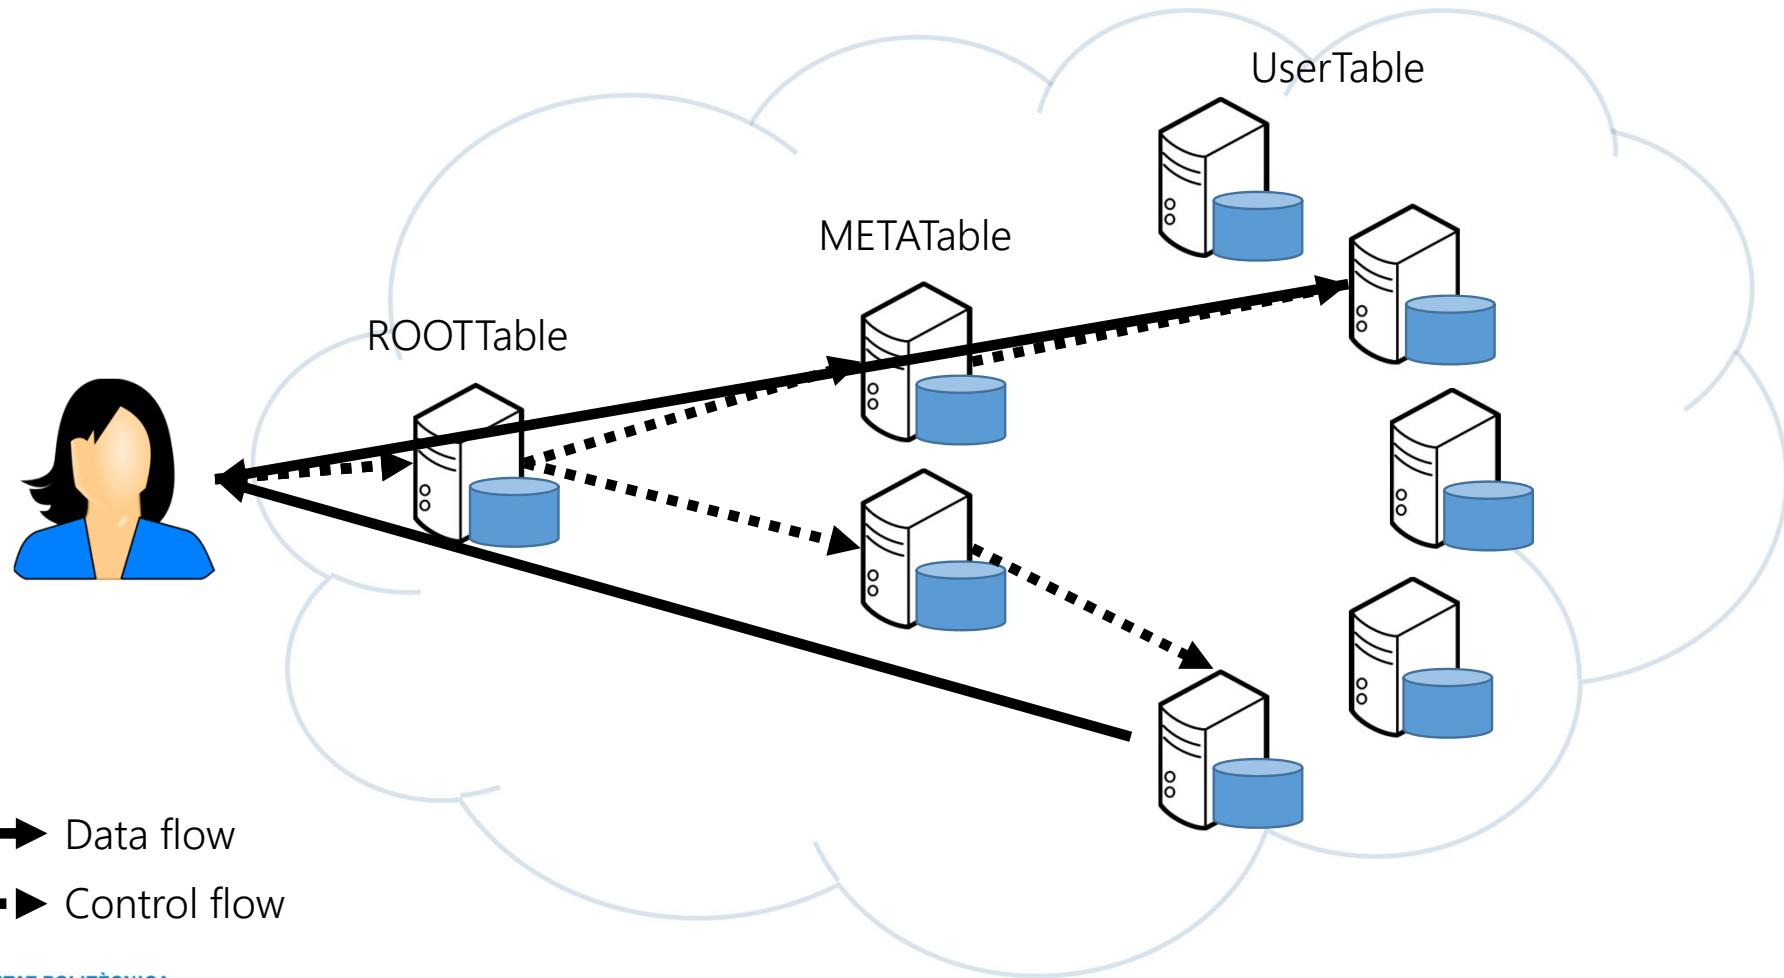
- Atomicity
 - Only single row guarantees (even across families)
 - ... since different families of a row are not distributed
- Consistency
 - Replication relies on HDFS
- Isolation
 - Locking mechanism at row level
 - Does not guarantee snapshot isolation (only read committed)
 - Rows (all families) are separately consistent
 - Sets of rows may not be consistent as a whole
- Durability
 - Write Ahead Log implementation

<https://hbase.apache.org/acid-semantics.html>


Distributed query processing

Challenge IV

Global execution



Local execution



Closing

Summary

- Key-value abstraction
- HBase functional components
- Directory caching
 - Mistake compensation
- Data distribution structures
 - LSM-Tree
 - Linear hash
 - Consistent hash

References

- P. O'Neil et al. *The log-structured merge-tree (LSM-tree)*. Acta Informatica, 33(4). Springer, 1996
- F. Chang et al. *Bigtable: A Distributed Storage System for Structured Data*. OSDI'06
- S. Abiteboul et al. *Web Data Management*. Cambridge University Press, 2011. <http://webdam.inria.fr/Jorge>
- N. Neeraj. *Mastering Apache Cassandra*. Packt, 2015
- O. Romero et al. *Tuning small analytics on Big Data: Data partitioning and secondary indexes in the Hadoop ecosystem*. Information Systems, 54. Elsevier, 2016
- A. Petrov. *Algorithms Behind Modern Storage Systems*. Communications of the ACM 61(8), 2018