

MapReduce I

Big Data Management

Knowledge objectives

1. Enumerate several use cases of MapReduce
2. Explain 6 benefits of using MapReduce
3. Describe what the MapReduce is in the context of a DDBMS
4. Recognize the signature of Map and Reduce functions
5. Justify to which extent MapReduce is generic

Understanding objectives

1. Simulate the execution of a simple MapReduce algorithm from the user (agnostic of implementation details) perspective

Application objectives

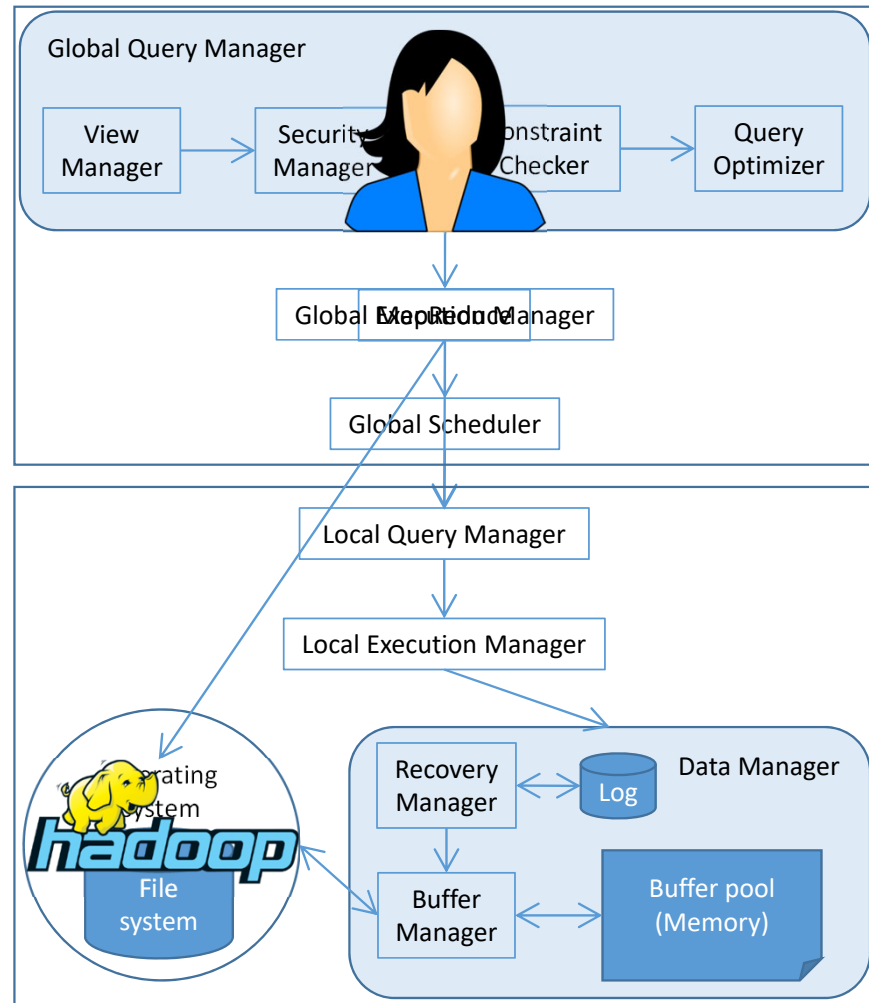
1. Identify the usefulness of MapReduce in a given use case
2. Define the key in the output of the map for a simple problem
3. Provide the pseudo-code of map and reduce functions for a simple problem

Distributed processing framework

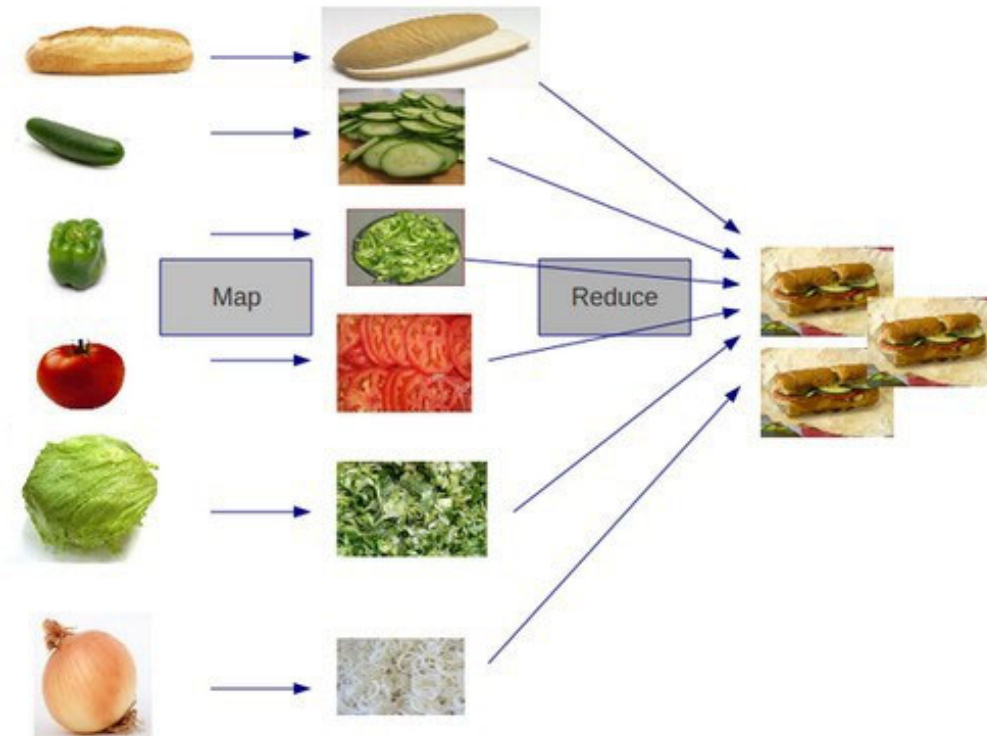
Origins

- Based on Google development
 - Conceived to compute the page Rank
- Data processing framework
 - Facilitate scalability
 - Hidden parallelism
 - Transparent distribution
 - Exploit data locality
 - Balance workload
 - Resilience to failure
 - Fine grained fault tolerance
- Useful in any domain

MapReduce as a DDBMS component

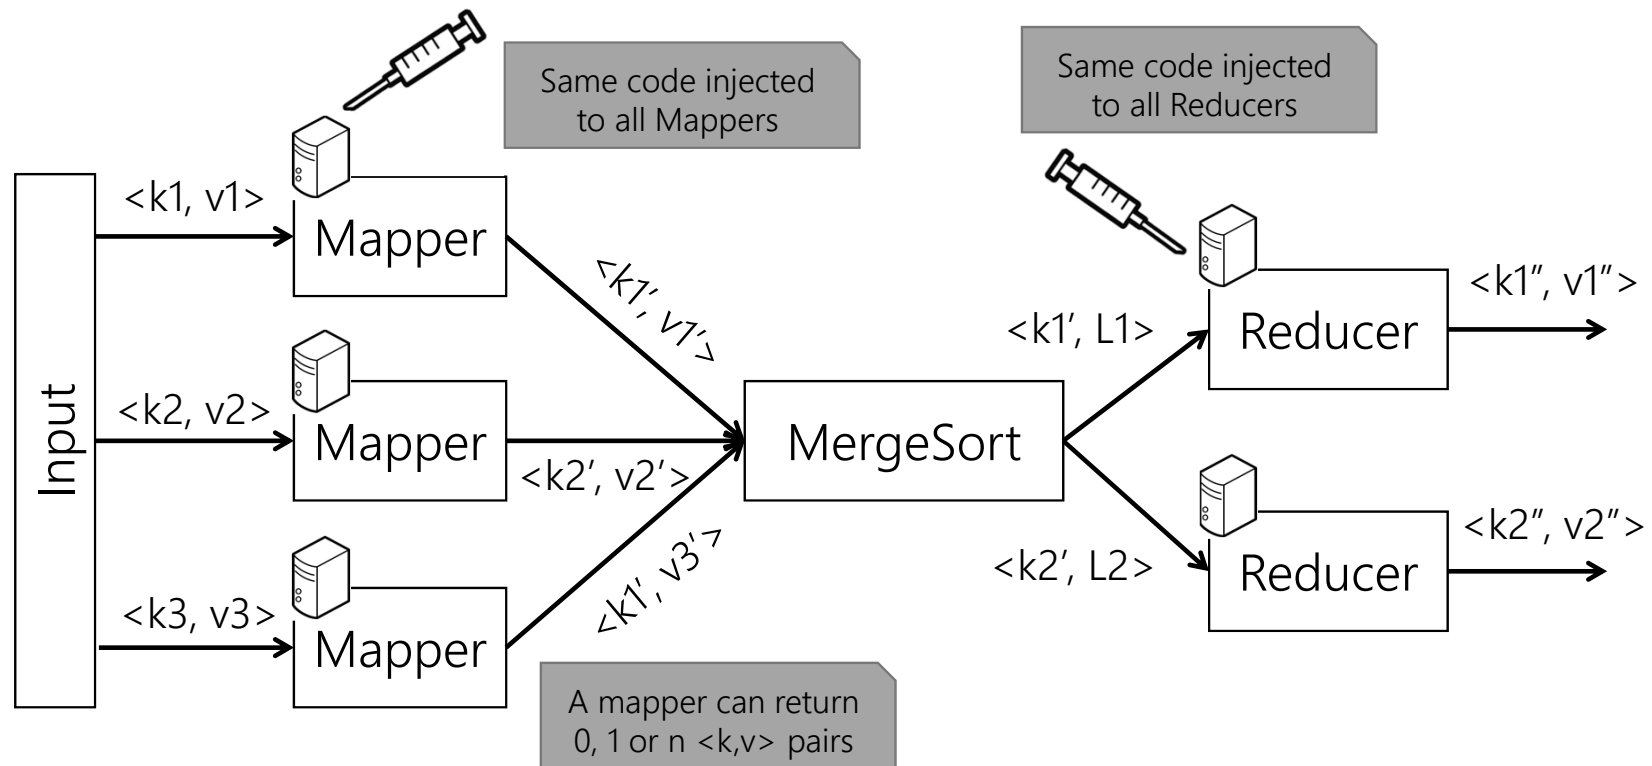


Chain production



Components and use

The MapReduce framework



The MapReduce framework in detail

1. Input: read input from a DFS
2. Map: for each input $\langle \text{key}_{\text{in}}, \text{value}_{\text{in}} \rangle$
 - generate zero-to-many $\langle \text{key}_{\text{map}}, \text{value}_{\text{map}} \rangle$
3. Partition: assign sets of $\langle \text{key}_{\text{map}}, \text{value}_{\text{map}} \rangle$ to reducer machines
4. Shuffle: data are shipped to reducer machines using a DFS
5. Sort&Merge: reducers sort their input data by key
6. Reduce: for each key_{map}
 - the set $\text{value}_{\text{map}}$ is processed to produce zero-to-many $\langle \text{key}_{\text{red}}, \text{value}_{\text{red}} \rangle$
7. Output: writes the result of reducers to the DFS

Formal definition

- Single input
 - Data are represented as $\langle \text{key}, \text{value} \rangle$ pairs
 - Value can be anything (structured or not)
- Functional programming
 - Map phase, for each input $\langle \text{key}, \text{value} \rangle$ a function f is applied that returns a multiset of new $\langle \text{key}, \text{value} \rangle$ pairs:

$$f(\langle k, v \rangle) \mapsto \{ \langle k_1, v_1 \rangle, \dots, \langle k_n, v_n \rangle \}$$

- Reduce phase, all pairs with the same key are grouped and a function g is applied, which returns also a multiset of new $\langle \text{key}, \text{value} \rangle$ pairs:

$$g(\langle k, \{v_1, \dots, v_n\} \rangle) \mapsto \{ \langle k_1, v_1 \rangle, \dots, \langle k_m, v_m \rangle \}$$

MapReduce examples

Word count

Word count example

```
The Project Gutenberg Ebook of The Outline of Science, Vol. 1 (of 4), by
J. Arthur Thomson

This eBook is for the use of anyone anywhere at no cost and with
almost no restrictions whatsoever. You may copy it, give it away or
re-use it under the terms of the Project Gutenberg License included
with this eBook or online at www.gutenberg.org

Title: The Outline of Science, Vol. 1 (of 4)
      A Plain Story Simply Told

Author: J. Arthur Thomson

Release Date: January 22, 2007 [EBook #20417]

Language: English

Character set encoding: ASCII

*** START OF THIS PROJECT GUTENBERG EBOOK OUTLINE OF SCIENCE ***

Produced by Brian Janes, Leonard Johnson and the Online
Distributed Proofreading Team at http://www.pgdp.net

[Illustration: THE GREAT SCARLET SOLAR PROMINENCES, WHICH ARE SUCH A
NOTABLE FEATURE OF THE SOLAR PHENOMENA, ARE IMMENSE OUTBURSTS OF FLAMING
HYDROGEN RISING SOMETIMES TO A HEIGHT OF 500,000 MILES]

THE
OUTLINE OF SCIENCE

A PLAIN STORY SIMPLY TOLD

EDITED BY
J. ARTHUR THOMSON
REGIUS PROFESSOR OF NATURAL HISTORY IN THE
UNIVERSITY OF ABERDEEN

WITH OVER 800 ILLUSTRATIONS
OF WHICH ABOUT 40 ARE IN COLOUR
```

<#line, text>

Map

Merge-Short

Reduce

The
Project
Gutenberg
Ebook
of
The
Outline
Of
Science,
Vol.
1
(of
4),
by

1
1
1
1
1
1
1
1
1
1
1
1
1
1
1

The

57631

["The", [1,1,1,...]]

WordCount Code Example

```
public void map(LongWritable key, Text value) {  
    String line = value.toString();  
    StringTokenizer tokenizer = new StringTokenizer(line);  
    while (tokenizer.hasMoreTokens()) {  
        write(new Text(tokenizer.nextToken()), new IntWritable(1));  
    }  
}
```

```
public void reduce(Text key, Iterable<IntWritable> values) {  
    int sum = 0;  
    for (IntWritable val : values) {  
        sum += val.get();  
    }  
    write(key, new IntWritable(sum));  
}
```

WordCount Code Example

```
public void map(Key Value) {
```

Blackbox

```
    write(Key Value);  
}
```

```
public void reduce(Key Value) {
```

Blackbox

```
    write(Key Value);  
}
```


MapReduce examples

Common friends

Friends in common example

- In a social network (e.g., Facebook), we aim to compute the friends in common for every pair of users
 - This is a value that does not frequently change, so it can be precomputed

- Friends are stored as

Person -> [List of friends]

- $A \rightarrow B\ C\ D$
- $B \rightarrow A\ C\ D\ E$
- $C \rightarrow A\ B\ D\ E$
- $D \rightarrow A\ B\ C\ E$
- $E \rightarrow B\ C\ D$

Friends in common – Map task

- For every friend in the list, the mapper will generate a $\langle k, v \rangle$
 - Key: the input key concatenated with one friend in alphabetical order
 - Value: the whole list of friends
- Keys will be sorted, a pair of friends go to the same reducer

A \rightarrow B C D

(A B) \rightarrow B C D

(A C) \rightarrow B C D

(A D) \rightarrow B C D

B \rightarrow A C D E

(A B) \rightarrow A C D E

(B C) \rightarrow A C D E

(B D) \rightarrow A C D E

(B E) \rightarrow A C D E

C \rightarrow A B D E

(A C) \rightarrow A B D E

(B C) \rightarrow A B D E

(C D) \rightarrow A B D E

(C E) \rightarrow A B D E

...

Friends in common – Reduce task

- Reducers receive two lists of friends per pair of people

(A B) → (B C D) (A C D E)

(A C) → (B C D) (A B D E)

(A D) → (B C D) (A B C E)

- The reduce function intersects the lists of values and generates the same key

(A B) → (C D)

(A C) → (B D)

(A D) → (B C)

...

- Now, when D visits A's profile we can lookup (A D) to see their common friends

Relational algebra in MapReduce

MapReduce Genericity

- Supported in many store systems
 - HBase, MongoDB, CouchDB, etc.
- Programming paradigm is computationally complete
 - Any data process can be adapted to it
 - Some tasks better adapt to it than others
 - Not necessarily efficient
 - Optimization is very limited because of lack of expressivity
- Signature is closed
 - Iterations can be chained
 - Fault tolerance is not guaranteed in between
 - Resources are released to be just requested again
- Criticized for being too low-level
 - APIs for Ruby, Python, Java, C++, etc.
 - Attempts to build declarative languages on top
 - SQL-like
 - HiveQL
 - Cassandra Query Language (CQL)

Relational operations: Projection

$$\pi_{a_{i_1}, \dots, a_{i_n}}(T) \Rightarrow \begin{cases} \text{map}(\text{key } k, \text{value } v) \mapsto [(\text{prj}_{a_{i_1}, \dots, a_{i_n}}(k \oplus v), 1)] \\ \text{reduce}(\text{key } ik, \text{vset } ivs) \mapsto [(ik)] \end{cases}$$

Relational operations: Cross Product

$$T \times S \Rightarrow \begin{cases} \text{map}(\text{key } k, \text{value } v) \mapsto \\ \begin{cases} [(h_T(k) \bmod D, k \oplus v)] & \text{if } \text{input}(k \oplus v) = T, \\ [(0, k \oplus v), \dots, (D-1, k \oplus v)] & \text{if } \text{input}(k \oplus v) = S. \end{cases} \\ \text{reduce}(\text{key } ik, \text{vset } ivs) \mapsto \\ \left[\text{crossproduct}(T_{ik}, S) \mid \right. \\ \quad T_{ik} = \{iv \mid iv \in ivs \wedge \text{input}(iv) = T\}, \\ \quad \left. S = \{iv \mid iv \in ivs \wedge \text{input}(iv) = S\} \right] \end{cases}$$

Closing

Summary

- MapReduce usefulness and benefits
- MapReduce programming model
 - Expressivity
- Relational algebra in MapReduce

References

- J. Dean et al. *MapReduce: Simplified Data Processing on Large Clusters*. OSDI'04
- A. Pavlo et al. *A Comparison of Approaches to Large-Scale Data Analysis*. SIGMOD, 2009
- J. Dittrich et al. *Hadoop++: Making a Yellow Elephant Run Like a Cheetah (Without It Even Noticing)*. Proc. VLDB Endow. 3(1-2), 2010
- M. Stonebraker et al. *MapReduce and parallel DBMSs: friends or foes?* Communication of ACM 53(1), 2010
- S. Abiteboul et al. *Web data management*. Cambridge University Press, 2011
- A. Rajaraman et al. *Mining massive data sets*. Cambridge University Press, 2012
- P. Sadagale and M. Fowler. *NoSQL distilled*. Addison-Wesley, 2013