

# MapReduce I

Big Data Management

# Knowledge objectives

1. Enumerate several use cases of MapReduce
2. Explain 6 benefits of using MapReduce
3. Describe what the MapReduce is in the context of a DDBMS
4. Recognize the signature of Map and Reduce functions
5. Justify to which extent MapReduce is generic

# Understanding objectives

1. Simulate the execution of a simple MapReduce algorithm from the user (agnostic of implementation details) perspective

# Application objectives

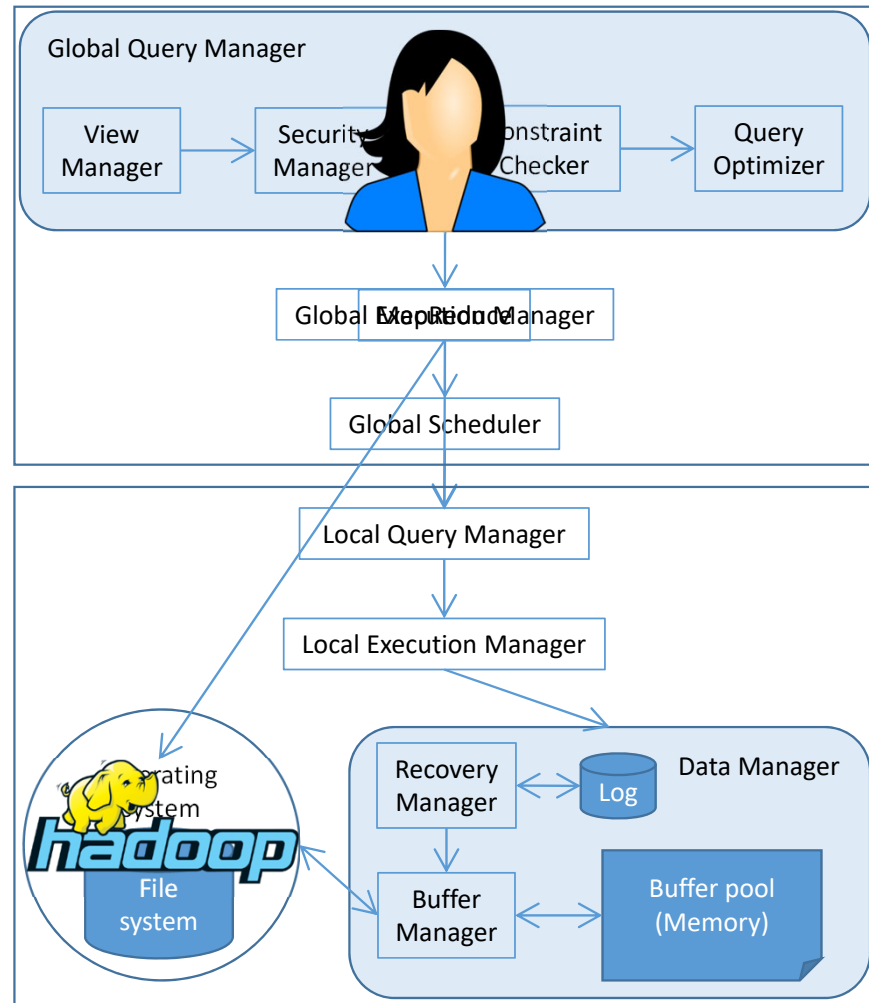
1. Identify the usefulness of MapReduce in a given use case
2. Define the key in the output of the map for a simple problem
3. Provide the pseudo-code of map and reduce functions for a simple problem

# Distributed processing framework

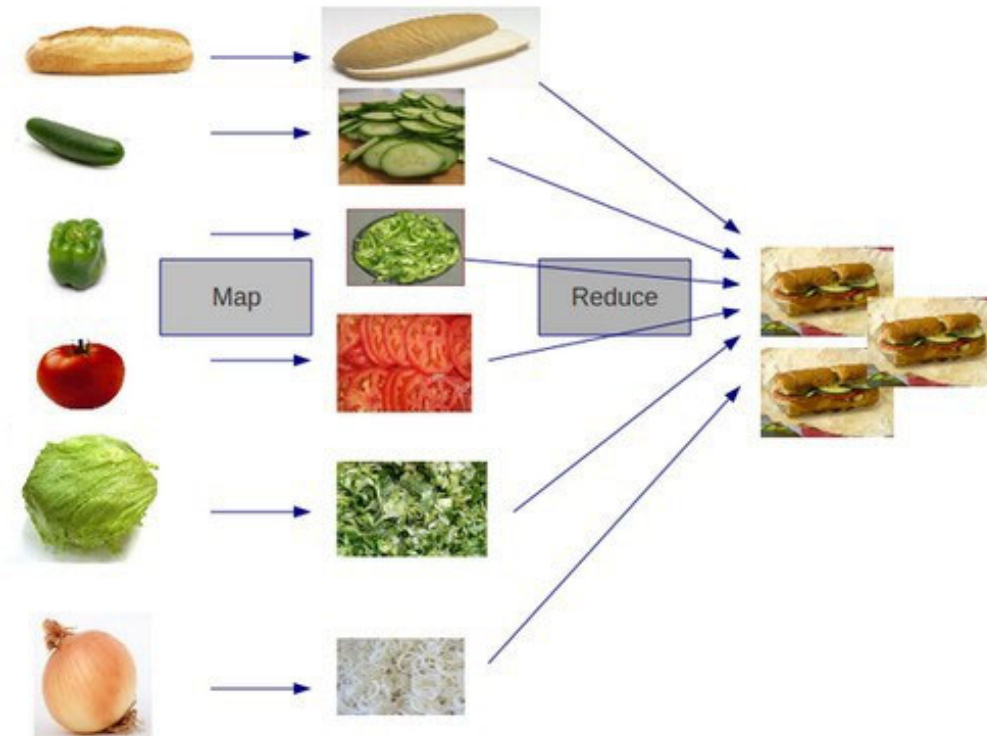
# Origins

- Based on Google development
  - Conceived to compute the page Rank
- Data processing framework
  - Facilitate scalability
  - Hidden parallelism
  - Transparent distribution
    - Exploit data locality
    - Balance workload
  - Resilience to failure
    - Fine grained fault tolerance
- Useful in any domain

# MapReduce as a DDBMS component



# Chain production





# Components and use

# Formal definition

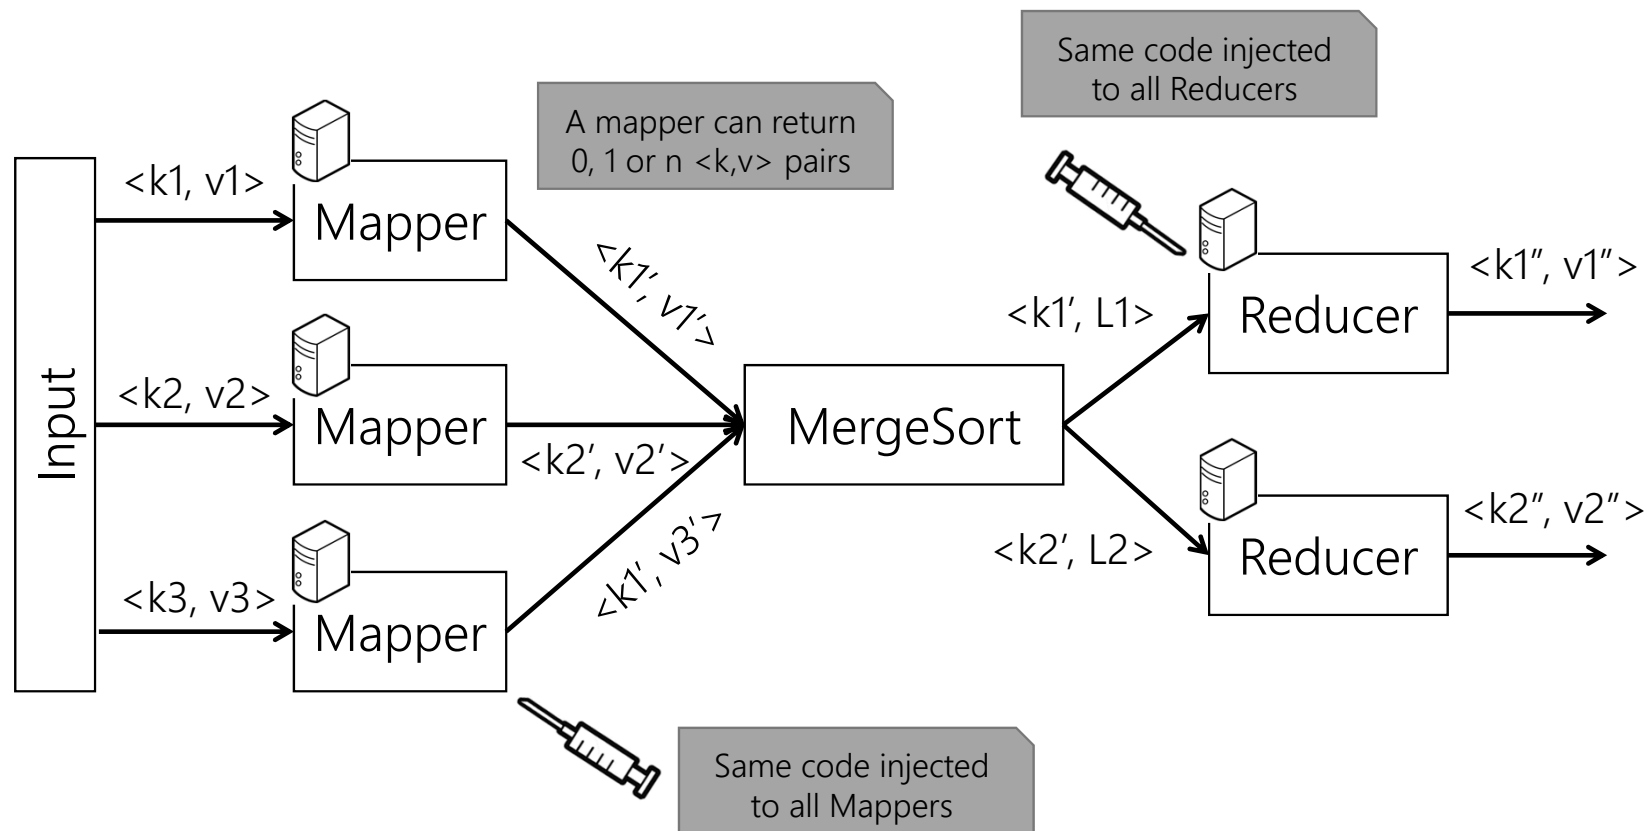
- Single input
  - Data are represented as  $\langle \text{key}, \text{value} \rangle$  pairs
    - Value can be anything (structured or not)
- Functional programming
  - Map phase, for each input  $\langle \text{key}, \text{value} \rangle$  a function  $f$  is applied that returns a multiset of new  $\langle \text{key}, \text{value} \rangle$  pairs:

$$f(\langle k, v \rangle) \mapsto \{ \langle k_1, v_1 \rangle, \dots, \langle k_n, v_n \rangle \}$$

- Reduce phase, all pairs with the same key are grouped and a function  $g$  is applied, which returns also a multiset of new  $\langle \text{key}, \text{value} \rangle$  pairs:

$$g(\langle k, \{v_1, \dots, v_n\} \rangle) \mapsto \{ \langle k_1, v_1 \rangle, \dots, \langle k_m, v_m \rangle \}$$

# The MapReduce framework



# The MapReduce framework in detail

1. Input: read input from a DFS
2. Map: for each input  $\langle \text{key}_{\text{in}}, \text{value}_{\text{in}} \rangle$ 
  - generate zero-to-many  $\langle \text{key}_{\text{map}}, \text{value}_{\text{map}} \rangle$
3. Partition: assign sets of  $\langle \text{key}_{\text{map}}, \text{value}_{\text{map}} \rangle$  to reducer machines
4. Shuffle: data are shipped to reducer machines using a DFS
5. Sort&Merge: reducers sort their input data by key
6. Reduce: for each  $\text{key}_{\text{map}}$ 
  - the set  $\text{value}_{\text{map}}$  is processed to produce zero-to-many  $\langle \text{key}_{\text{red}}, \text{value}_{\text{red}} \rangle$
7. Output: writes the result of reducers to the DFS

# MapReduce examples

# Word count example

```
The Project Gutenberg Ebook of The Outline of Science, Vol. 1 (of 4), by
J. Arthur Thomson

This eBook is for the use of anyone anywhere at no cost and with
almost no restrictions whatsoever. You may copy it, give it away or
re-use it under the terms of the Project Gutenberg License included
with this eBook or online at www.gutenberg.org

Title: The Outline of Science, Vol. 1 (of 4)
      A Plain Story Simply Told

Author: J. Arthur Thomson

Release Date: January 22, 2007 [EBook #20417]

Language: English

Character set encoding: ASCII

*** START OF THIS PROJECT GUTENBERG EBOOK OUTLINE OF SCIENCE ***

Produced by Brian Janes, Leonard Johnson and the Online
Distributed Proofreading Team at http://www.pgdp.net

[Illustration: THE GREAT SCARLET SOLAR PROMINENCES, WHICH ARE SUCH A
NOTABLE FEATURE OF THE SOLAR PHENOMENA, ARE IMMENSE OUTBURSTS OF FLAMING
HYDROGEN RISING SOMETIMES TO A HEIGHT OF 500,000 MILES]

THE
OUTLINE OF SCIENCE

A PLAIN STORY SIMPLY TOLD

EDITED BY
J. ARTHUR THOMSON
REGIUS PROFESSOR OF NATURAL HISTORY IN THE
UNIVERSITY OF ABERDEEN

WITH OVER 800 ILLUSTRATIONS
OF WHICH ABOUT 40 ARE IN COLOUR
```

<#line, text>

Map

Merge-Short

Reduce

The  
Project  
Gutenberg  
Ebook  
of  
The  
Outline  
Of  
Science,  
Vol.  
1  
(of  
4),  
by

1  
1  
1  
1  
1  
1  
1  
1  
1  
1  
1  
1  
1  
1  
1

The

57631

["The", [1,1,1,...]]

# WordCount Code Example

```
public void map(Key Value) {
```

Blackbox

```
    write(Key Value);  
}
```

```
public void reduce(Key Values) {
```

Blackbox

```
    write(Key Value);  
}
```

# Friends in common example

- In a social network (e.g., Facebook) we aim to compute the friends in common
  - This is a value that does not frequently change, so it can be precomputed

- Friends are stored as

Person -> [List of friends]

- $A \rightarrow B\ C\ D$
- $B \rightarrow A\ C\ D\ E$
- $C \rightarrow A\ B\ D\ E$
- $D \rightarrow A\ B\ C\ E$
- $E \rightarrow B\ C\ D$



# Friends in common – Map task

- For every friend in the list, the mapper will generate a  $\langle k, v \rangle$ 
  - Key: the input key and one friend in alphabetical order
  - Value: the list of friends
- Keys will be sorted, a pair of friends go to the same reducer

A  $\rightarrow$  B C D

(A B)  $\rightarrow$  B C D

(A C)  $\rightarrow$  B C D

(A D)  $\rightarrow$  B C D

B  $\rightarrow$  A C D E

(A B)  $\rightarrow$  A C D E

(B C)  $\rightarrow$  A C D E

(B D)  $\rightarrow$  A C D E

(B E)  $\rightarrow$  A C D E

C  $\rightarrow$  A B D E

(A C)  $\rightarrow$  A B D E

(B C)  $\rightarrow$  A B D E

(C D)  $\rightarrow$  A B D E

(C E)  $\rightarrow$  A B D E

...

# Friends in common – Reduce task

- Reducers receive two lists of friends per pair of people

(A B) → (B C D) (A C D E)

(A C) → (B C D) (A B D E)

(A D) → (B C D) (A B C E)

- The reduce function intersects the lists of values and generates the same key

(A B) → (C D)

(A C) → (B D)

(A D) → (B C)

...

- Now, when D visits A's profile we can lookup (A D) to see their common friends

# Relational algebra in MapReduce

# Relational operations: Projection

$$\pi_{a_{i_1}, \dots, a_{i_n}}(T) \Rightarrow \begin{cases} \text{map}(\text{key } k, \text{value } v) \mapsto [(\text{prj}_{a_{i_1}, \dots, a_{i_n}}(k \oplus v), 1)] \\ \text{reduce}(\text{key } ik, \text{vset } ivs) \mapsto [(ik)] \end{cases}$$

# Relational operations: Cross Product

$$T \times S \Rightarrow \begin{cases} \text{map}(\text{key } k, \text{value } v) \mapsto \\ \begin{cases} [(h_T(k) \bmod D, k \oplus v)] & \text{if } \text{input}(k \oplus v) = T, \\ [(0, k \oplus v), \dots, (D-1, k \oplus v)] & \text{if } \text{input}(k \oplus v) = S. \end{cases} \\ \text{reduce}(\text{key } ik, \text{vset } ivs) \mapsto \\ \left[ \text{crossproduct}(T_{ik}, S) \mid \right. \\ \quad T_{ik} = \{iv \mid iv \in ivs \wedge \text{input}(iv) = T\}, \\ \quad \left. S = \{iv \mid iv \in ivs \wedge \text{input}(iv) = S\} \right] \end{cases}$$

# Genericity

- The MapReduce programming paradigm is computationally complete
  - Any data process can be adapted to it
    - Some tasks better adapt to it than others
    - Not necessarily efficient
      - Optimization is very limited because of lack of expressivity
- MapReduce's signature is closed
  - MapReduce iterations can be chained
    - Fault tolerance is not guaranteed in between
    - Resources are released to be just requested again
- Supported in many store systems
  - HBase, MongoDB, CouchDB, etc.
- Criticized for being too low-level
  - APIs for Ruby, Python, Java, C++, etc.
  - Attempts to build declarative languages on top
    - SQL-like
      - HiveQL
      - Cassandra Query Language (CQL)

# Closing

# Summary

- MapReduce usefulness and benefits
- MapReduce programming model
  - Expressivity
- Relational algebra in MapReduce



# References

- J. Dean et al. *MapReduce: Simplified Data Processing on Large Clusters*. OSDI'04
- A. Pavlo et al. *A Comparison of Approaches to Large-Scale Data Analysis*. SIGMOD, 2009
- J. Dittrich et al. *Hadoop++: Making a Yellow Elephant Run Like a Cheetah (Without It Even Noticing)*. Proc. VLDB Endow. 3(1-2), 2010
- M. Stonebraker et al. *MapReduce and parallel DBMSs: friends or foes?* Communication of ACM 53(1), 2010
- S. Abiteboul et al. *Web data management*. Cambridge University Press, 2011
- A. Rajaraman et al. *Mining massive data sets*. Cambridge University Press, 2012
- P. Sadagale and M. Fowler. *NoSQL distilled*. Addison-Wesley, 2013