Big Data Architectures

Big Data Management

Knowledge objectives

1. Explain the problema of a spaghetti architecture
2. Explain the need of the Lambda architecture
3. Explain the difference between the Kappa and Lambda architectures
4. Justify the need of a Data Lake
5. Identify the difficulties of a Data Lake
6. Explain the need of each component in the Bolster architecture
7. Map the components of Bolster to a RDBMS architecture

Application Objectives

1. Given a use case, define its software architecture

Problem definition

Data Management (I)

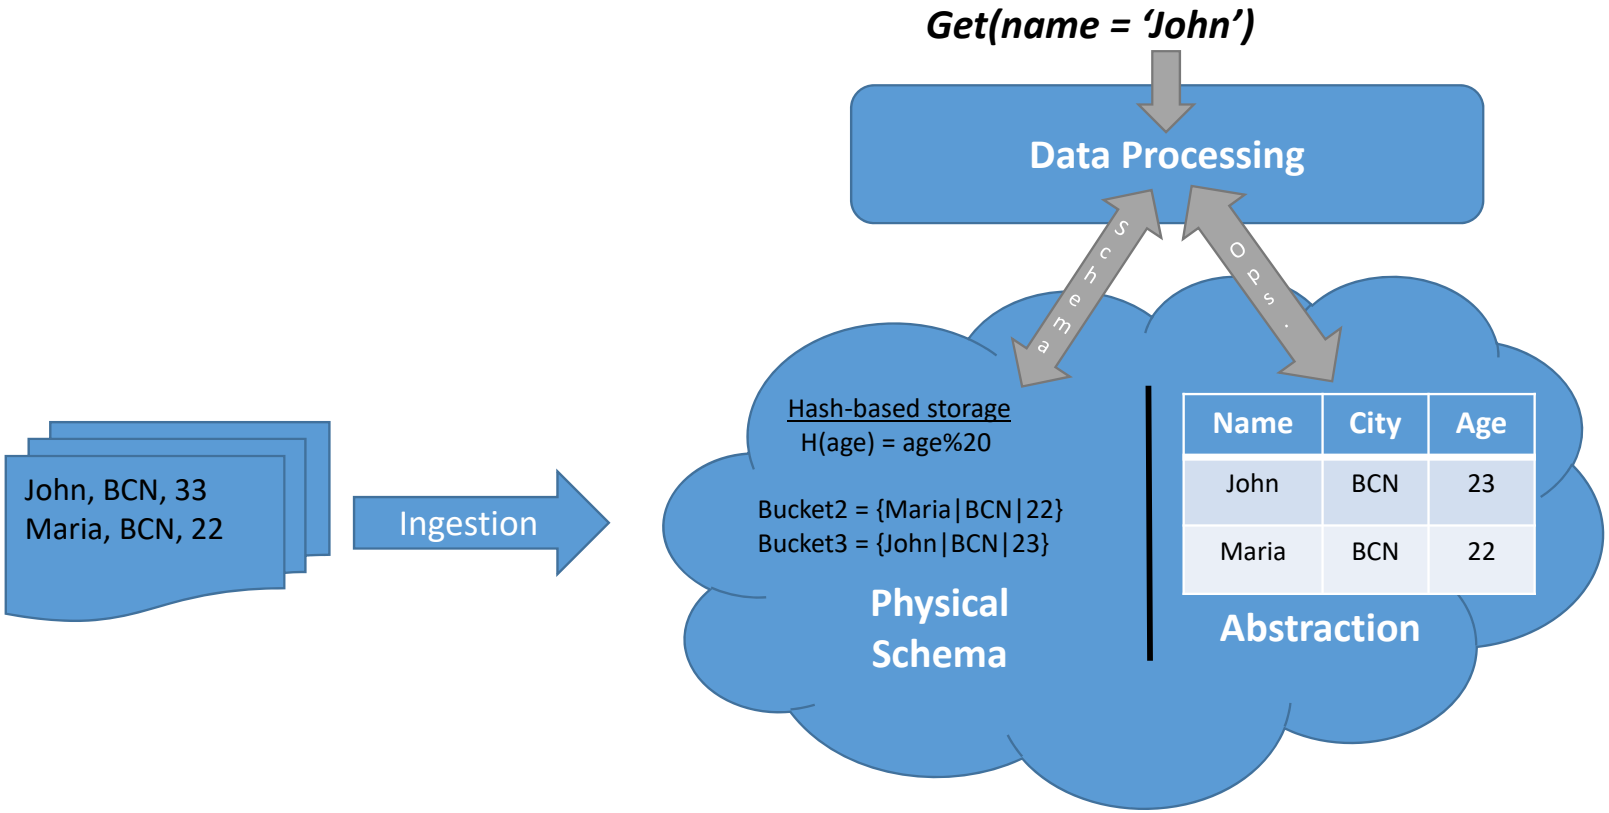

Data Management (II)

Data management refers to the features a DBMS must provide:

- Ingestion: means provided to insert /upload data
  - E.g., ORACLE SQL*Loader
- Storage: format/structures used to persist data
  - E.g., hash, B-tree, heap file
- Modelling: arrangement of data within the available structures
  - E.g., normalization, partitioning
- Processing: means provided to manipulate data
  - E.g., PL/SQL
- Querying/fetching: means provided to allow users to retrieve data
  - E.g., SQL, relational algebra

In Big Data settings, they are the same concepts but assuming NOSQL underneath

- Typically, a distributed system
- Possibly with an alternative data model to the Relational one
- Implementing ad-hoc architectural solutions

Big Data Architectures

- Question the main principles of traditional DB architectures
- Implement from scratch the whole stack
  - Ingestion, Storage, Modeling, Processing, and Querying
- Use new trendy technological features
  - Primary indexes to implement the global catalog
    - Distributed B+
    - Consistent Hashing
  - In-memory processing
  - Columnar block iteration: vertical fragmentation + fixed-size values + RLE compression
    - Heavily exploited by column-oriented databases
    - Good for read-only workloads
  - Sequential reads for large workloads
    - Key design
    - Take the most out of databases by boosting sequential reads
      - Enables pre-fetching
      - Option to maximize the effective read ratio (by a good db design)

The Multi-Project Approach

- The DBMS tasks can be spread over different systems
  - Independent
  - Heterogeneous
- Hadoop as paradigmatic case:
  - Ingestion: Kafka
  - Storage: HDFS + Hbase
  - Modeling: Hive + HCatalog
  - Processing: Spark
  - Querying: Spark SQL

Big Data Landscape

Spaghetti architecture

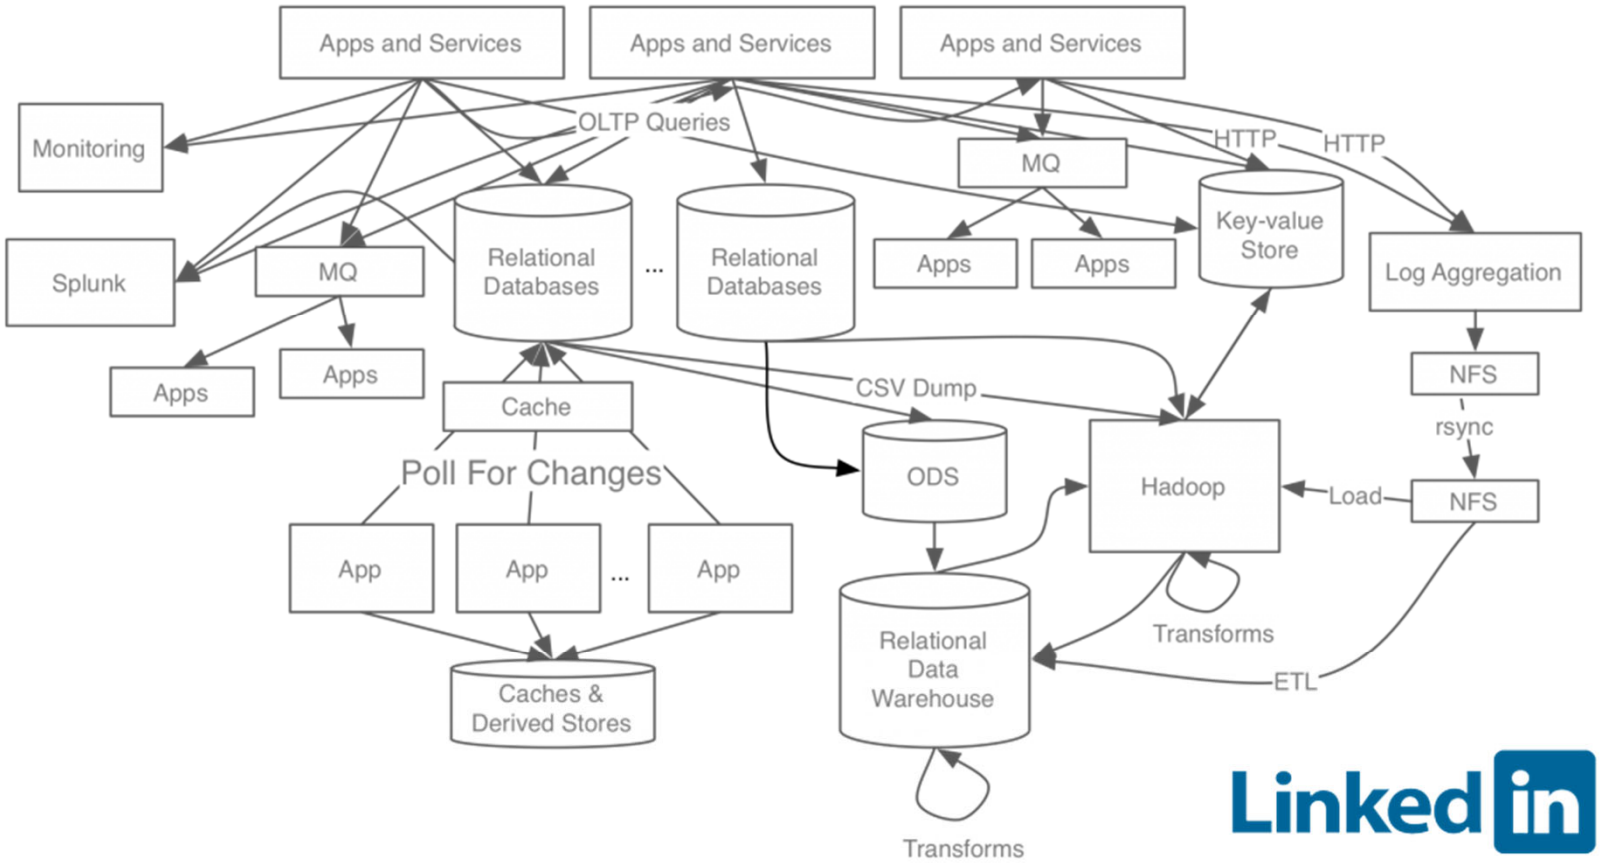

https://www.confluent.io/blog/event-streaming-platform-1

New Processing Architectural Patterns

Data-centered architecture

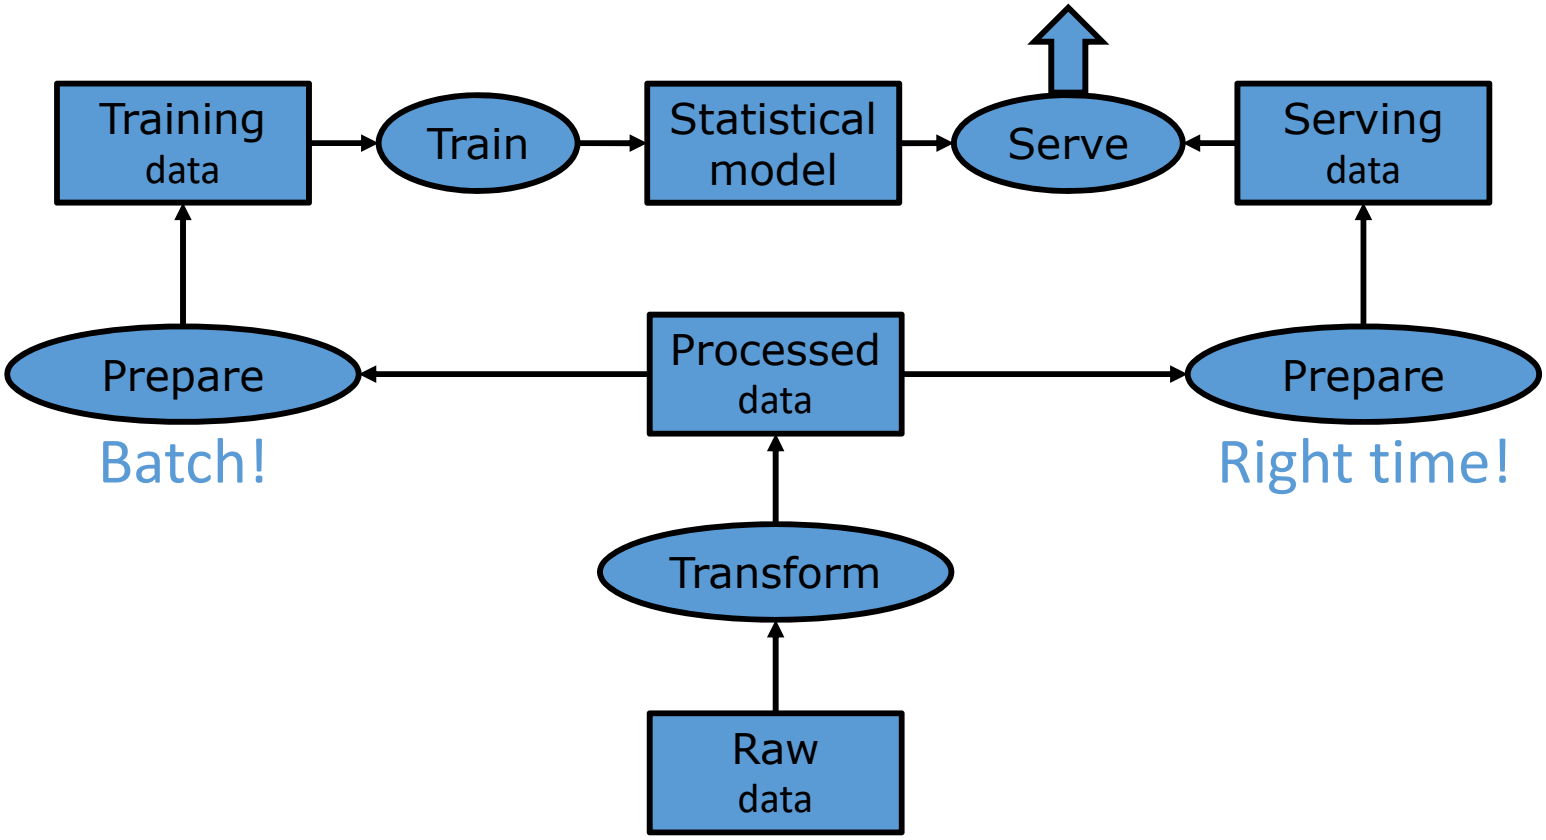

λ-Architecture

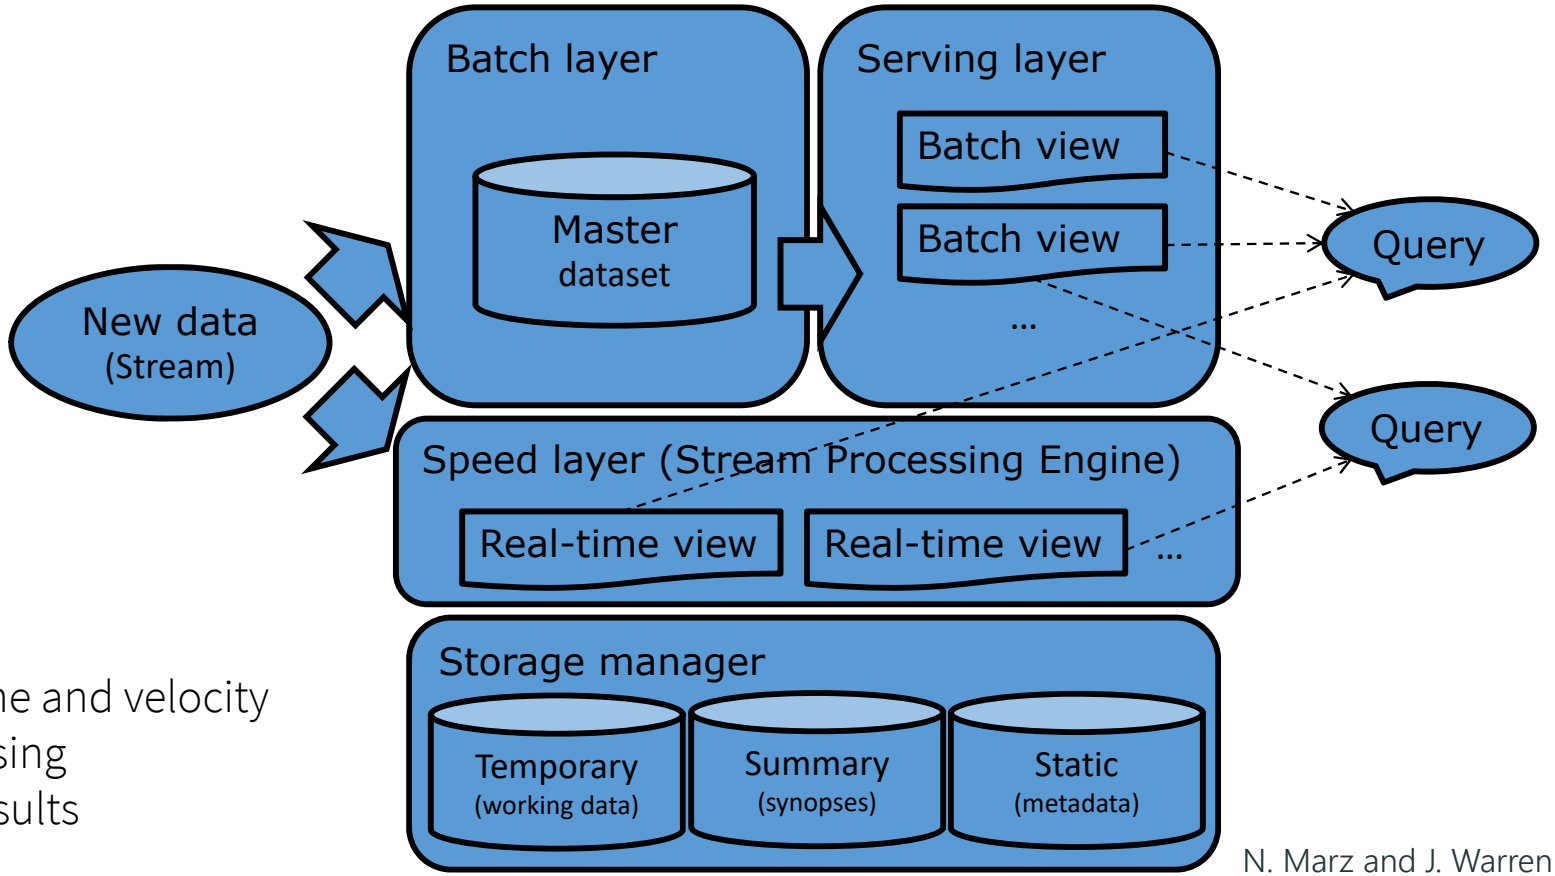

IDEA: Accommodate volume and velocity
Real time Vs. Batch processing
Precise Vs. Approximate results

N. Marz and J. Warren

κ-architecture

- Data is considered to be a never-ending stream

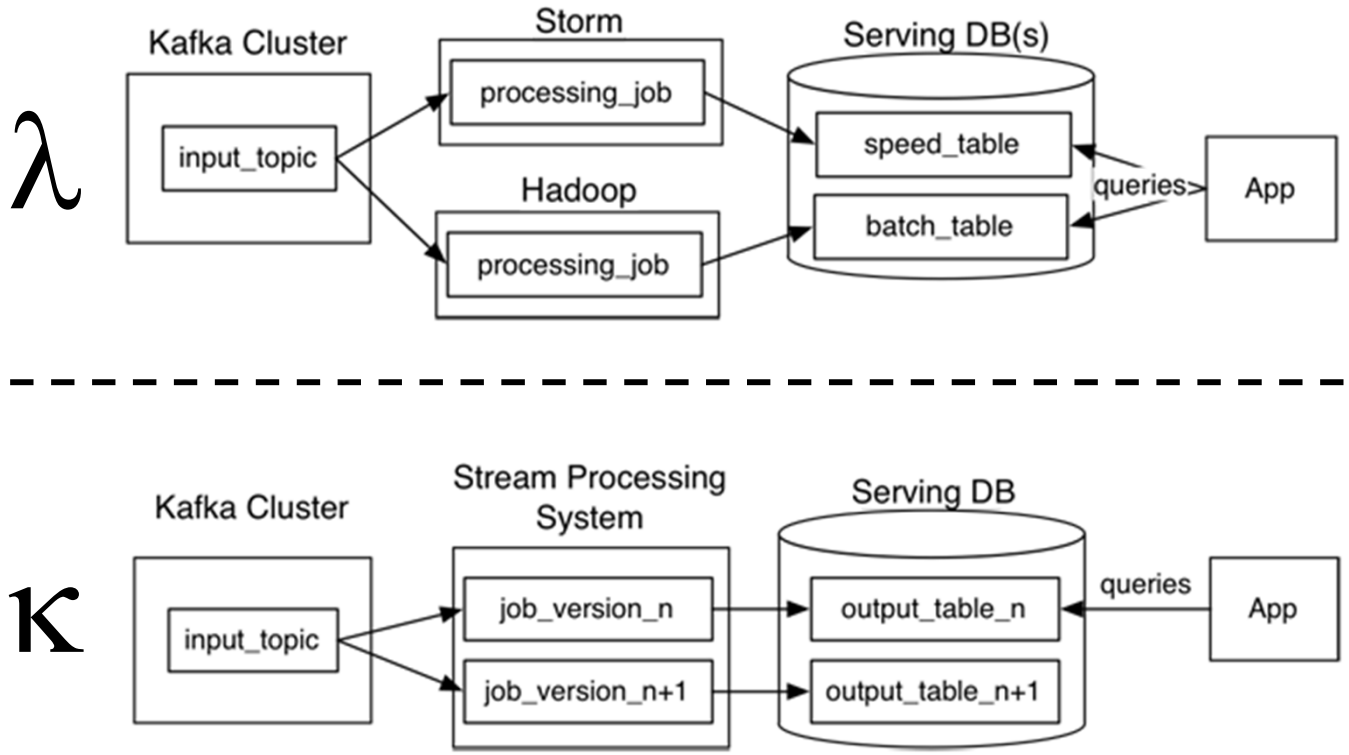

Jay Kreps and O'Reilley

New Storage Architectural Pattern

From data warehousing to data lakes

The Data Lake

- IDEA: Load-first, Model-Later
- Modeling at load time restricts the potential analysis that can be done later (Big Analytics)
- Store raw data and create on-demand views to handle with precise analysis needs

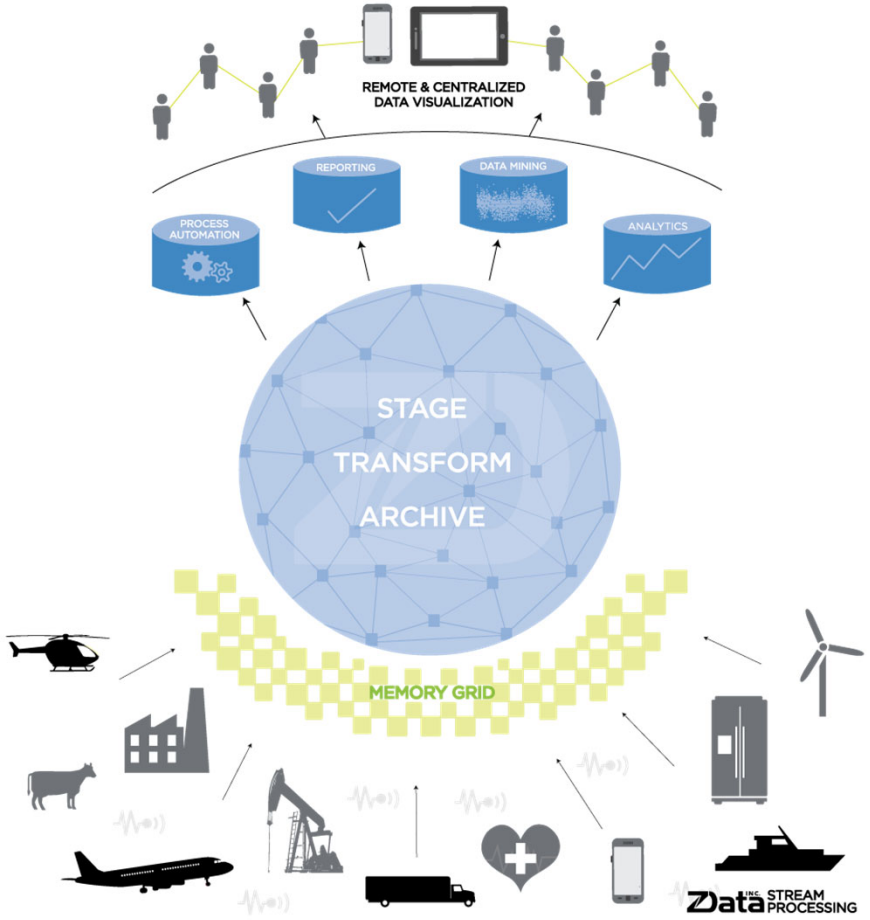

Model-First (Load-Later)

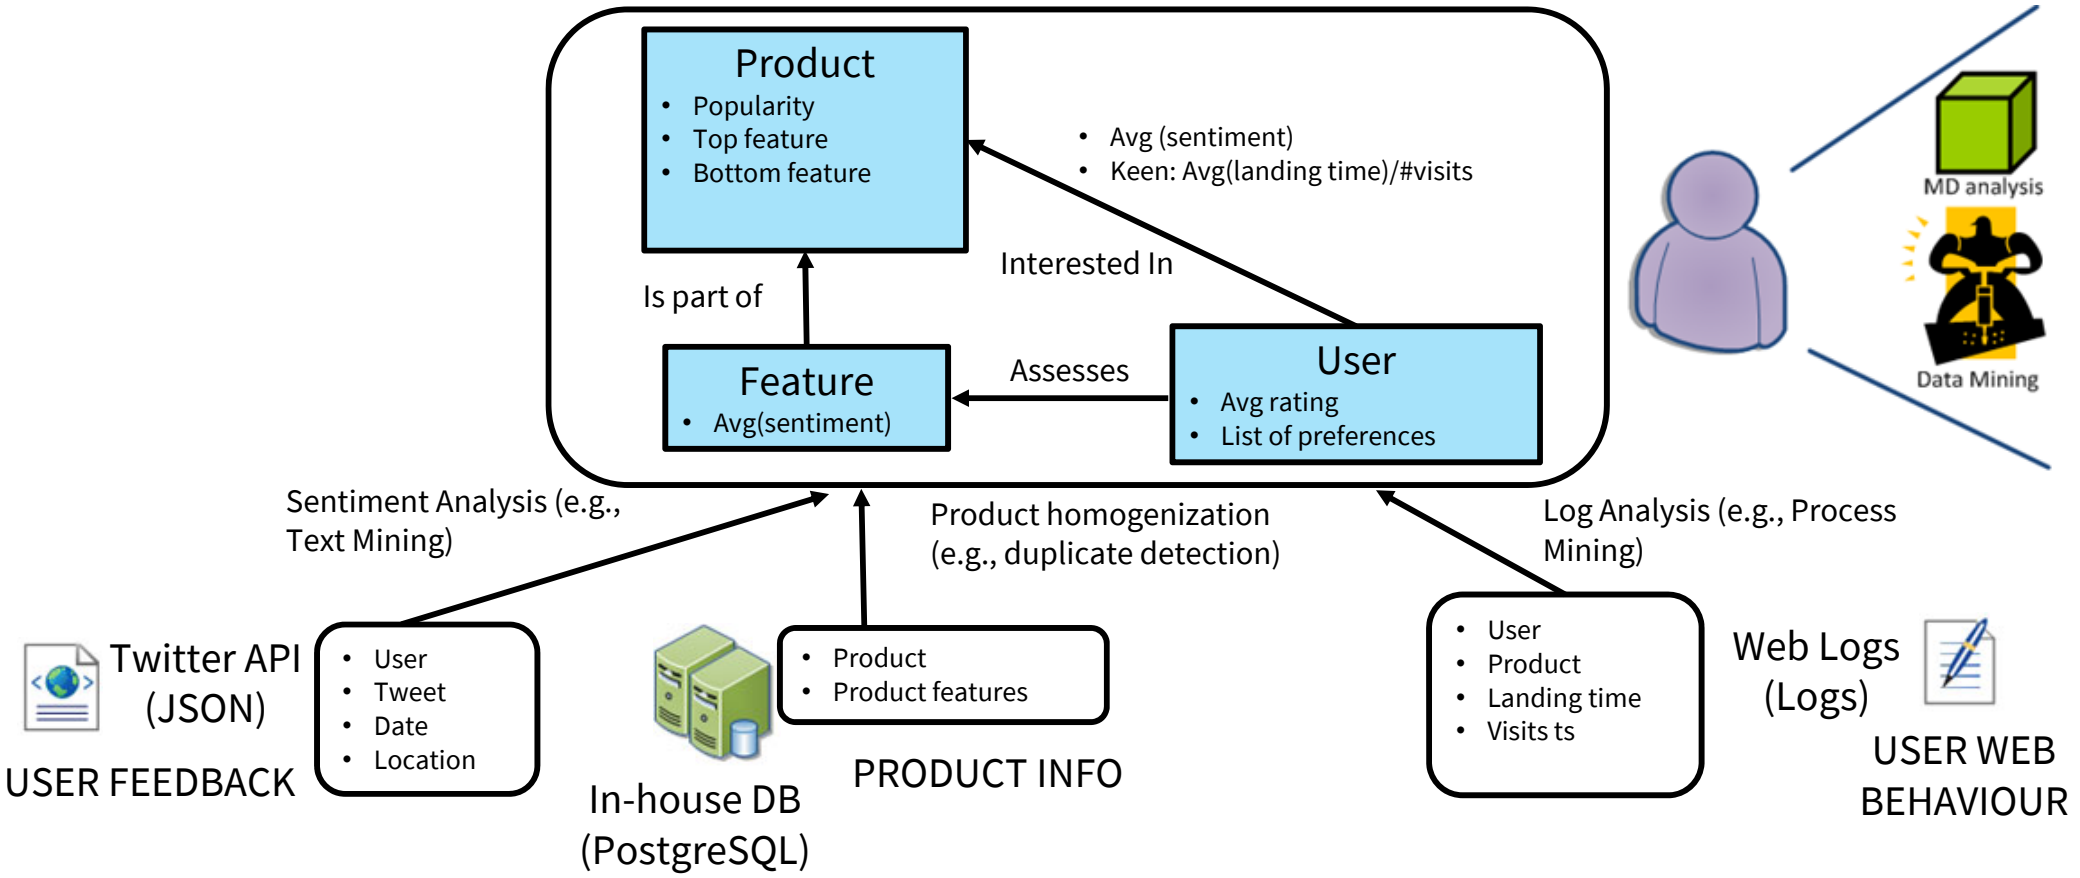

Drawbacks of Model-First (Load-Later)

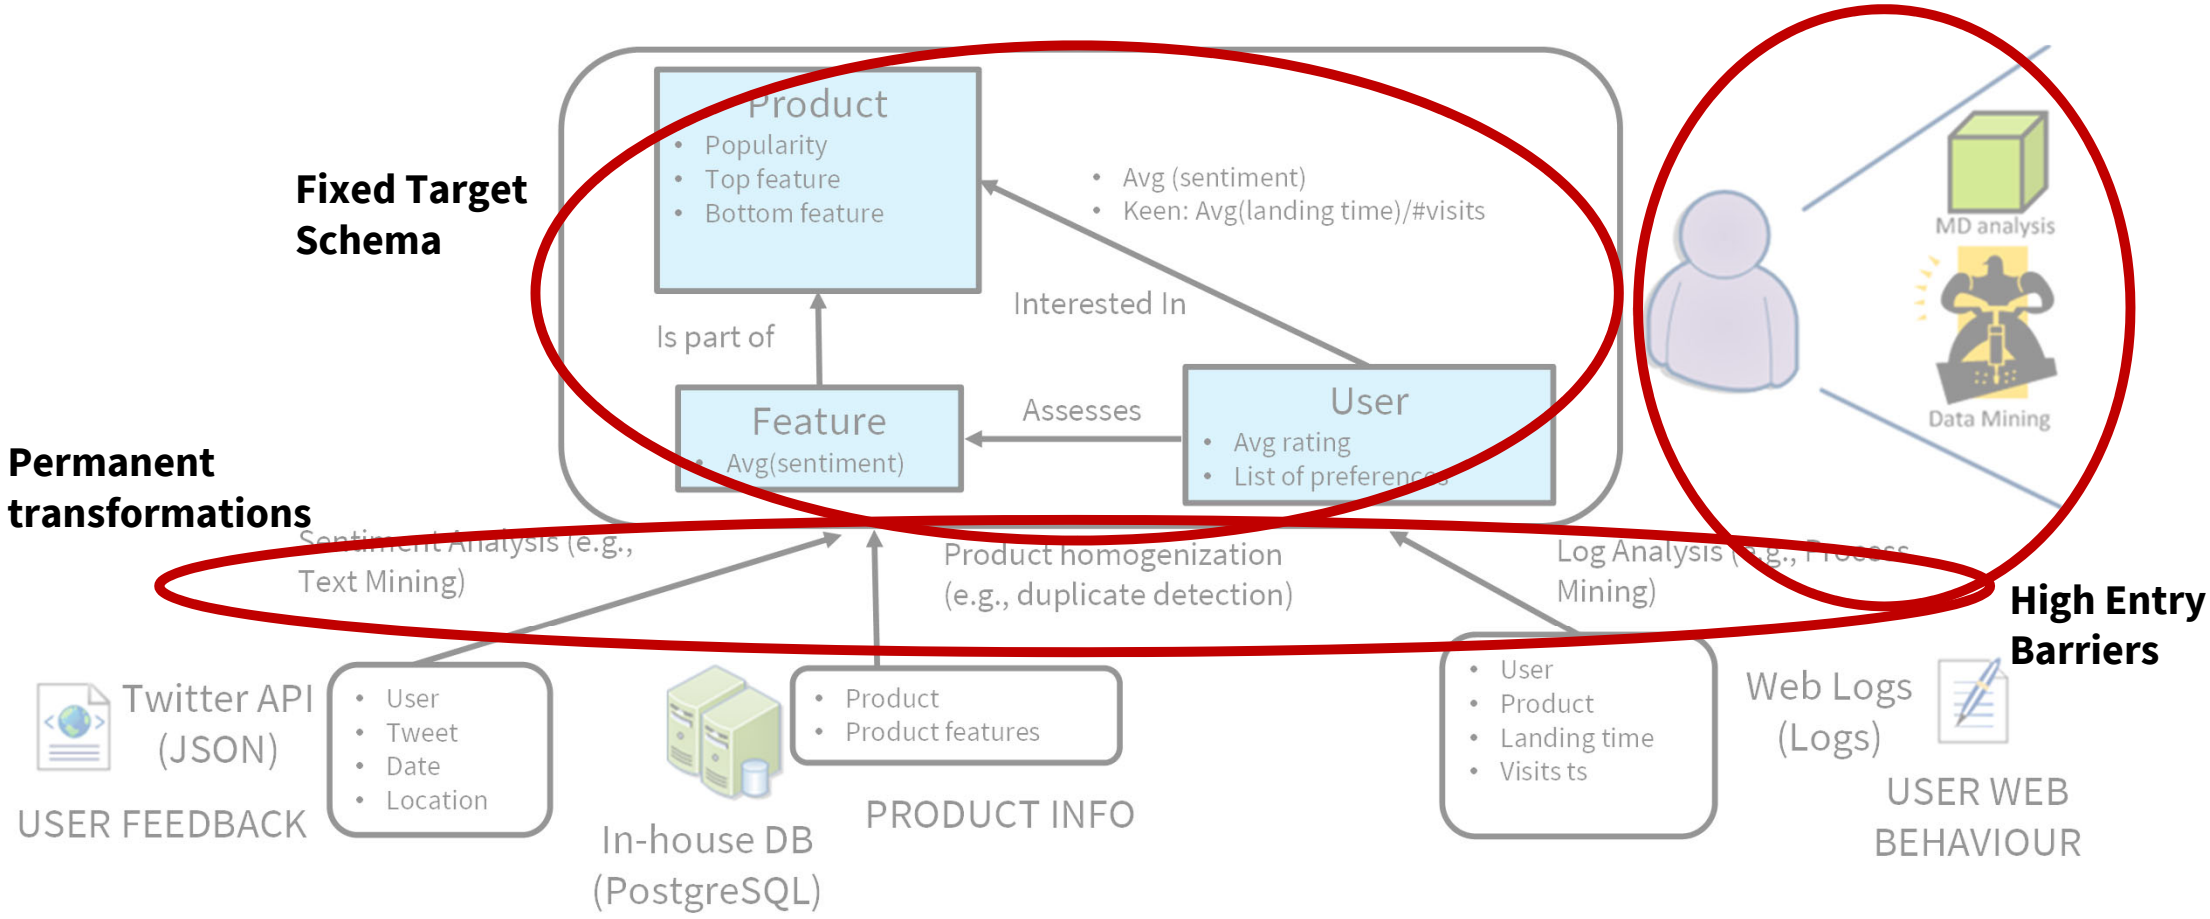

Load-First (Model-Later)

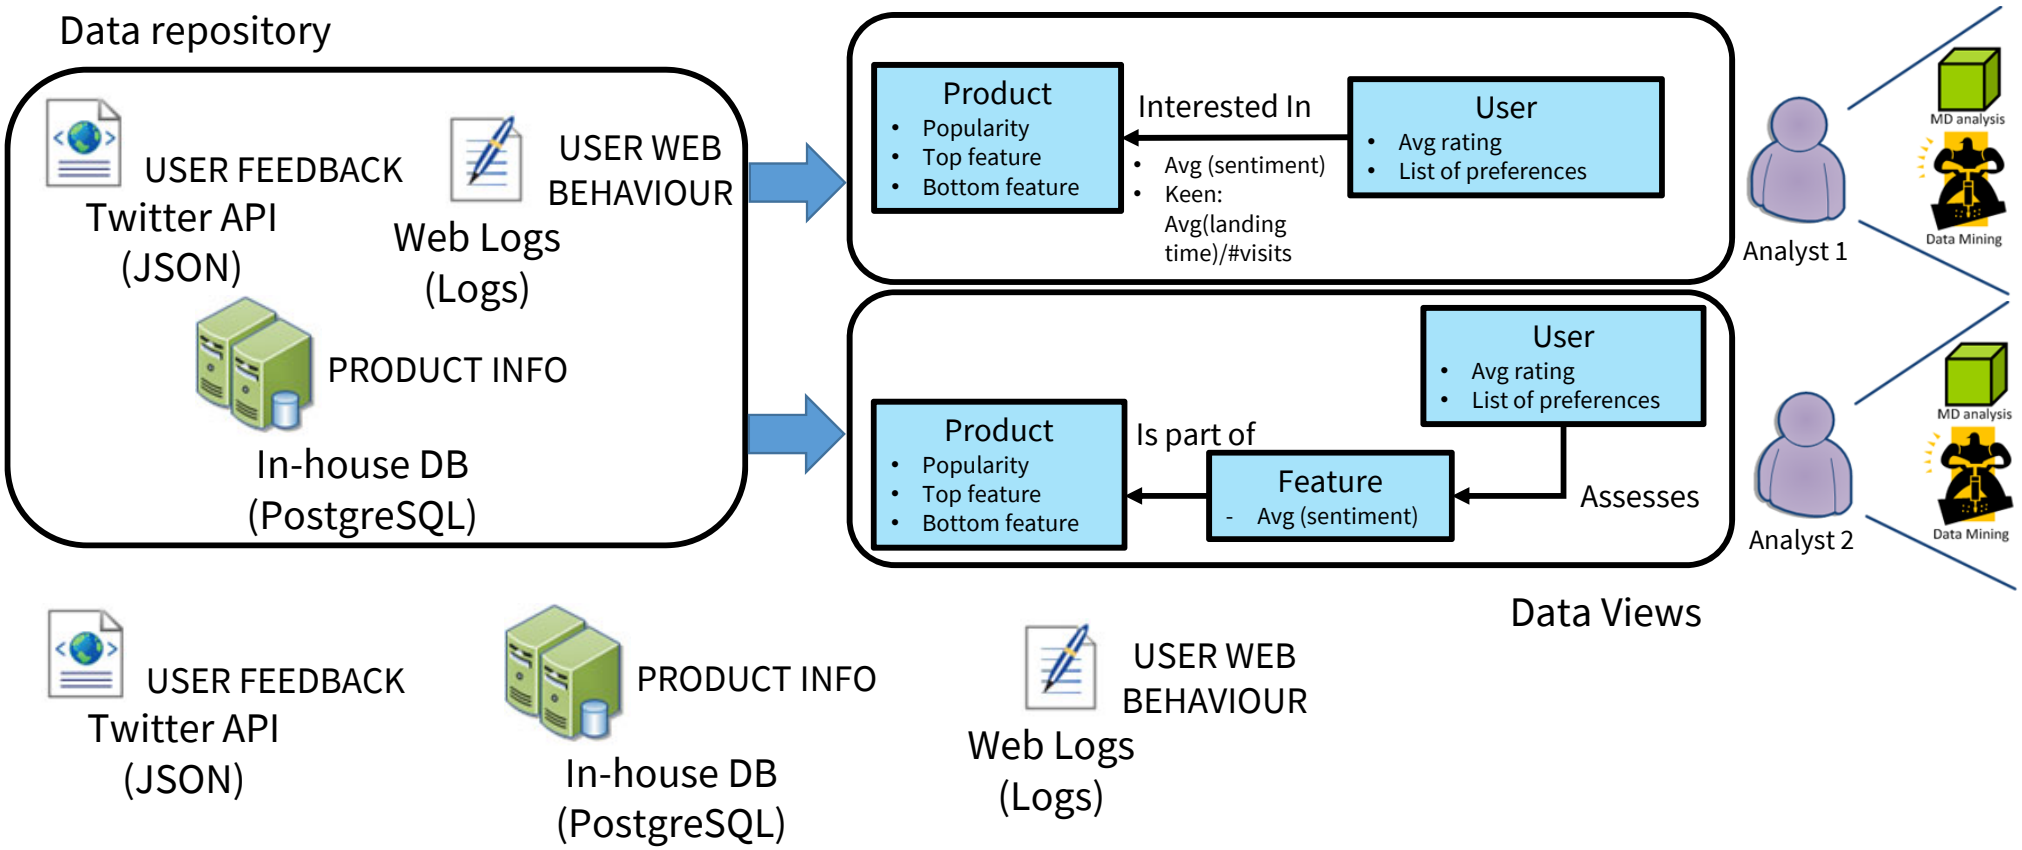

Drawbacks of Load-First (Model-Later)

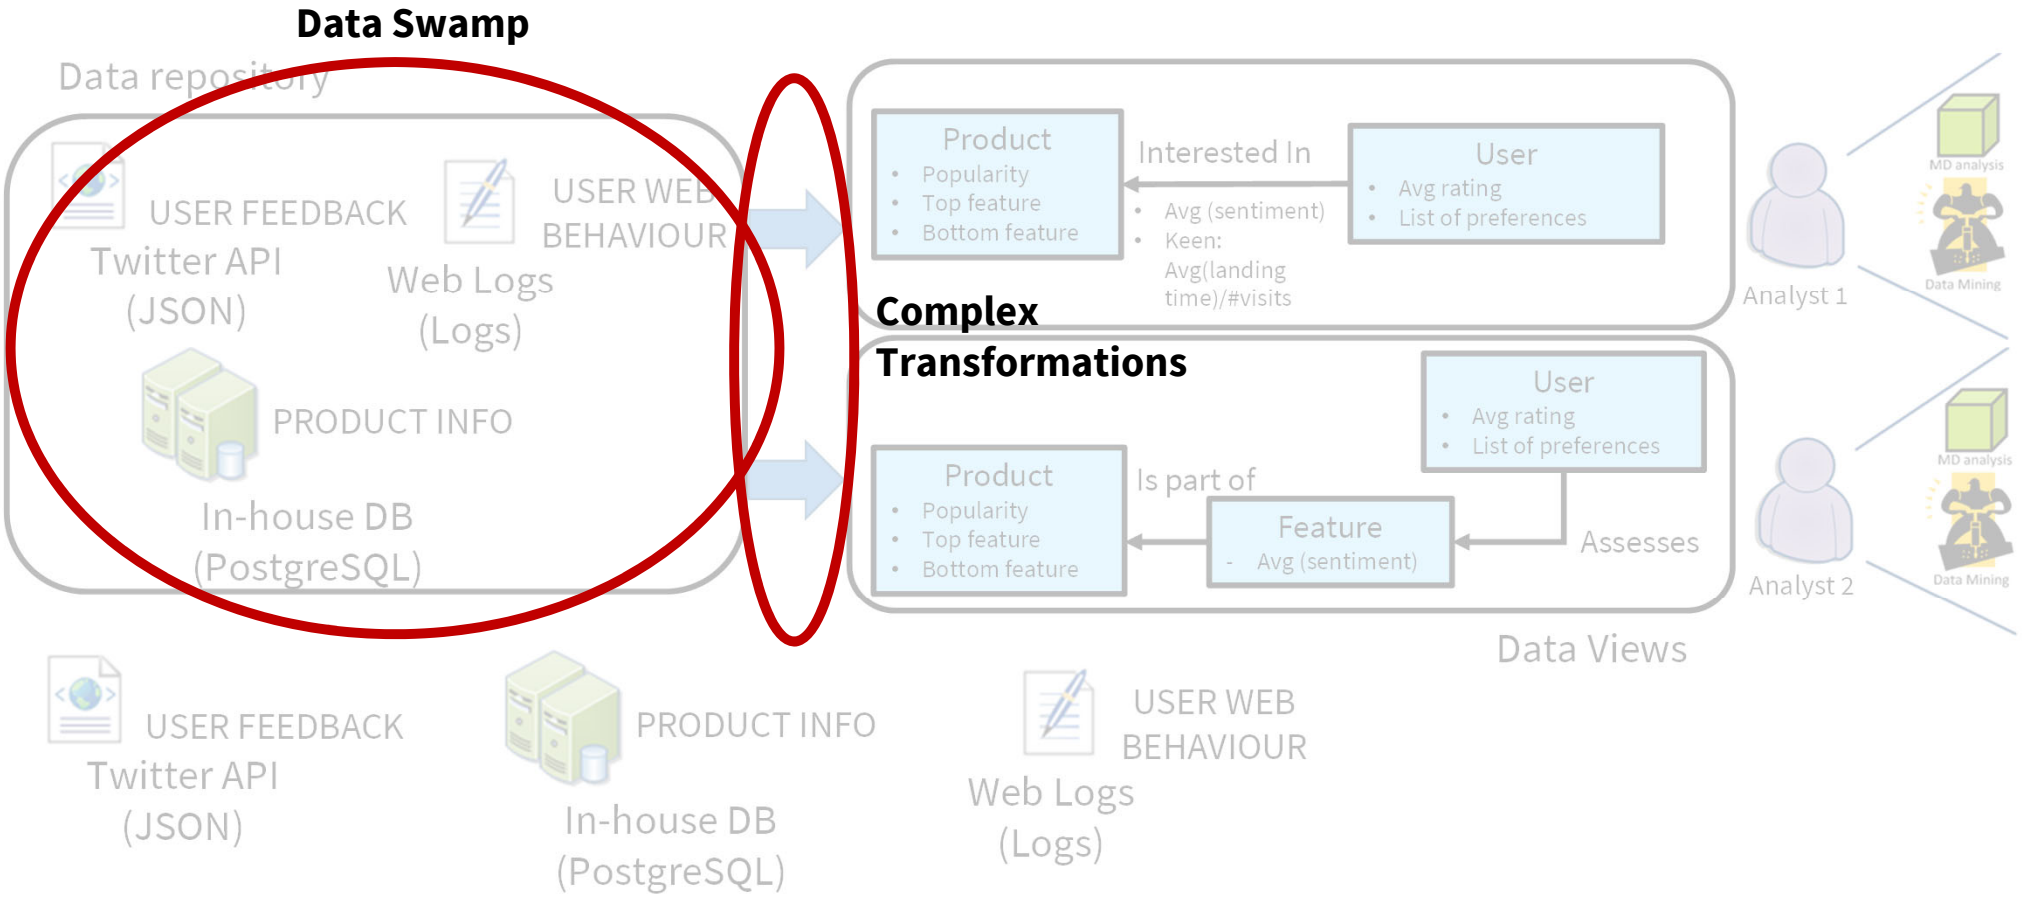

Stonebraker (2014)

Towards semantic-awareness

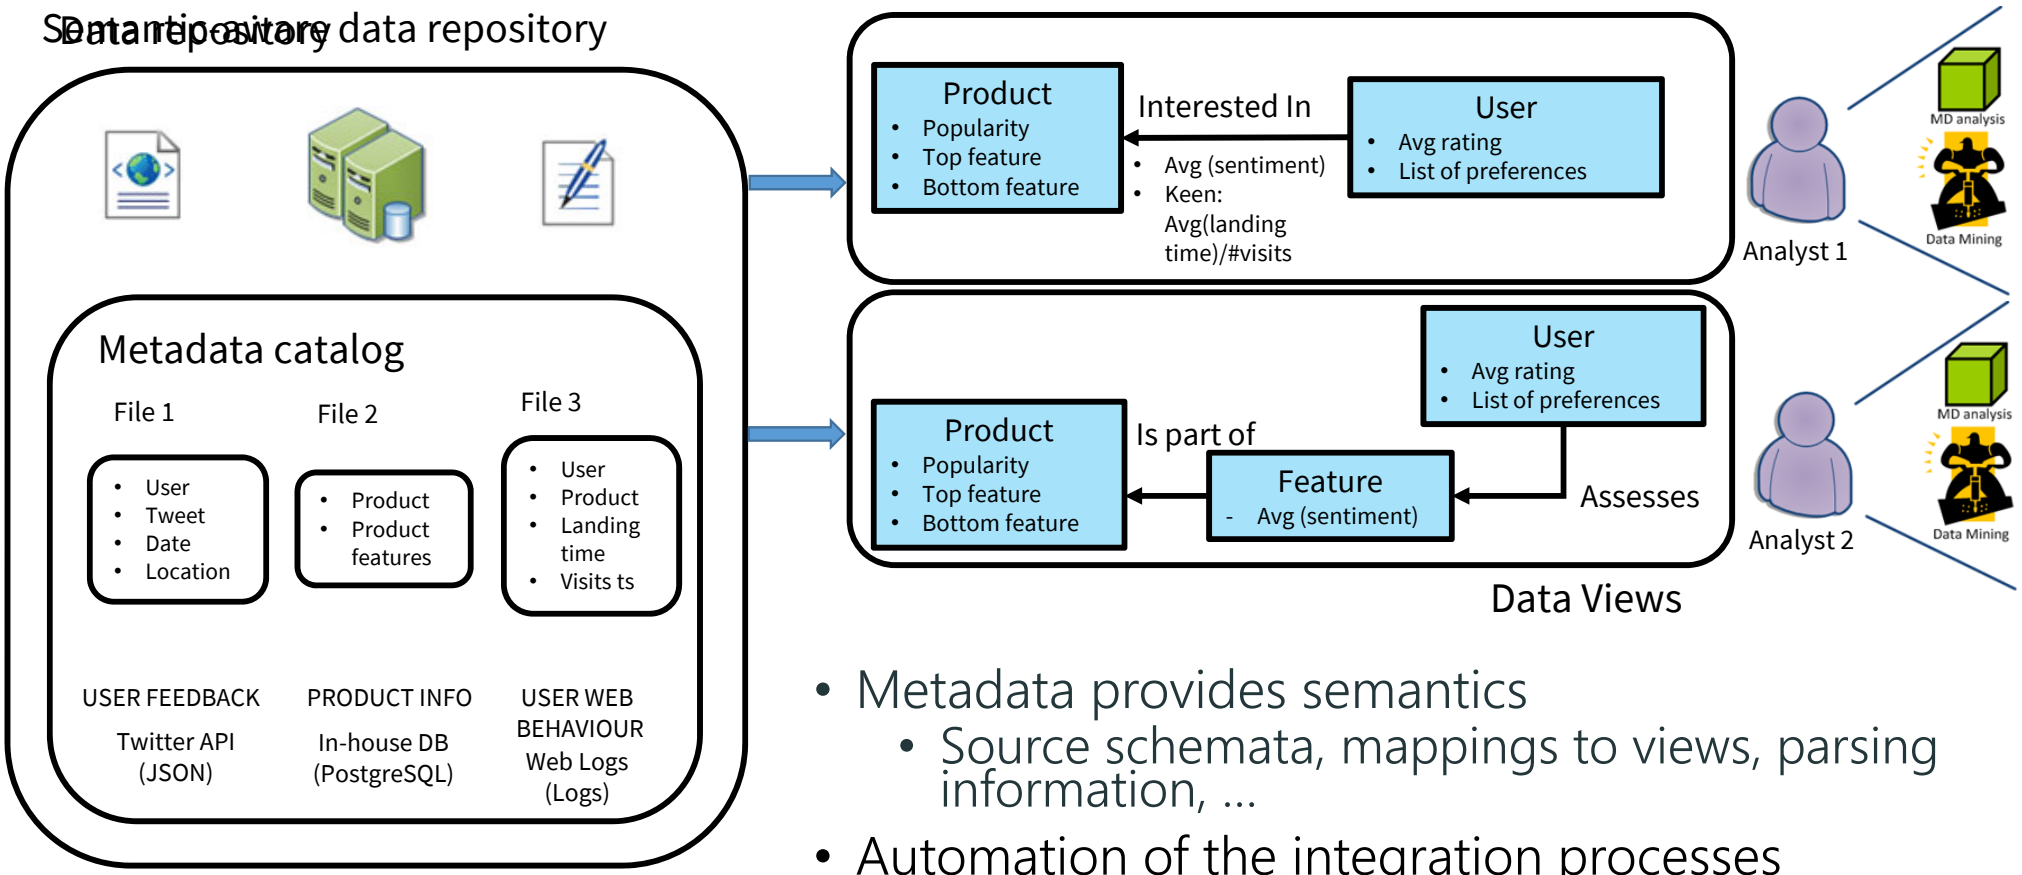

- Metadata provides semantics
  - Source schemata, mappings to views, parsing information, …
- Automation of the integration processes

From IT-Centered to User-Centered

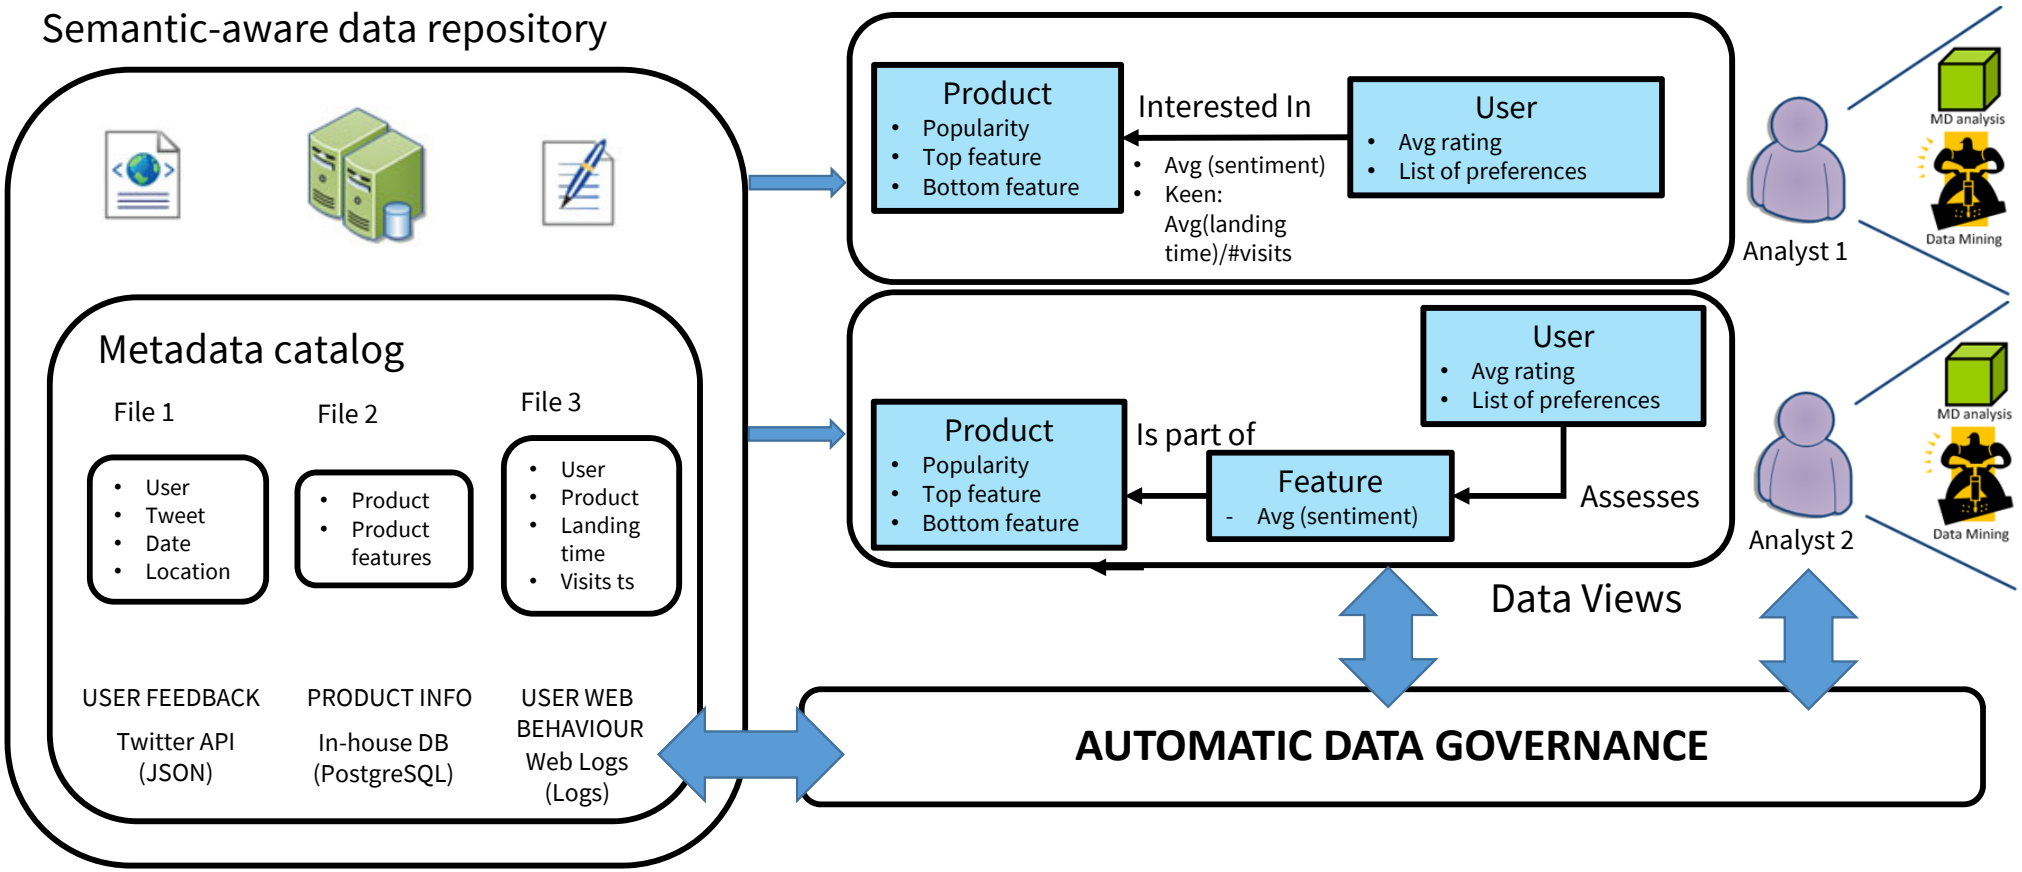

The Missing Link: Metadata

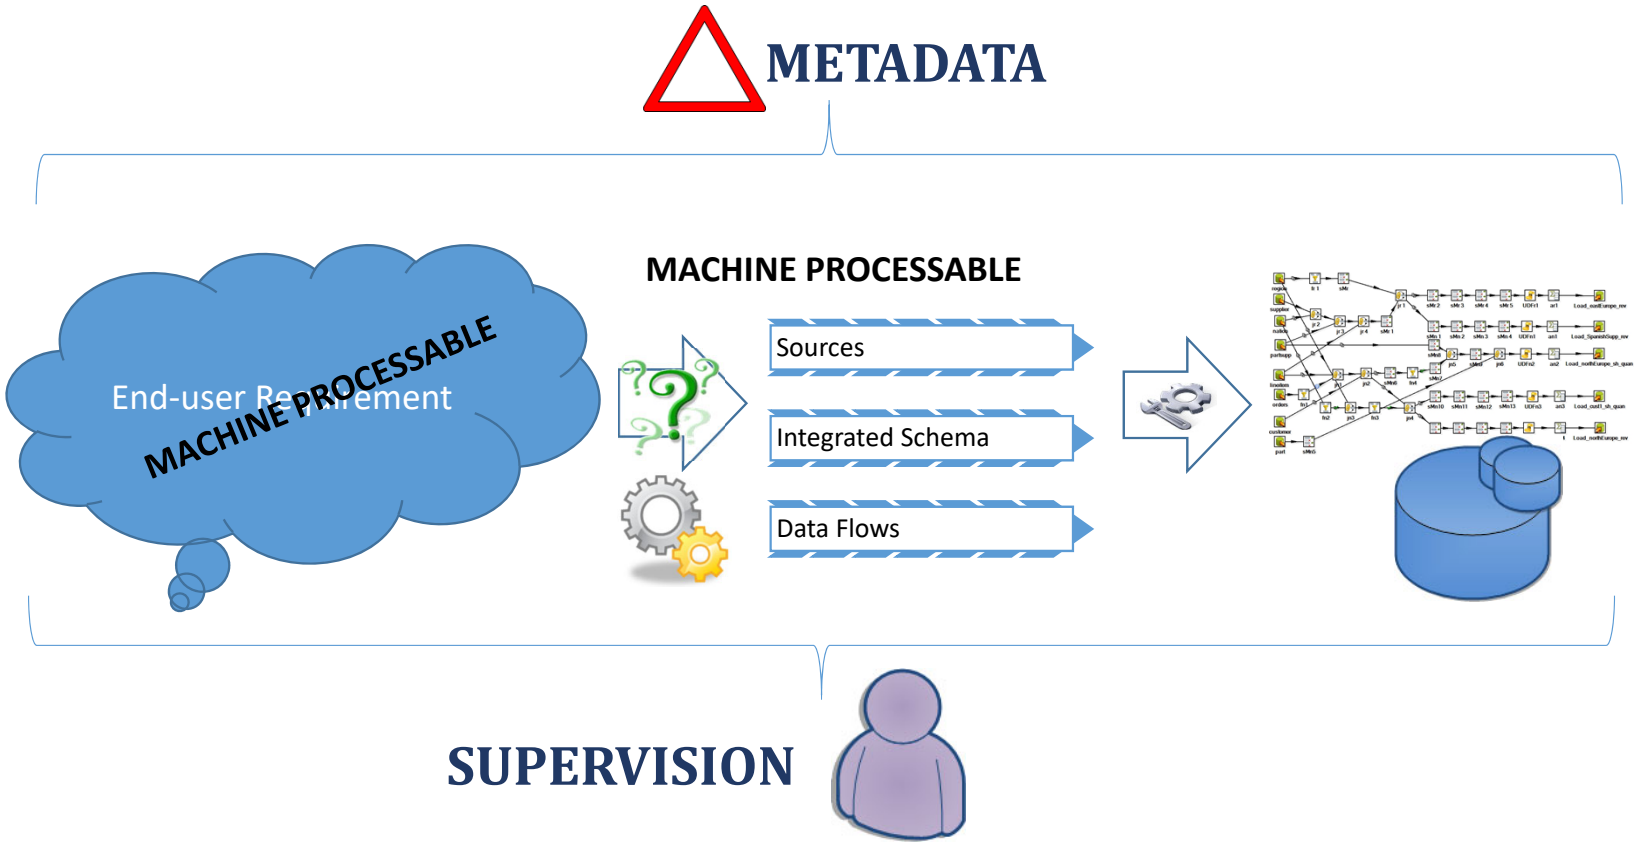

Database Management System view

Centralized DBMS Architecture

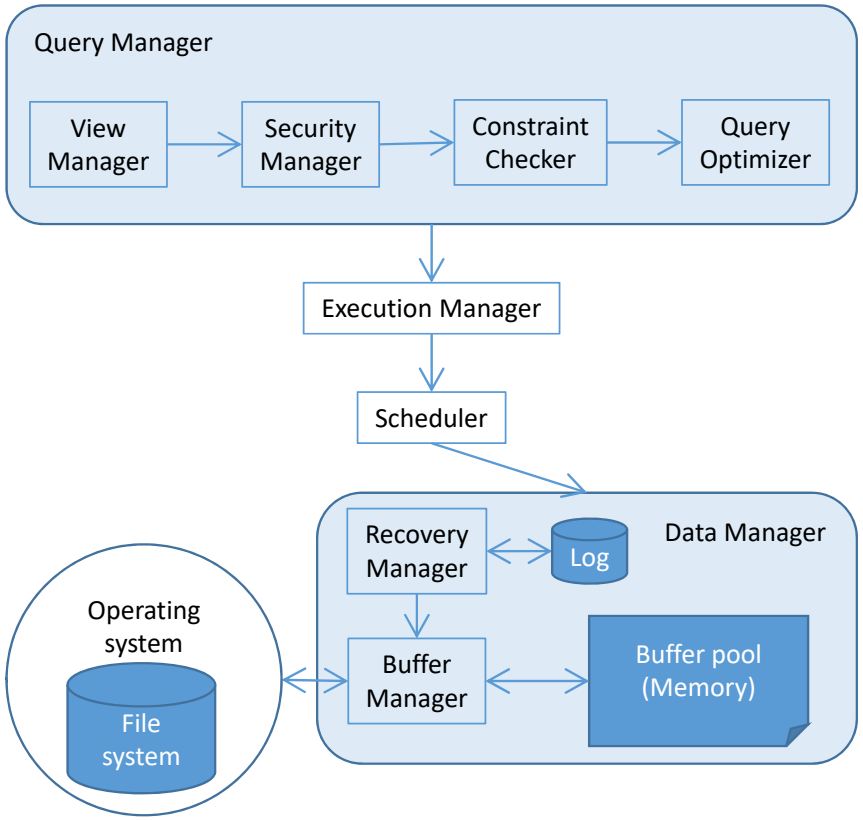

Centralized DBMS Architecture

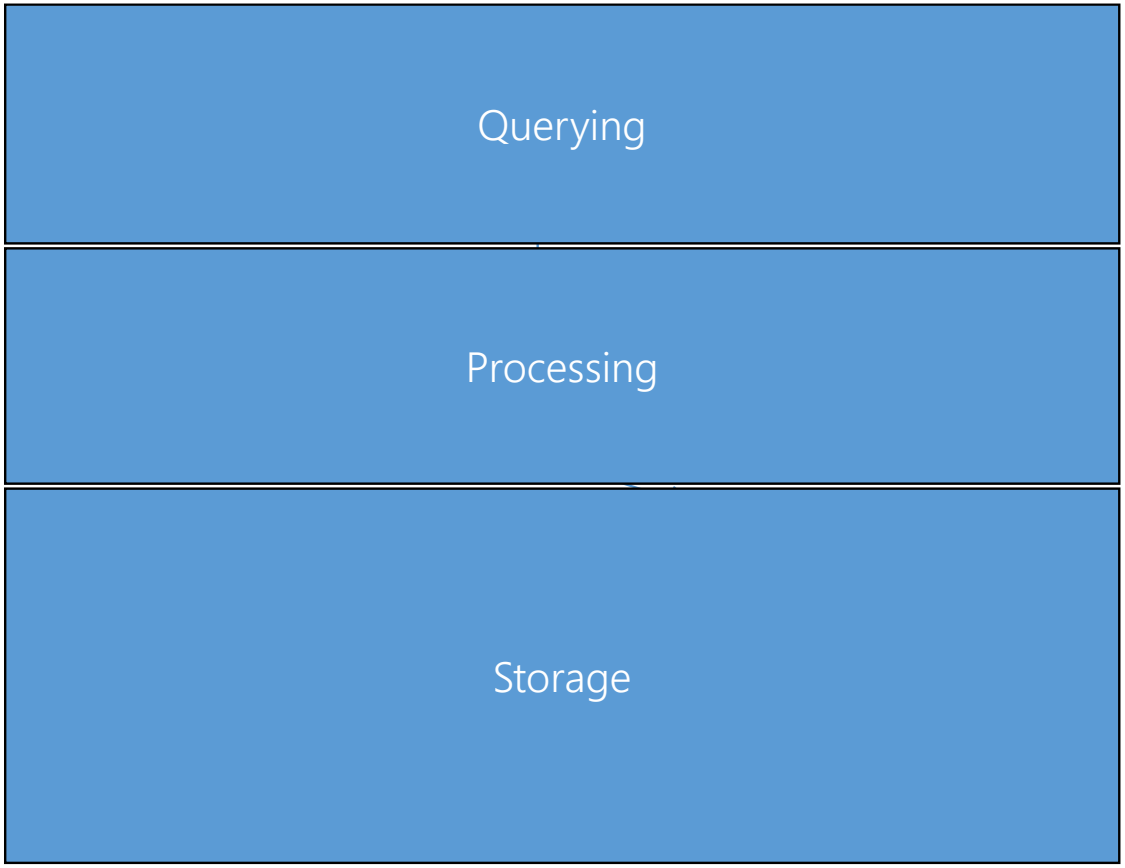

Distributed DBMS Architecture

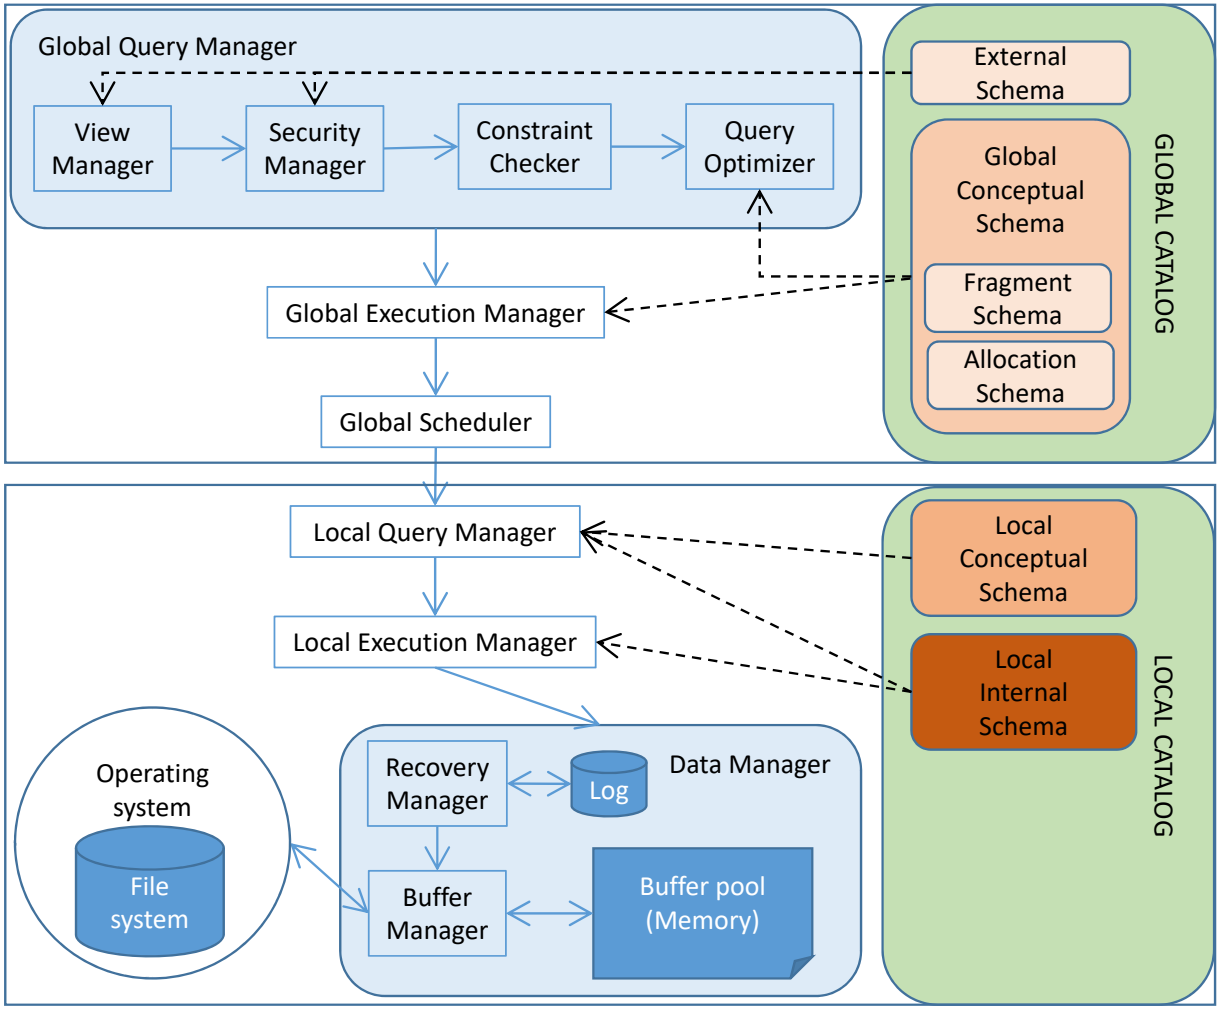

Distributed DBMS Architecture

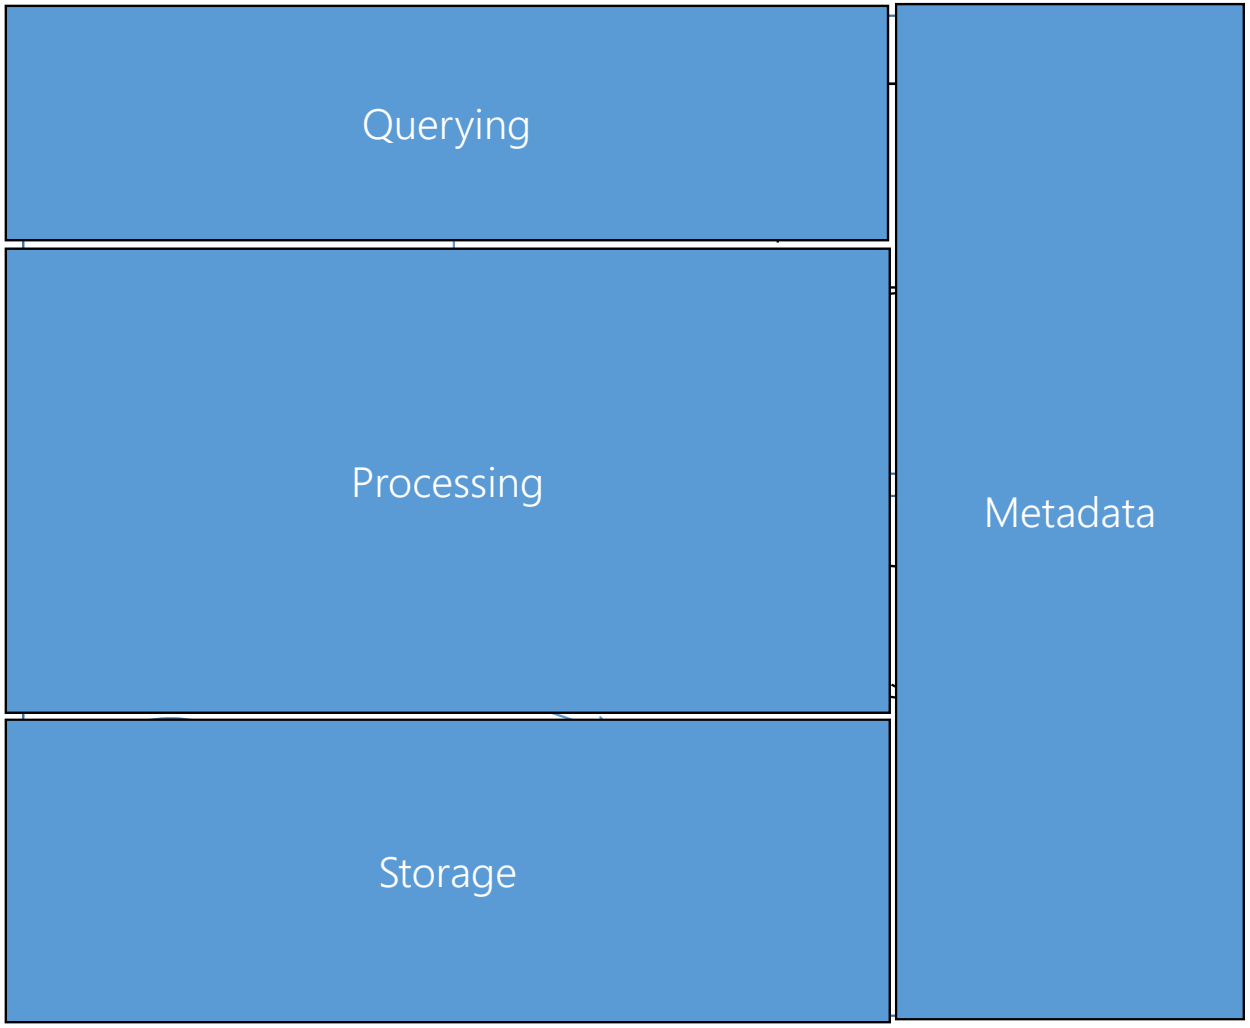

Bolster

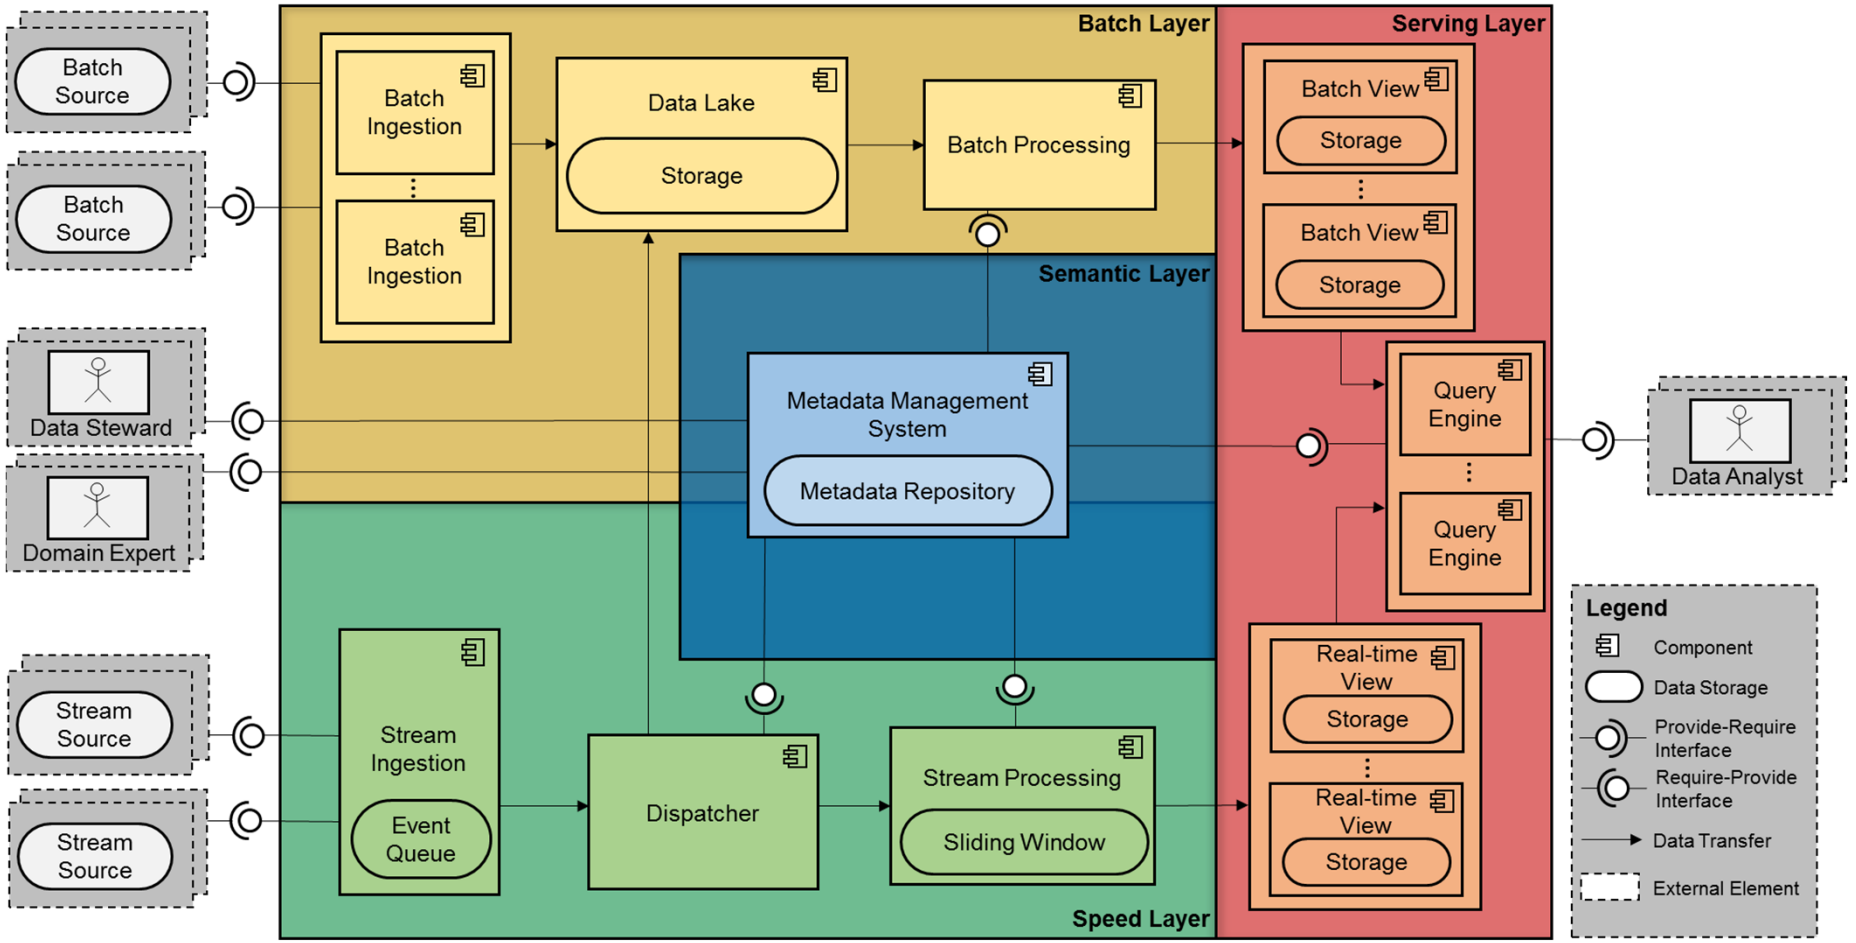
Batch Layer
Serving Layer
Semantic Layer
Speed Layer
Batch Source
Batch Source
Data Steward
Domain Expert
Stream Source
Stream Source
Batch Ingestion
Batch Ingestion
Data Lake
Storage
Batch Processing
Batch View
Storage
Batch View
Storage
Metadata Management System
Metadata Repository
Query Engine
Query Engine
Data Analyst
Stream Ingestion
Event Queue
Dispatcher
Stream Processing
Sliding Window
Real-time View
Storage
Real-time View
Storage
Legend
Component
Data Storage
Provide-Require Interface
Require-Provide Interface
Data Transfer
External Element

Bolster

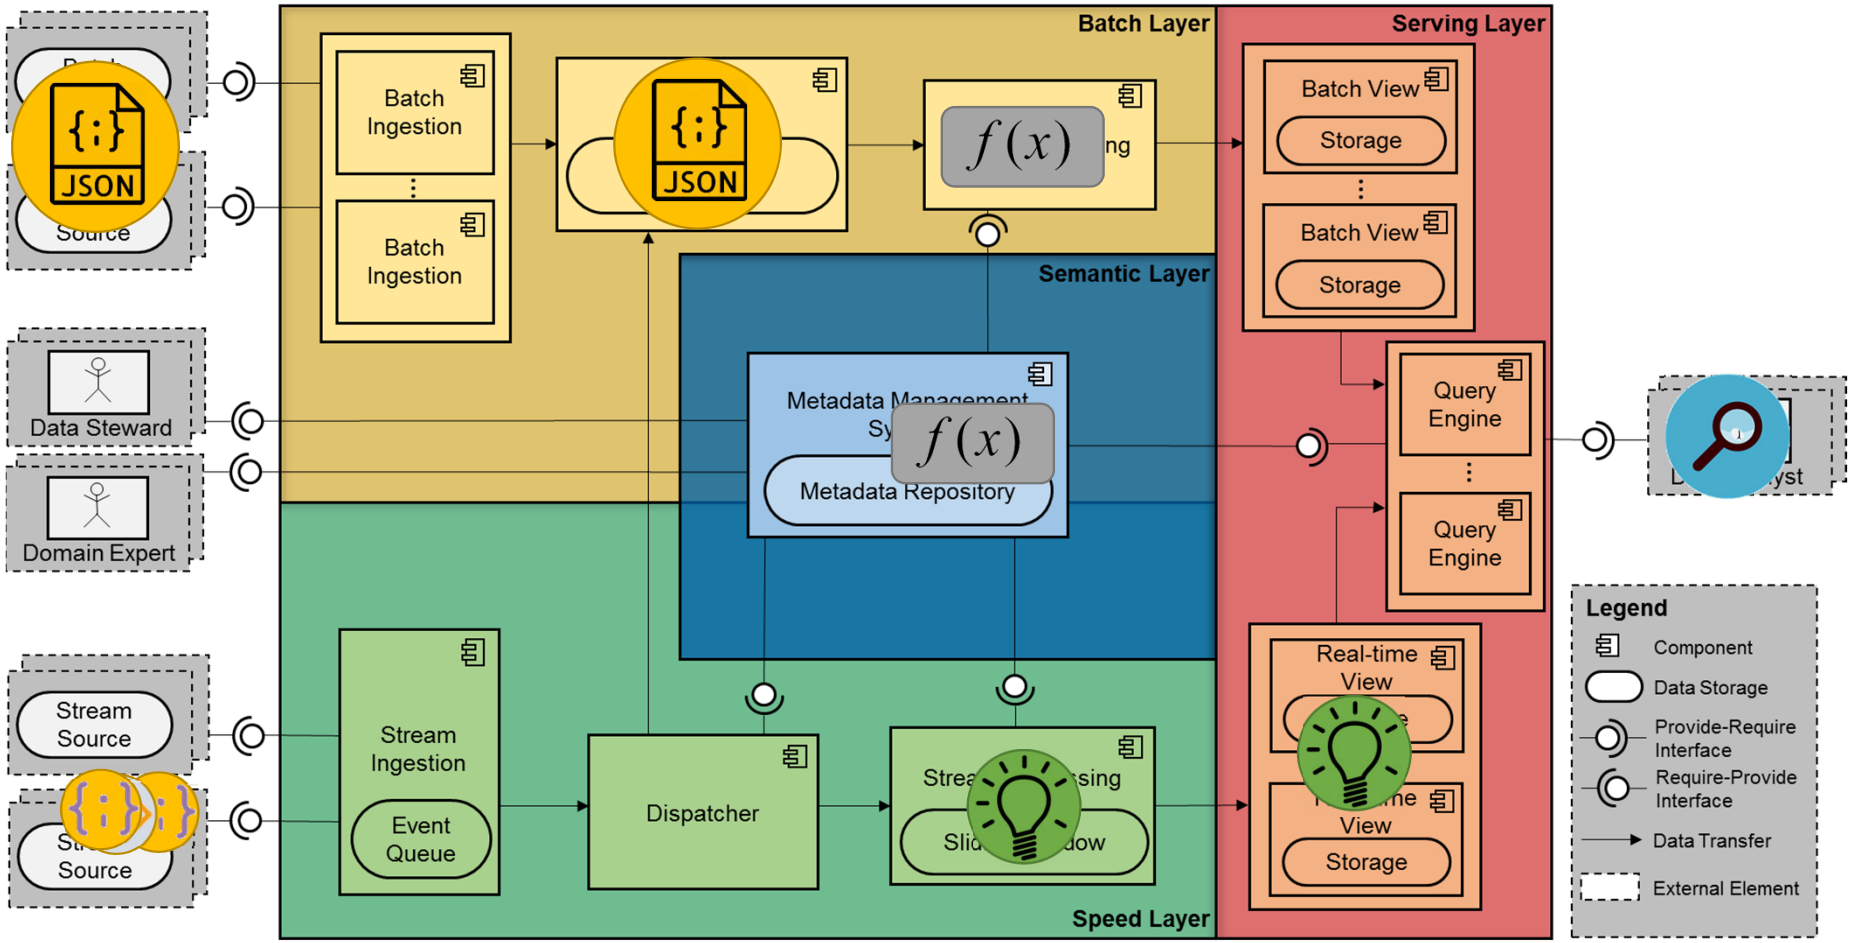

S. Nadal et al.

Bolster

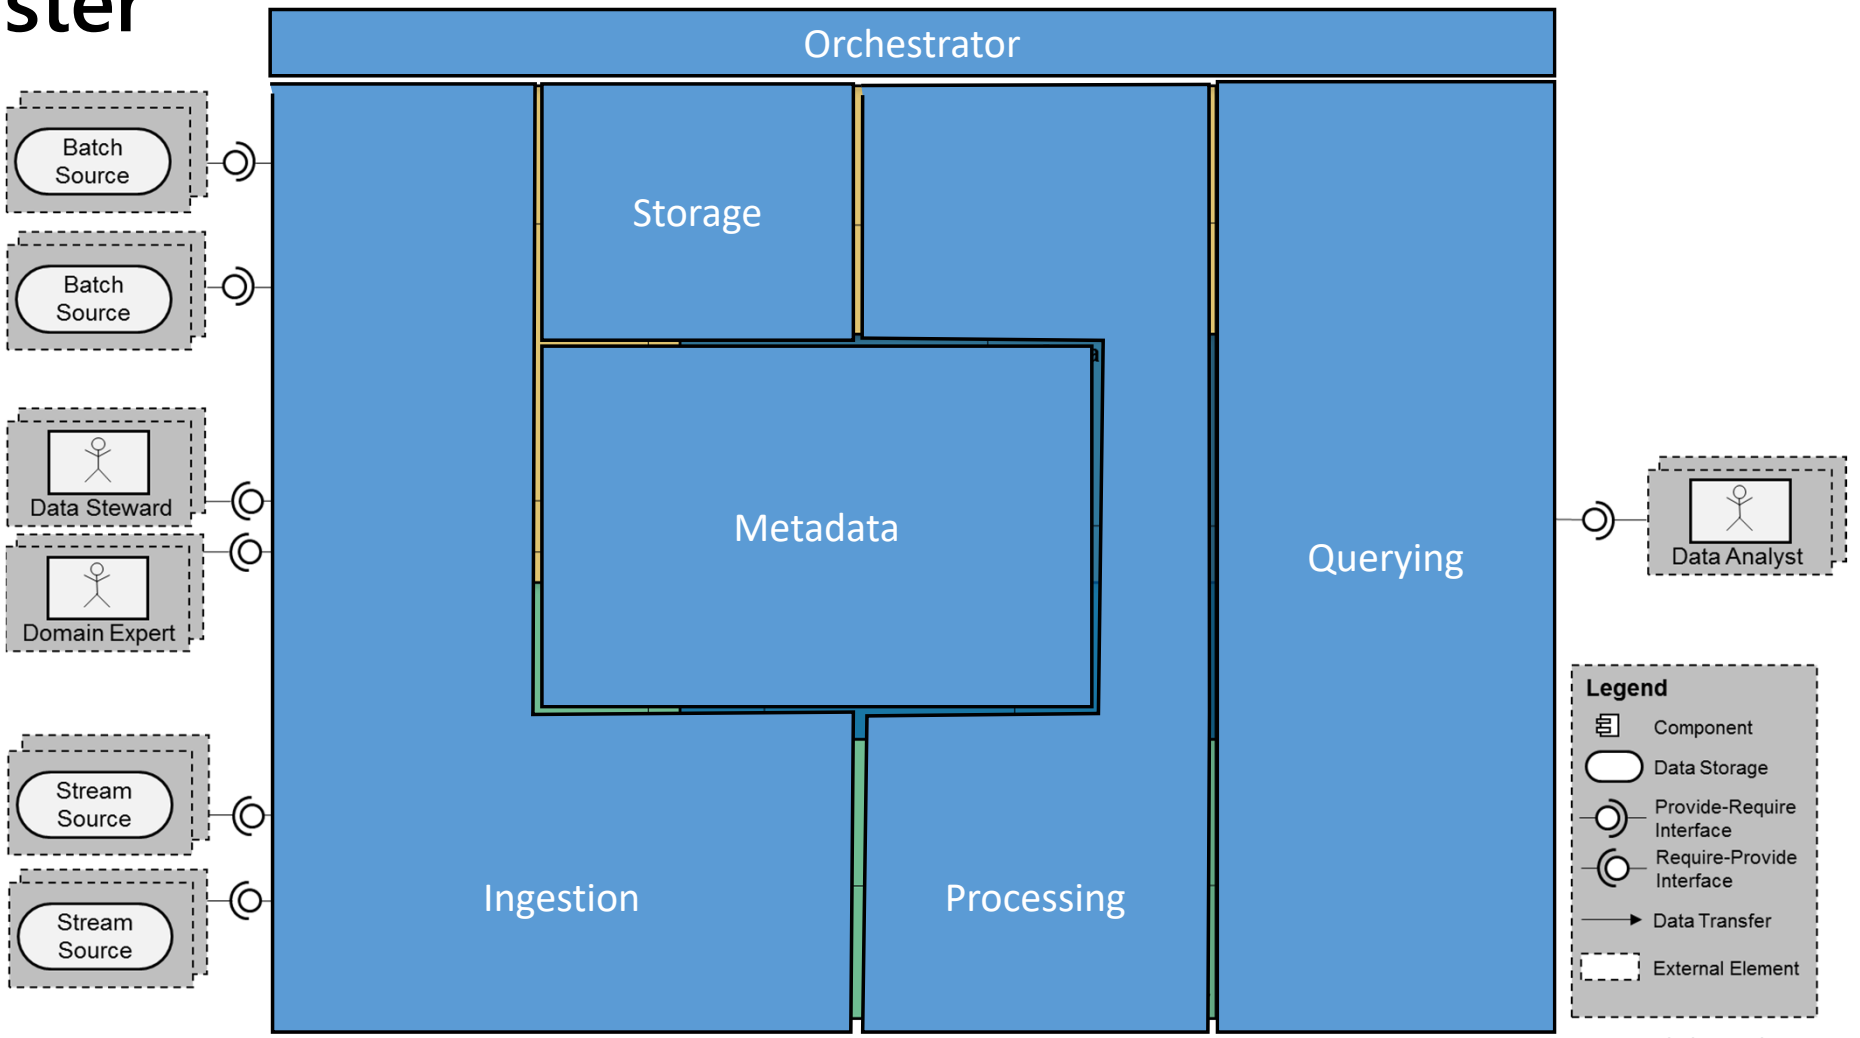

S. Nadal et al.

Instantiation

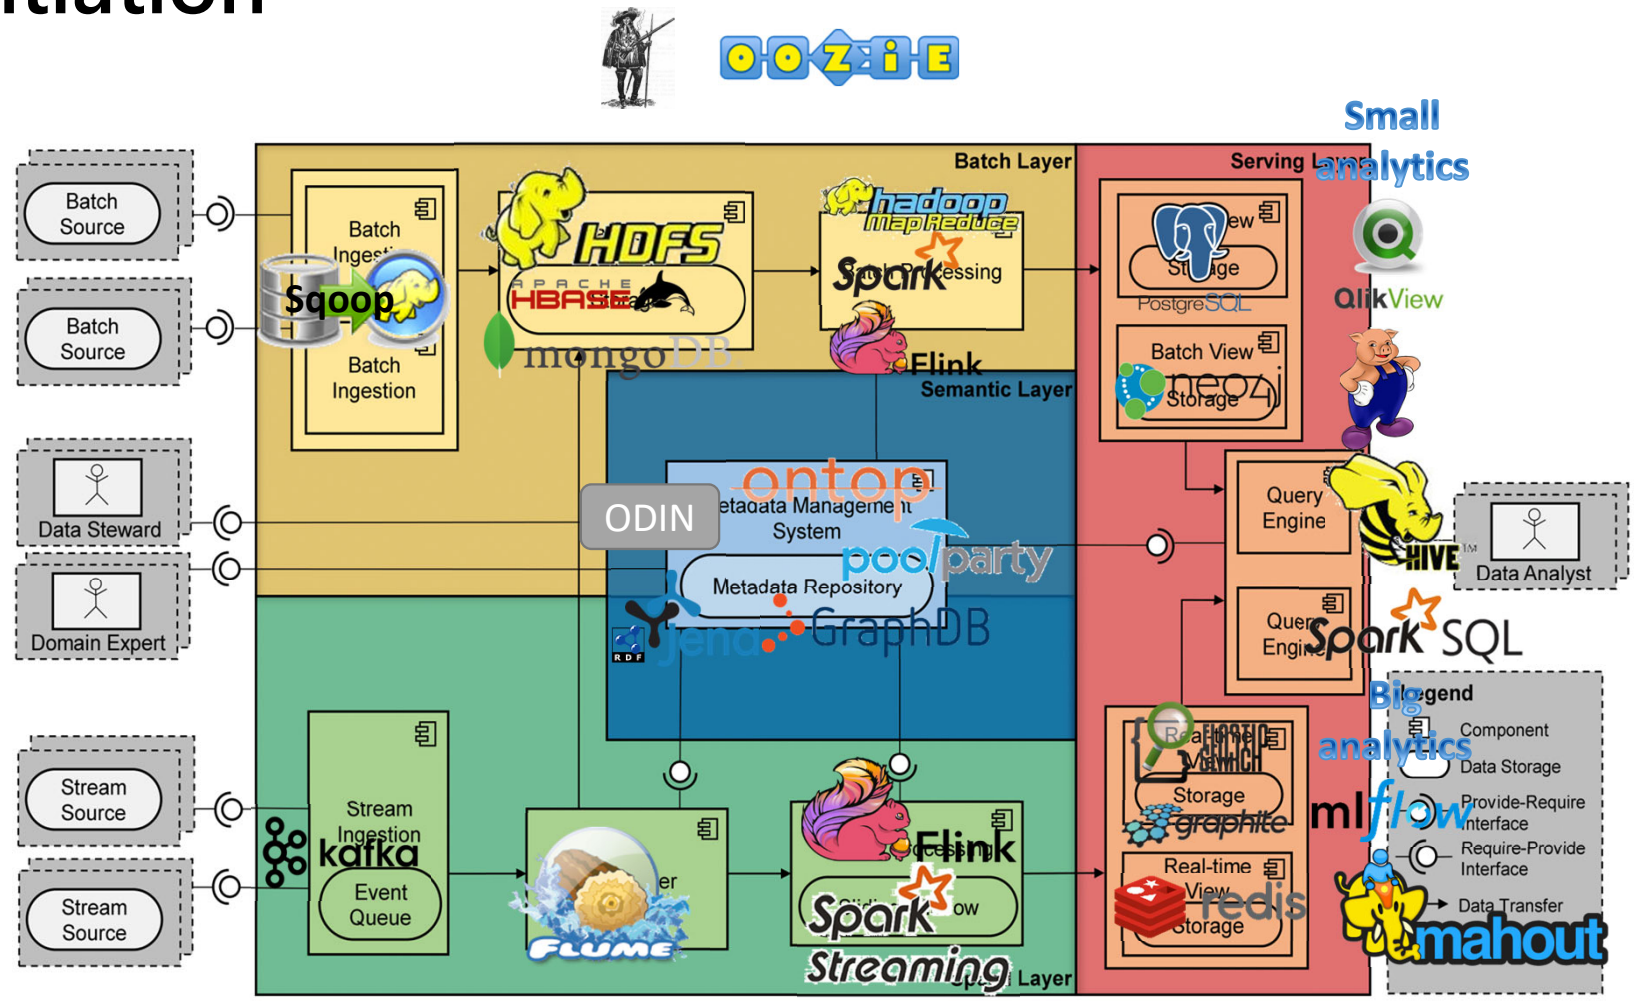

S. Nadal et al.

Workflow Orchestrators

- Current workflow orchestrators are rather poor: Oozie
- But there are attempts for smarter approaches: the ideas behind Musketeer deserves special attention

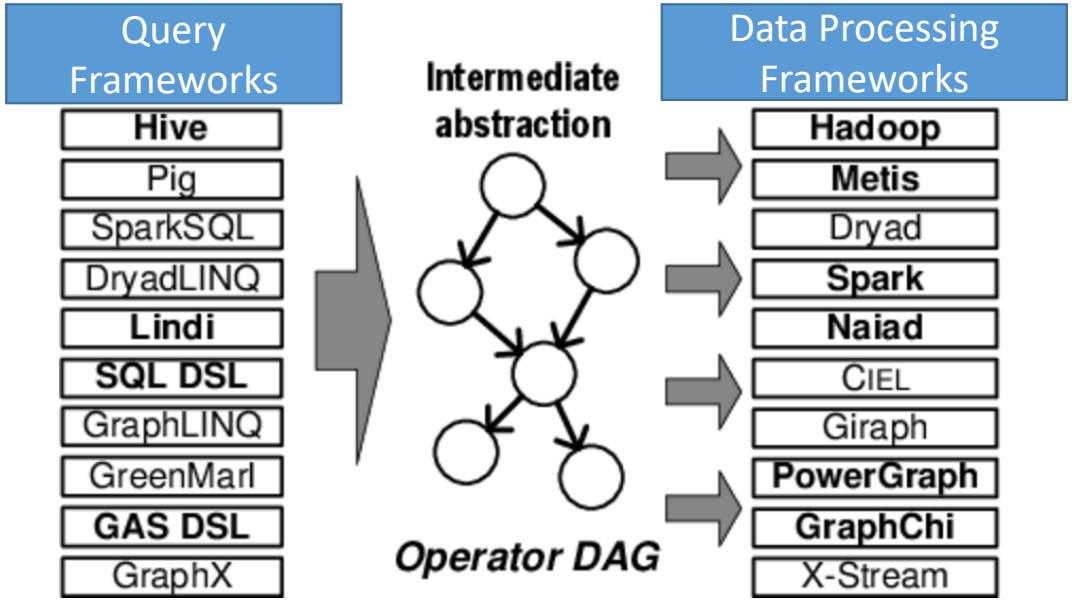

Shortly, does a similar job to global query optimizers of traditional distributed RDBMS

Closing

Summary

- New architectural solutions
  - Lambda
  - Kappa
  - Polyglot systems
- Data Lakes
  - The need of metadata
- Reference architectures
  - Bolster
  - Quarry

References

- T. Erl. *Service-oriented architecture: concepts, technology and design*. Prentice Hall, 2005
- D. McCreary and A. Kelly. *Making Sense of NoSQL*. Manning, 2014
- M. Grover et al. *Hadoop Application Architectures*. O'Reilly, 2015
- N. Marz and J. Warren. *Big Data: Principles and best practices of scalable realtime data systems*. Manning Publications Co., 2015
- S. Nadal et al. *A Software Reference Architecture for Semantic-Aware Big Data Systems*. Information and Software Technology 90. Elsevier, 2017
- S. Nadal et al. *ODIN: A Dataspace Management System*. International Semantic Web Conference 2019
- S. Nadal. Metadata-Driven Data Integration (PhD Thesis). 2019
- P. Jovanovic et al. *Quarry: A User-centered Big Data Integration Platform*. Information Systems Frontier, 2021

Resources

- http://hadoop.apache.org
- http://www.cloudera.com
- http://hortonworks.com